CS 305 Lab Tutorial
Lab 8 TCP Sliding Window & WebSocket

Dept. Computer Science and Engineering
Southern University of Science and Technology

SUSTech
Southern University
of Science and Technology

Thanks to Wei WANG & HHQ ZHANG

Part A.1 Sliding window system(1)

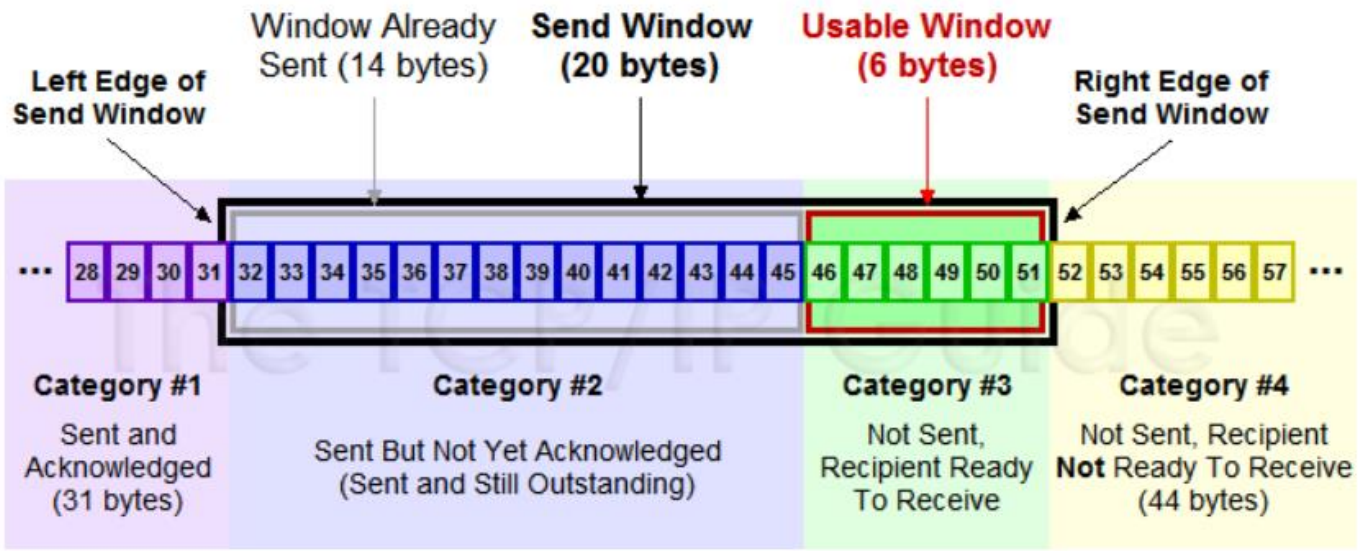

The ***send window*** is the key to the entire TCP sliding window system: it represents the maximum number of unacknowledged bytes a device is allowed to have outstanding at once.
The **usable window** is the number of bytes that the sender is still allowed to send at any point in time; it is equal to the size of the send window less the number of unacknowledged bytes already transmitted.

http://www.TCPipguide.com/free/t_TCPSlidingWindowAcknowledgmentSystemForDataTranspo-6.htm

Sliding window system(2)

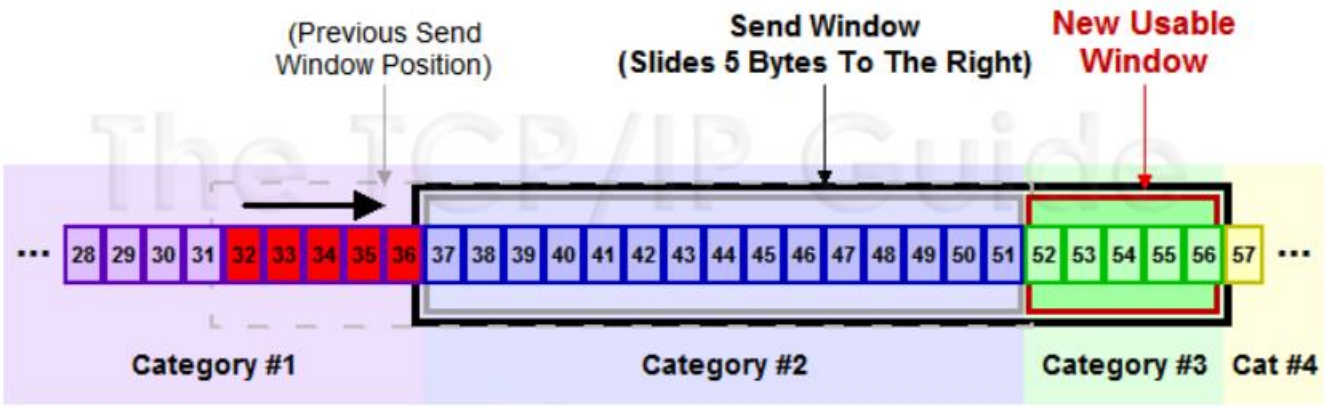

When the sending device **receives new acknowledgment**, it will be able to transfer some of the bytes from Category #2 to Category #1, since they have now been acknowledged. When it does so, something interesting will happen. Since five bytes have been acknowledged, and the window size didn't change, the sender is allowed to send five more bytes. In effect, the window shifts, or ***slides***, over to the right in the timeline.

At the same time five bytes move from Category #2 to Category #1, five bytes move from Category #4 to Category #3, **creating a new usable window for subsequent transmission**.

http://www.TCPipguide.com/free/t_TCPSlidingWindowAcknowledgmentSystemForDataTranspo-6.htm

ACK number, Sequence number, len

No.	Time	Source	Destination	Protocol	Info
94	8.574280	gaia.cs.umass.…	192.168.88.149	TCP	http(80) → 54861 [ACK] Seq=65701 Ack=333 Win=30336 Len=1460
95	8.576343	gaia.cs.umass.…	192.168.88.149	TCP	http(80) → 54861 [ACK] Seq=67161 Ack=333 Win=30336 Len=1460
96	8.576345	gaia.cs.umass.…	192.168.88.149	TCP	http(80) → 54861 [ACK] Seq=68621 Ack=333 Win=30336 Len=1460
97	8.576345	gaia.cs.umass.…	192.168.88.149	TCP	http(80) → 54861 [ACK] Seq=70081 Ack=333 Win=30336 Len=1460
98	8.576516	192.168.88.149	gaia.cs.umass.…	TCP	54861 → http(80) [ACK] Seq=333 Ack=71541 Win=65536 Len=0

```
Source Port: 54861 (54861)
Destination Port: http (80)
[Stream index: 0]
[TCP Segment Len: 0]
Sequence number: 333    (relative sequence number)
[Next sequence number: 333    (relative sequence number)]
Acknowledgment number: 71541    (relative ack number)
0101 .... = Header Length: 20 bytes (5)
Flags: 0x010 (ACK)
   000. .... .... = Reserved: Not set
   ...0 .... .... = Nonce: Not set
   .... 0... .... = Congestion Window Reduced (CWR): Not set
   .... .0.. .... = ECN-Echo: Not set
   .... ..0. .... = Urgent: Not set
   .... ...1 .... = Acknowledgment: Set
   .... .... 0... = Push: Not set
   .... .... .0.. = Reset: Not set
   .... .... ..0. = Syn: Not set
   .... .... ...0 = Fin: Not set
   [TCP Flags: ·····A····]
```

ack_num (71541) = seq (70081) +len (1460)

SUSTech
Southern University of Science and Technology

Changes of window

While the size of usable window turn to be 0, it means the sender will not send any segment at this moment. Wireshark mark the segment with “[Tcp Window Full]”

Seq（135781）+ len（1280） – ack（135781） == win（1280）

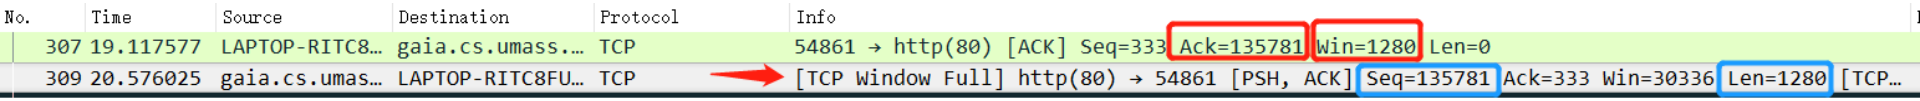

While the recv window turn to be zero, sender will stop to send packet, it will send “[TCP Keep-Alive]” to keep the TCP connection, waiting for the changing of recv window.

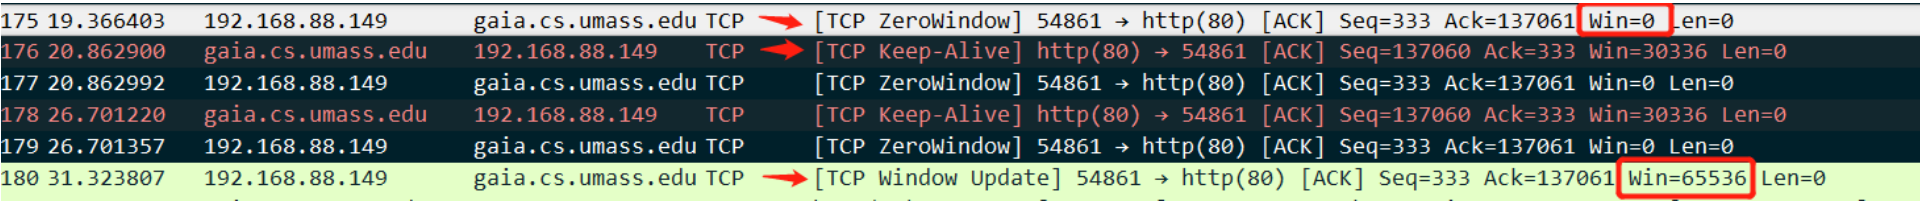

Part A.2 retransimission : GBN

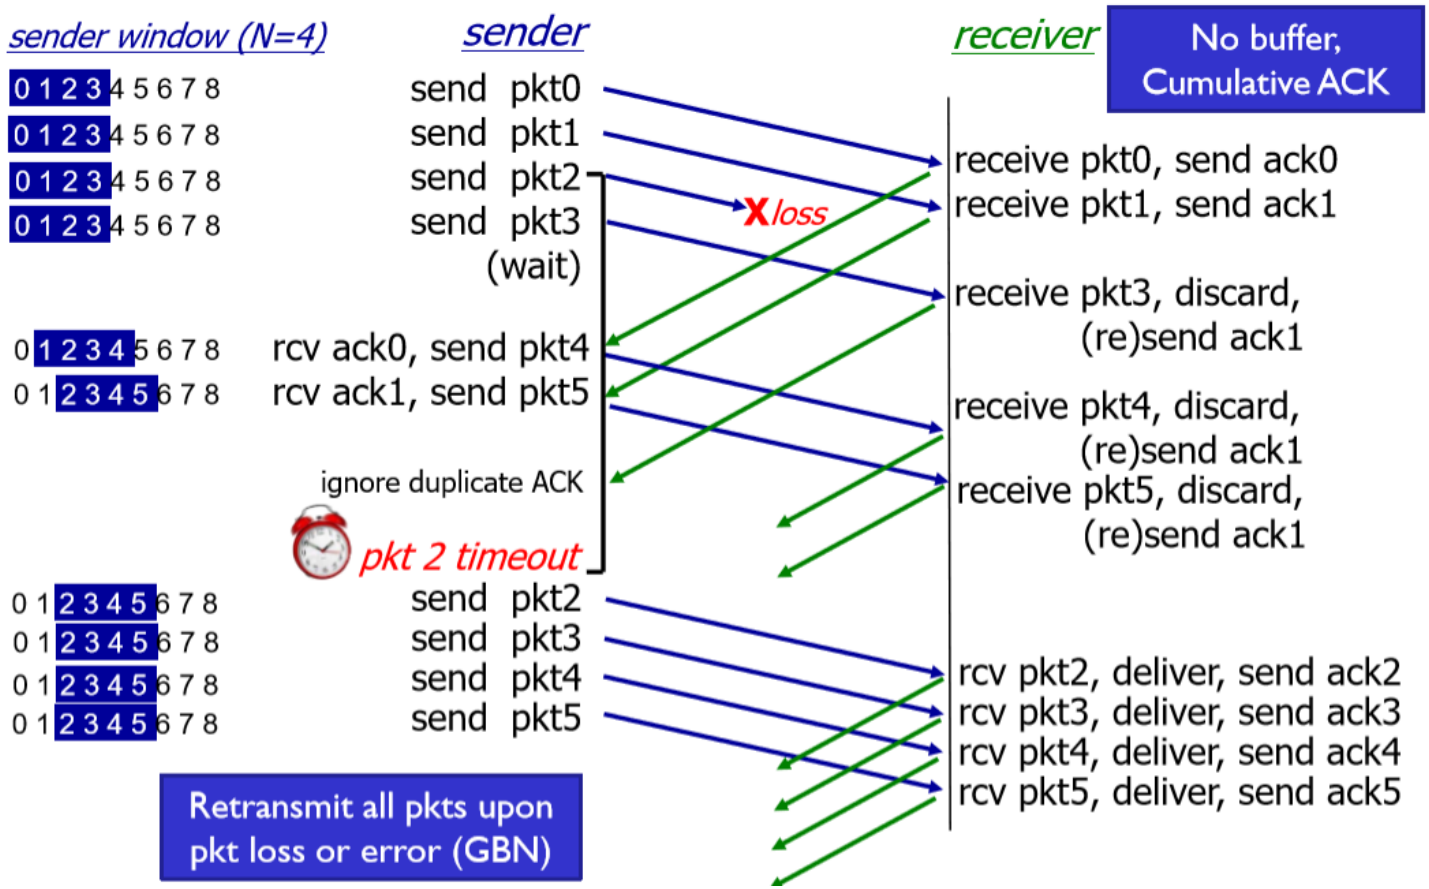

SR

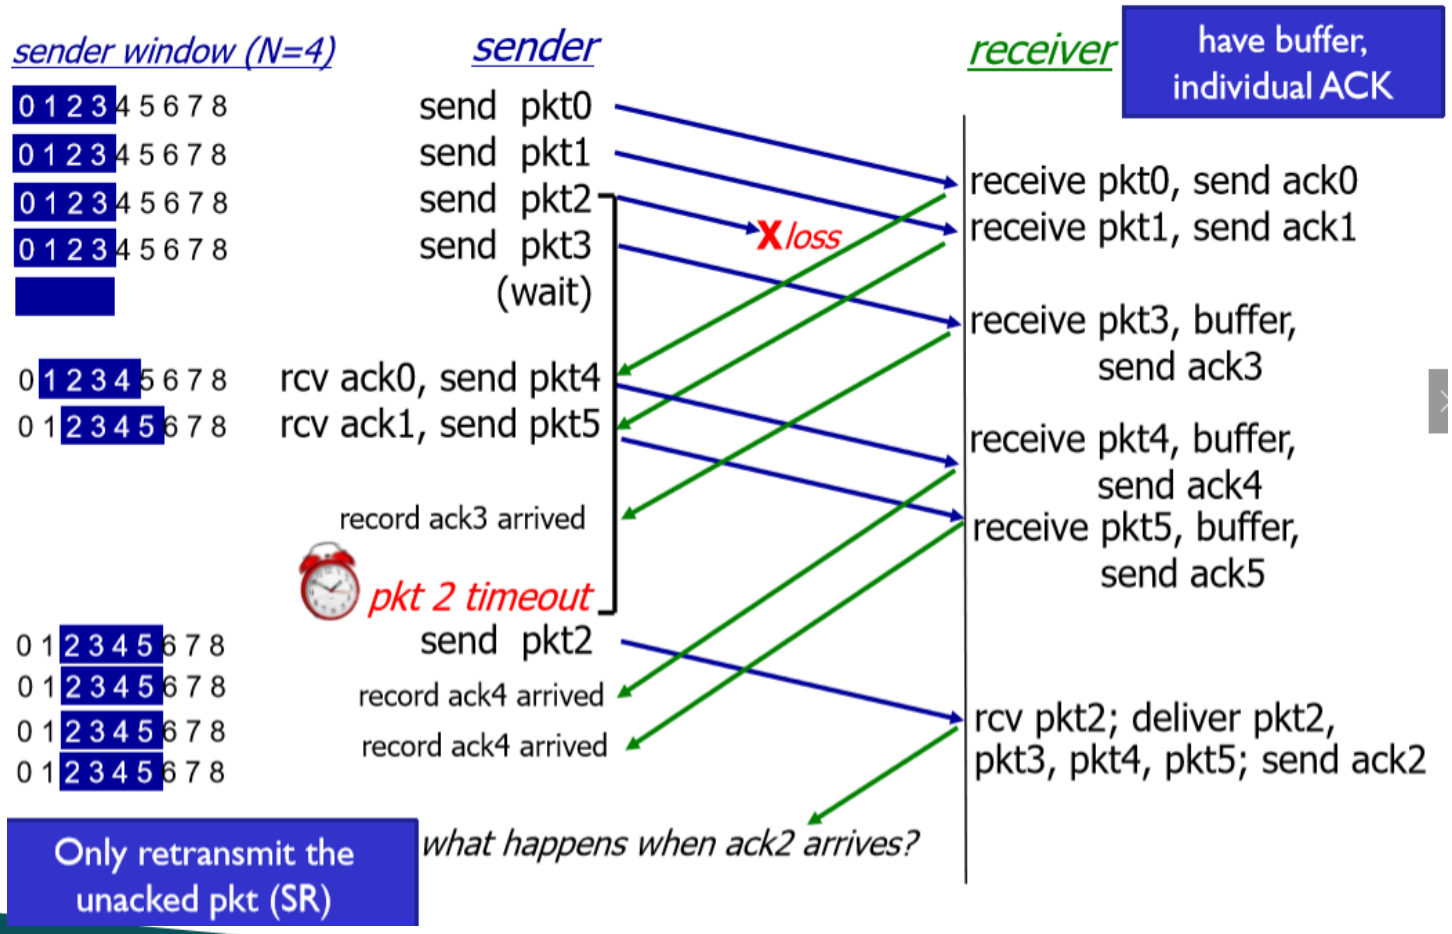
sender window (N=4)
sender
receiver
have buffer, individual ACK
0 1 2 3 4 5 6 7 8
send pkt0
send pkt1
send pkt2
send pkt3
(wait)
loss
receive pkt0, send ack0
receive pkt1, send ack1
receive pkt3, buffer, send ack3
rcv ack0, send pkt4
rcv ack1, send pkt5
receive pkt4, buffer, send ack4
receive pkt5, buffer, send ack5
record ack3 arrived
pkt 2 timeout
send pkt2
record ack4 arrived
record ack4 arrived
rcv pkt2; deliver pkt2, pkt3, pkt4, pkt5; send ack2
what happens when ack2 arrives?
Only retransmit the unacked pkt (SR)

TCP SACK

- A Selective Acknowledgment (**SACK**) mechanism, combined with a **selective repeat retransmission policy**, can help the sender retransmit only the missing data segments.
  - The receiving TCP **sends back SACK packets to the sender** informing the sender of data that has been received.
  - With selective acknowledgments, the data receiver can inform the sender about all segments that have arrived successfully, so the sender need **retransmit only the segments that have actually been lost.**

- The selective acknowledgment extension uses two TCP options:
  - **SACK-permitted** option
  - **SACK** option

https://tools.ietf.org/html/rfc2018

SACK-permitted option

- "**SACK-permitted**", which may **be sent in a SYN segment** to indicate that the SACK option can be used once the connection is established.

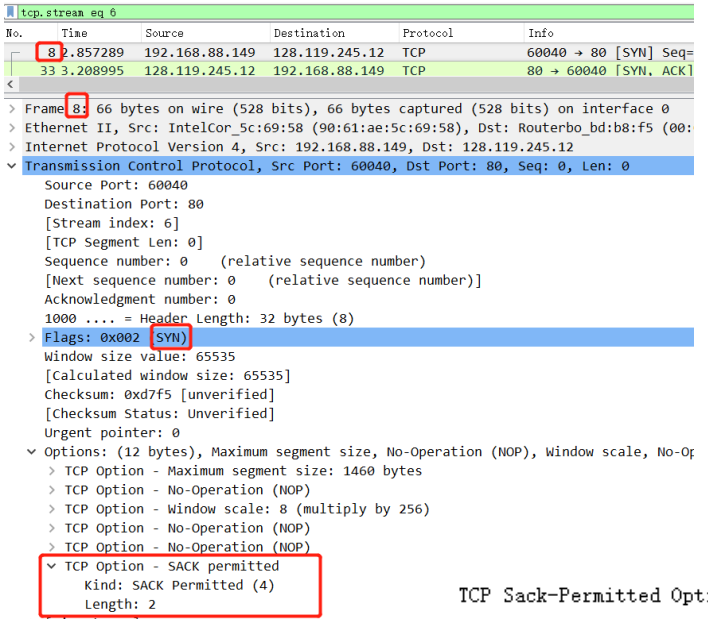

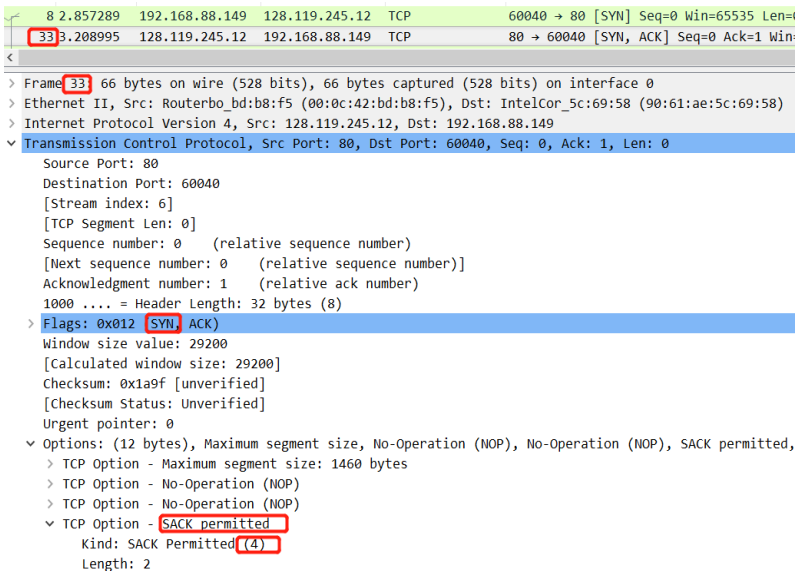

TCP Sack-Permitted Option:

Kind: 4

| Kind=4 | Length=2 |

Wireshark tips: TCP.option_kind==4

SACK option(1)

The **SACK option** is to be **sent by a data receiver to inform the data sender of non-contiguous blocks of data that have been received and queued**.

The data receiver awaits the receipt of data (perhaps by means of retransmissions) to fill the gaps in sequence space between received blocks.

When missing segments are received, **the data receiver acknowledges the data normally** by advancing the left window edge in the Acknowledgement Number Field of the TCP header.

The SACK option does not change the meaning of the Acknowledgement Number field.

This option contains a list of **some of the blocks of contiguous sequence space occupied by data that has been received and queued within the window.**

Each contiguous block of data queued at the data receiver is defined in the SACK option by **two 32-bit unsigned integers** in network byte order:

- **Left Edge** of Block This is the first sequence number of this block.
- **Right Edge** of Block This is the sequence number immediately following the last sequence number of this block.

Each block represents received bytes of data that are **contiguous and isolated**; that is, **the bytes just below the block, (Left Edge of Block - 1), and just above the block, (Right Edge of Block), have NOT been received**.

SACK option(2)

- SACK option , which may be sent over an established connection once permission has been given by SACK-permitted.

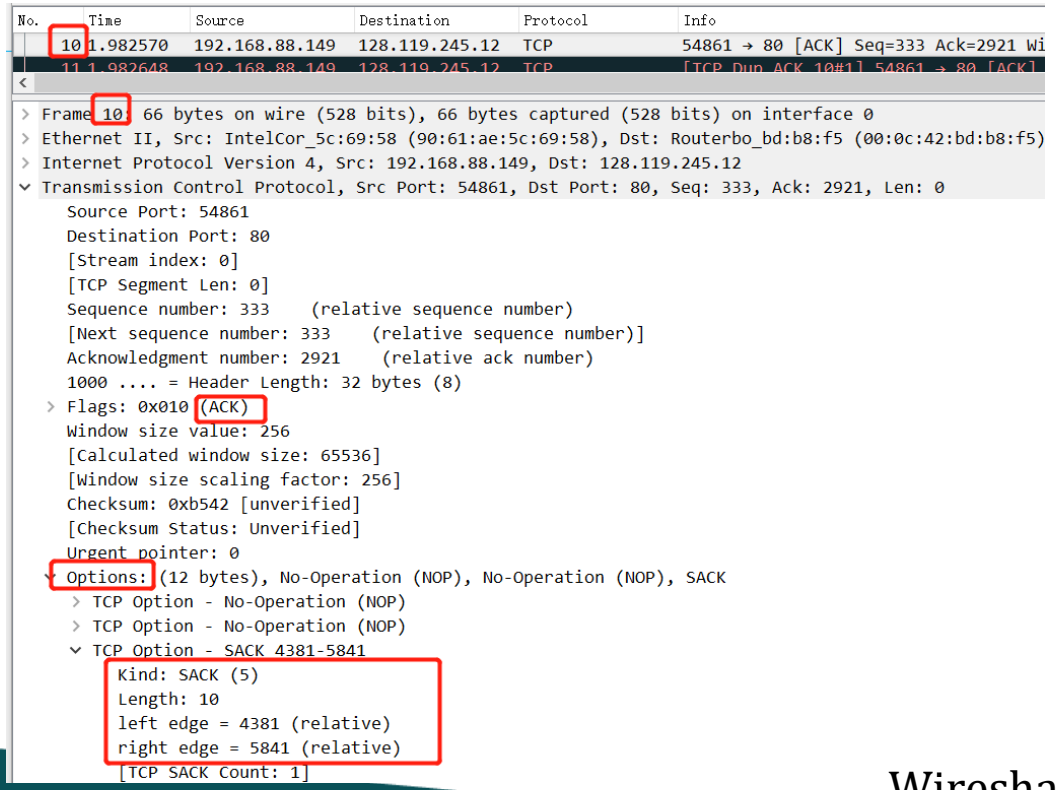

TCP SACK Option:

Kind: 5

Length: Variable

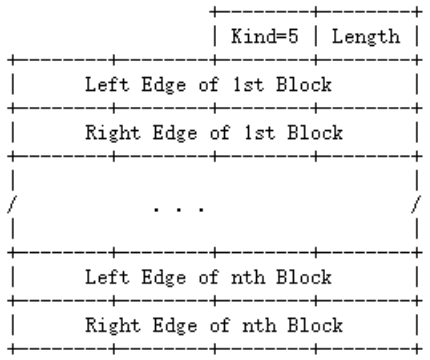

Wireshark tips: TCP.option_kind==5

SUSTech
Southern University of Science and Technology

SACK option(3)

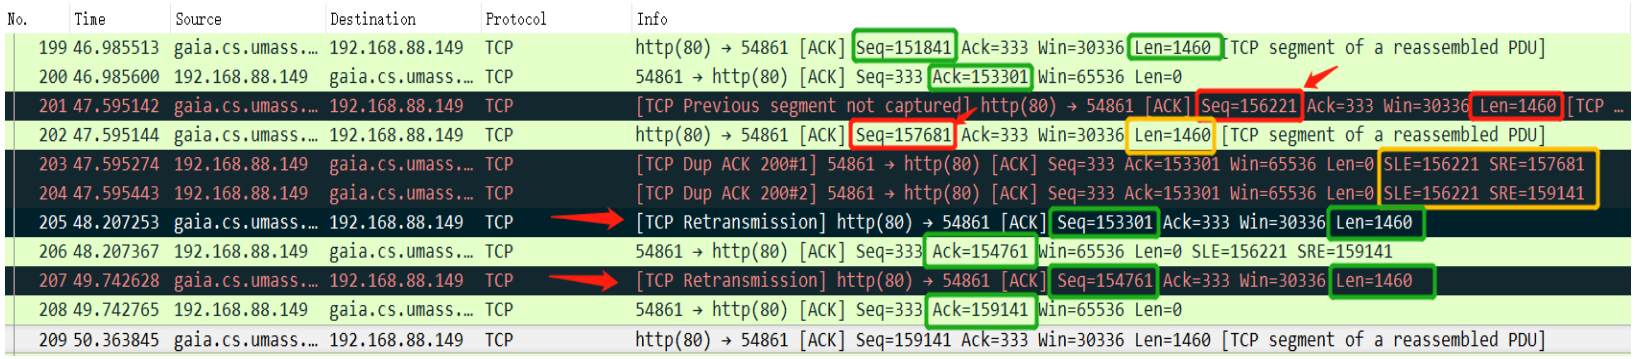

No.	Time	Source	Destination	Protocol	Info
199	46.985513	gaia.cs.umass.…	192.168.88.149	TCP	http(80) → 54861 [ACK] Seq=151841 Ack=333 Win=30336 Len=1460 [TCP segment of a reassembled PDU]
200	46.985600	192.168.88.149	gaia.cs.umass.…	TCP	54861 → http(80) [ACK] Seq=333 Ack=153301 Win=65536 Len=0
201	47.595142	gaia.cs.umass.…	192.168.88.149	TCP	[TCP Previous segment not captured] http(80) → 54861 [ACK] Seq=156221 Ack=333 Win=30336 Len=1460 [TCP …
202	47.595144	gaia.cs.umass.…	192.168.88.149	TCP	http(80) → 54861 [ACK] Seq=157681 Ack=333 Win=30336 Len=1460 [TCP segment of a reassembled PDU]
203	47.595274	192.168.88.149	gaia.cs.umass.…	TCP	[TCP Dup ACK 200#1] 54861 → http(80) [ACK] Seq=333 Ack=153301 Win=65536 Len=0 SLE=156221 SRE=157681
204	47.595443	192.168.88.149	gaia.cs.umass.…	TCP	[TCP Dup ACK 200#2] 54861 → http(80) [ACK] Seq=333 Ack=153301 Win=65536 Len=0 SLE=156221 SRE=159141
205	48.207253	gaia.cs.umass.…	192.168.88.149	TCP	[TCP Retransmission] http(80) → 54861 [ACK] Seq=153301 Ack=333 Win=30336 Len=1460
206	48.207367	192.168.88.149	gaia.cs.umass.…	TCP	54861 → http(80) [ACK] Seq=333 Ack=154761 Win=65536 Len=0 SLE=156221 SRE=159141
207	49.742628	gaia.cs.umass.…	192.168.88.149	TCP	[TCP Retransmission] http(80) → 54861 [ACK] Seq=154761 Ack=333 Win=30336 Len=1460
208	49.742765	192.168.88.149	gaia.cs.umass.…	TCP	54861 → http(80) [ACK] Seq=333 Ack=159141 Win=65536 Len=0
209	50.363845	gaia.cs.umass.…	192.168.88.149	TCP	http(80) → 54861 [ACK] Seq=159141 Ack=333 Win=30336 Len=1460 [TCP segment of a reassembled PDU]

#203 and #204 are SACK
#203 tells that 156221~157681 are **contiguous and isolated**
#204 tells that 156221~159141 are **contiguous and isolated**
#200 tells that the block before 153301 are acked
So
#205 retransmit 153301 ~ 153301+ 1460 -1 ,#206 ack it with 154761
#207 retransmit 154761 ~ 154761+1460 -1
#208 ack it with 159141 (for 157681~159140 are **contiguous but NOT isolated**)

Retransmission(1)

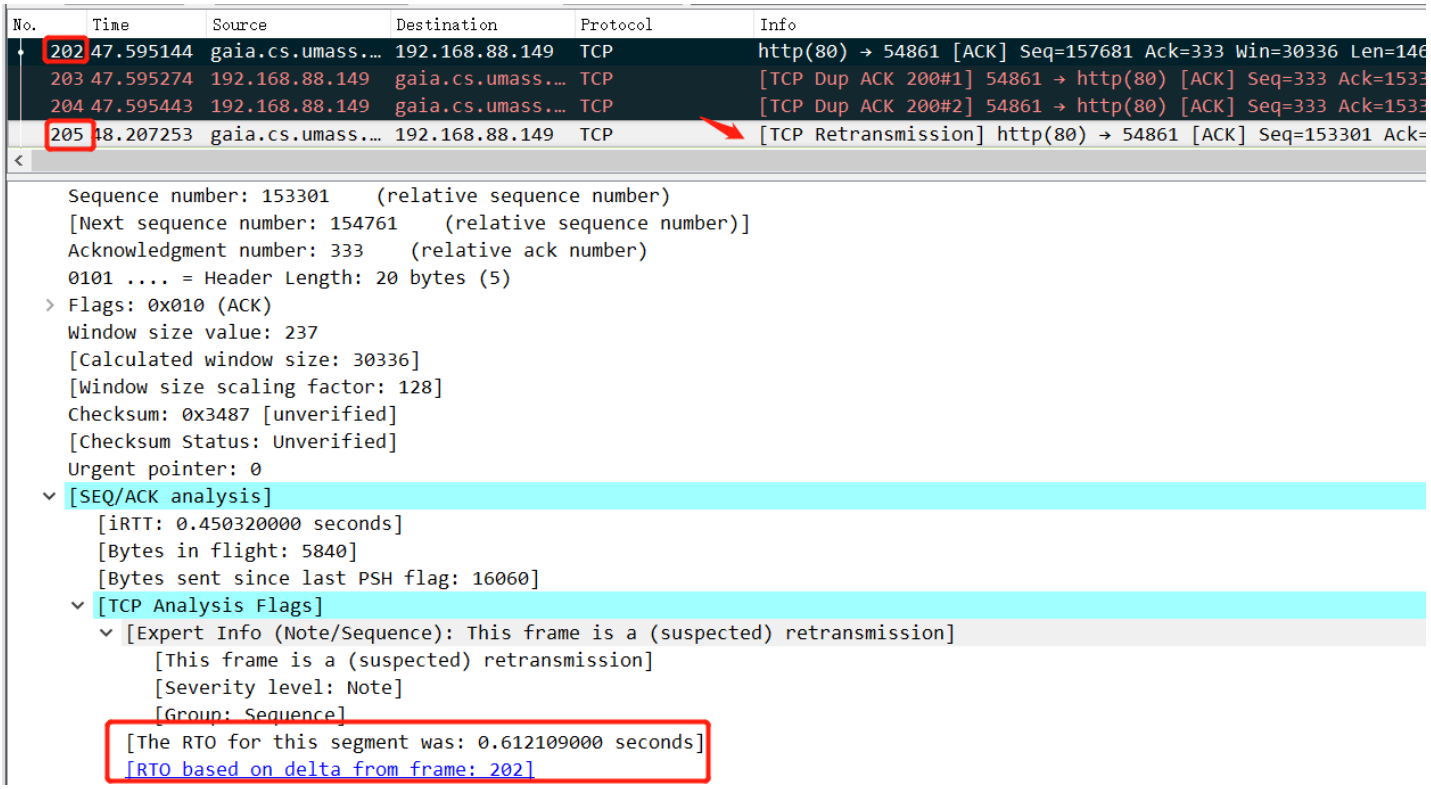

While RTO timeout , retransmission is triggered

Retransmission(2)

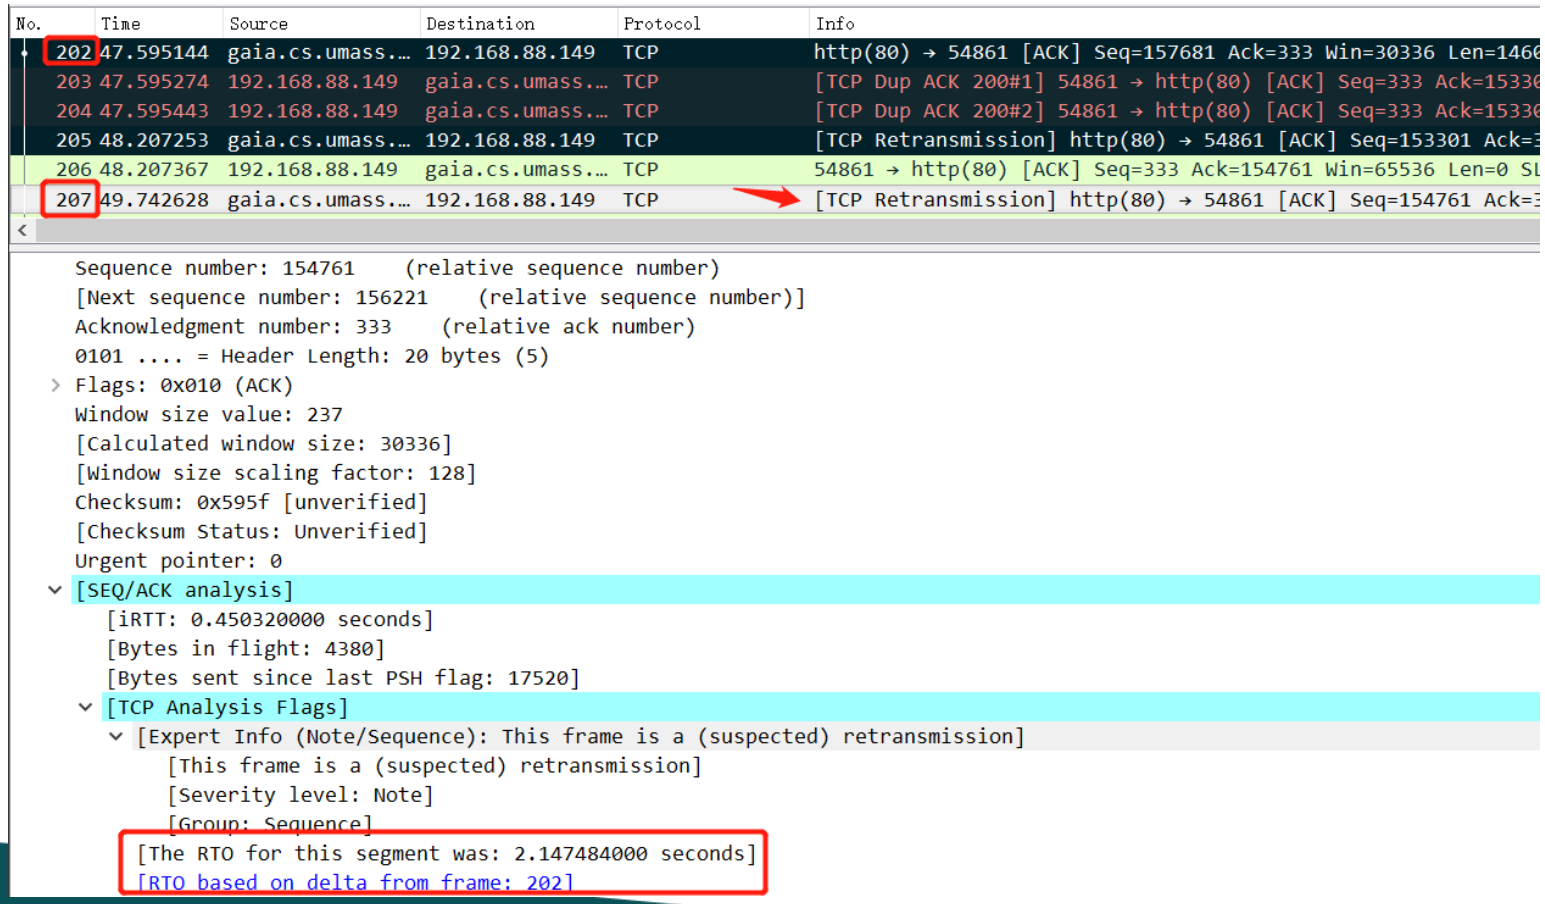

No.	Time	Source	Destination	Protocol	Info
202	47.595144	gaia.cs.umass.…	192.168.88.149	TCP	http(80) → 54861 [ACK] Seq=157681 Ack=333 Win=30336 Len=1460
203	47.595274	192.168.88.149	gaia.cs.umass.…	TCP	[TCP Dup ACK 200#1] 54861 → http(80) [ACK] Seq=333 Ack=15336
204	47.595443	192.168.88.149	gaia.cs.umass.…	TCP	[TCP Dup ACK 200#2] 54861 → http(80) [ACK] Seq=333 Ack=15336
205	48.207253	gaia.cs.umass.…	192.168.88.149	TCP	[TCP Retransmission] http(80) → 54861 [ACK] Seq=153301 Ack=
206	48.207367	192.168.88.149	gaia.cs.umass.…	TCP	54861 → http(80) [ACK] Seq=333 Ack=154761 Win=65536 Len=0 SL
207	49.742628	gaia.cs.umass.…	192.168.88.149	TCP	[TCP Retransmission] http(80) → 54861 [ACK] Seq=154761 Ack=

```
Sequence number: 154761    (relative sequence number)
[Next sequence number: 156221    (relative sequence number)]
Acknowledgment number: 333    (relative ack number)
0101 .... = Header Length: 20 bytes (5)
> Flags: 0x010 (ACK)
Window size value: 237
[Calculated window size: 30336]
[Window size scaling factor: 128]
Checksum: 0x595f [unverified]
[Checksum Status: Unverified]
Urgent pointer: 0
v [SEQ/ACK analysis]
    [iRTT: 0.450320000 seconds]
    [Bytes in flight: 4380]
    [Bytes sent since last PSH flag: 17520]
  v [TCP Analysis Flags]
    v [Expert Info (Note/Sequence): This frame is a (suspected) retransmission]
        [This frame is a (suspected) retransmission]
        [Severity level: Note]
        [Group: Sequence]
      [The RTO for this segment was: 2.147484000 seconds]
      [RTO based on delta from frame: 202]
```

SUSTech
Southern University
of Science and Technology

Fast retransmission

- TCP may generate an immediate acknowledgment (a duplicate ACK) when an out-of-order segment is received. This duplicate ACK should not be delayed. The purpose of this duplicate ACK is to let the other end know that a segment was received out of order, and to tell it what sequence number is expected.
- Since TCP does not know whether a duplicate ACK is caused by a lost segment or just a reordering of segments, it waits for a small number of duplicate ACKs to be received.
  - It is assumed that if there is just a reordering of the segments, there will be only one or two duplicate ACKs before the reordered segment is processed, which will then generate a new ACK.
  - If three or more duplicate ACKs are received in a row, it is a strong indication that a segment has been lost.
- TCP then performs a retransmission of what appears to be the missing segment, without waiting for a retransmission timer to expire.

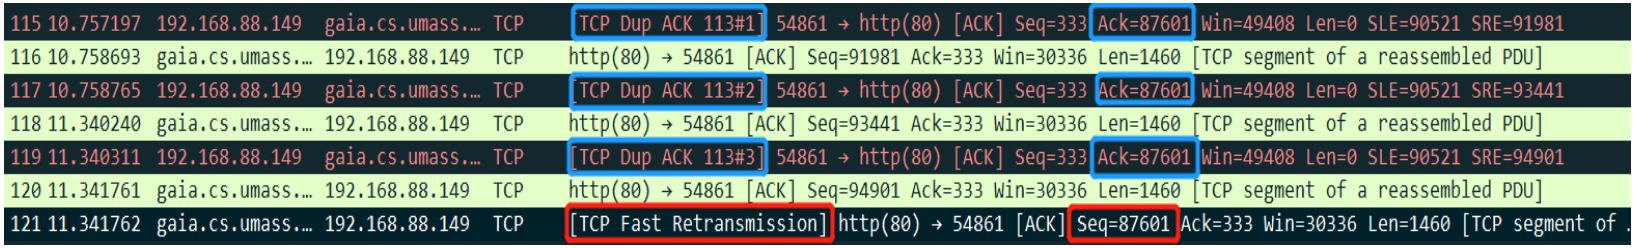

https://tools.ietf.org/html/rfc2001

Part B. WebSocket

- The WebSocket Protocol is designed to supersede existing bidirectional communication technologies that use HTTP as a transport layer to benefit from existing infrastructure (proxies, filtering, authentication).
- The WebSocket Protocol attempts to address the goals of existing bidirectional HTTP technologies in the context of the existing HTTP infrastructure.
  - support HTTP proxies and intermediaries
  - does not limit WebSocket to HTTP, and future implementations could use a simpler handshake over a dedicated port without reinventing the entire protocol.

HTTP polls for new messages

- Historically, creating web applications that need bidirectional communication between a client and a server (e.g., instant messaging and gaming applications) has required an abuse of HTTP to poll the server for updates while sending upstream notifications as distinct HTTP calls .

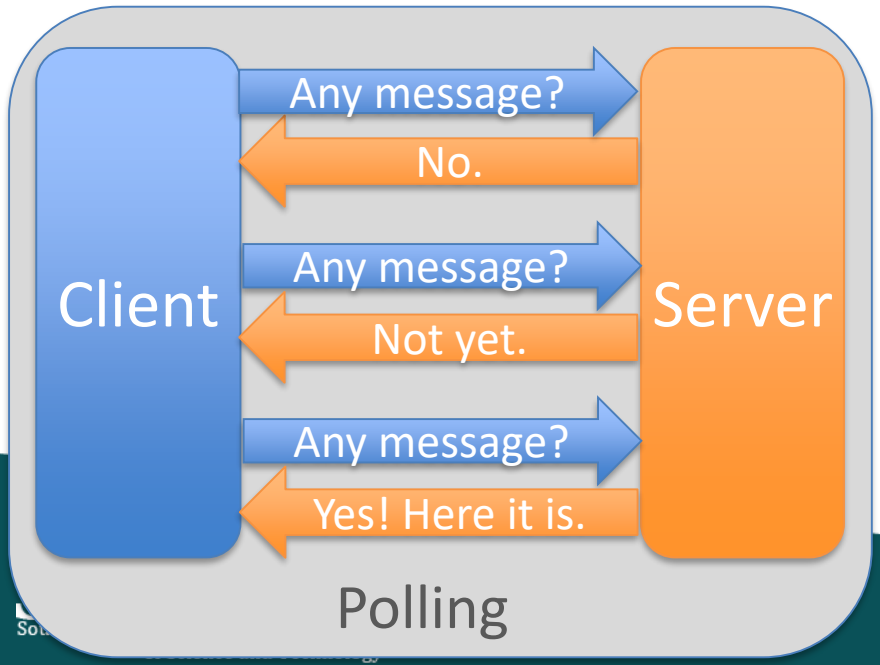

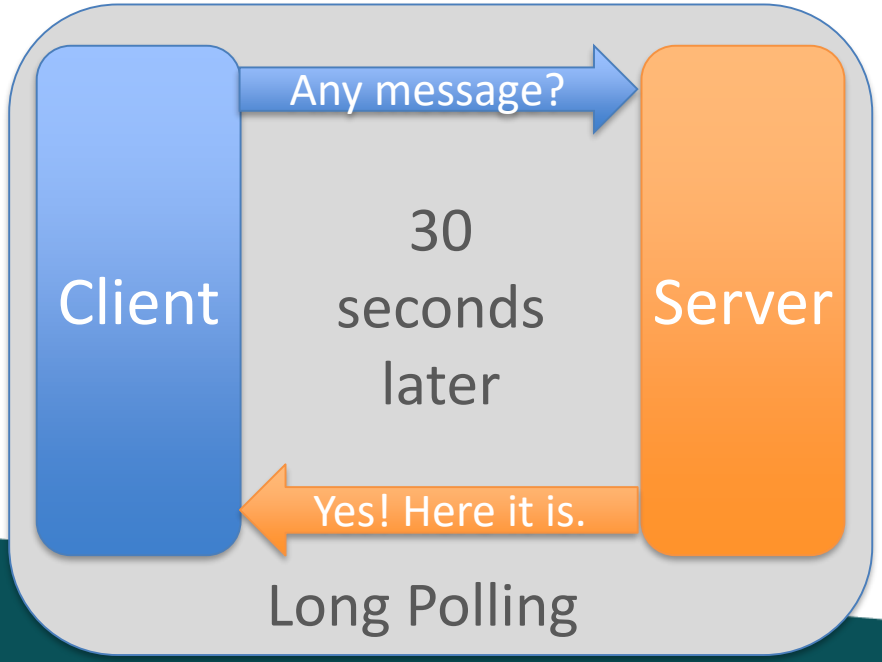

WebSocket protocol overview

- Two parts
  - Handshakes: Opening Handshake & Closing Handshake
  - Data transfer
- The WebSocket Protocol is an independent TCP-based protocol. Its only relationship to HTTP is that its handshake is interpreted by HTTP servers as an Upgrade request.
  - Uses port 80 for regular WebSocket connections
  - Uses port 443 for WebSocket connections tunneled over TLS
  - Can not establish a connection with servers of pre-existing protocols like SMTP and HTTP

URI

- A Uniform Resource Identifier (URI) is a compact sequence of characters that identifies an abstract or physical resource.
- ws-URI = "ws:" "//" host [":" port] path ["?" query]
  - ws://example.com/chat
- wss-URI = "wss:" "//" host [":" port] path ["?" query]

Opening handshake

```
GET /chat HTTP/1.1
Host: server.example.com
Upgrade: websocket
Connection: Upgrade
Sec-WebSocket-Key: dGhlIHNhbXBsZSBub25jZQ==
Origin: http://example.com
Sec-WebSocket-Protocol: chat, superchat
Sec-WebSocket-Version: 13
```

client

example.com

"dGhlIHNhbXBsZSBub25jZQ=="

Step1. concatenate "258EAFA5-E914-47DA-95CA-C5AB0DC85B11"

Step2. take the SHA-1 hash

Step3. base64-encoded the hash value

```
HTTP/1.1 101 Switching Protocols
Upgrade: websocket
Connection: Upgrade
Sec-WebSocket-Accept:
s3pPLMBiTxaQ9kYGzzhZRbK+xOo=
Sec-WebSocket-Protocol: chat
```

Data framing

Indicates that this is the final fragment in a message.

Defines whether the "Payload data" is masked.

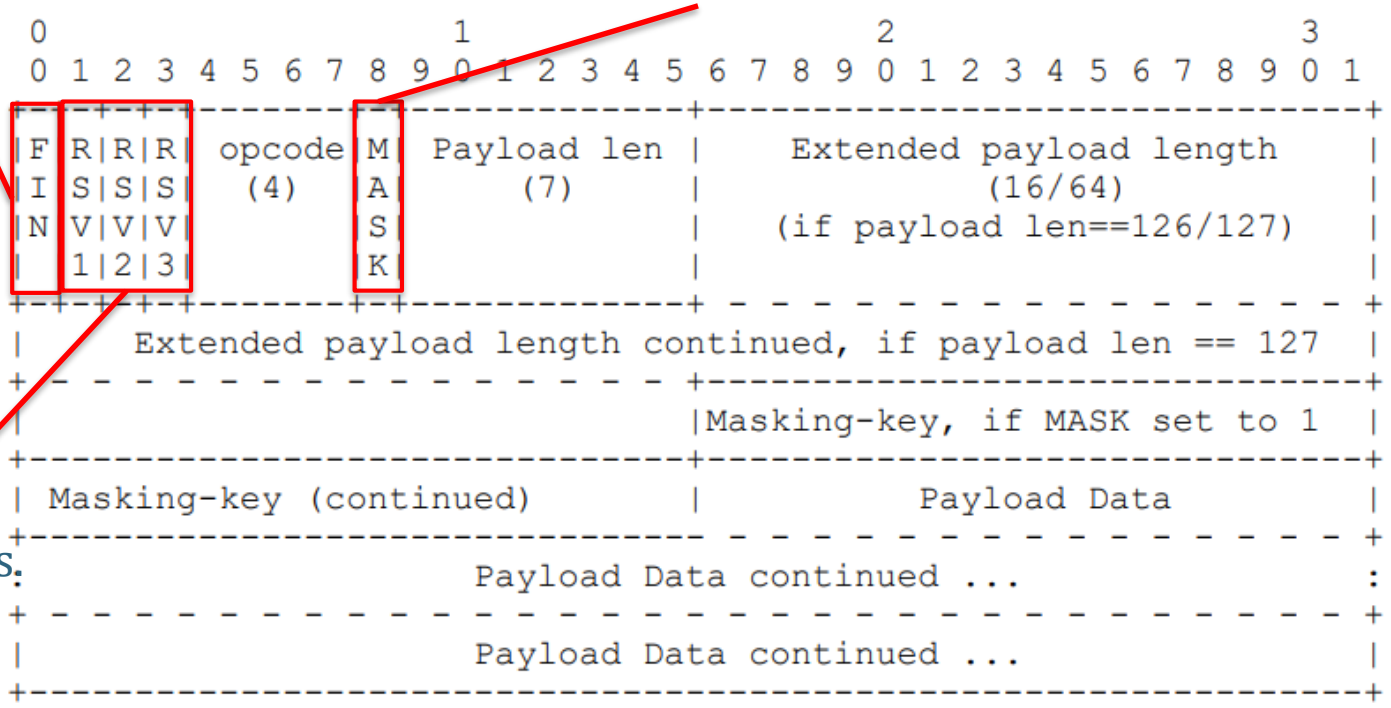

MUST be 0 unless an extension is negotiated that defines meanings for non-zero values.

SUSTech
Southern University of Science and Technology

Data framing (continued)

- Opcode (4 bits)
  - %x0(0000) denotes a continuation frame
  - %x1(0001) denotes a text frame
  - %x2(0010) denotes a binary frame
  - %x3-7(0011-0111) are reserved for further non-control frames
  - %x8(1000) denotes a connection close
  - %x9(1001) denotes a ping
  - %xA(1010) denotes a pong
  - %xB-F(1011-1111) are reserved for further control frames

Data framing (continued)

- Payload len
  - 7 bits: if 0~125 bytes
  - 7 + 16 bit: if the 7 bits equals to 126
  - 7 + 64 bits: if the 7 bits equals to 127
  - payload length = the length of the "Extension data" + the length of the "Application data"
- Masking key
  - 0 bits: if the mask bit is set to 1
  - 32 bits: if the mask bit is set to 0
  - The masking key is a 32-bit value chosen at random by the client.

Data framing examples

- A single-frame unmasked text message
  - 0x81 0x05 0x48 0x65 0x6c 0x6c 0x6f (contains "Hello")
- A single-frame masked text message
  - 0x81 0x85 0x37 0xfa 0x21 0x3d 0x7f 0x9f 0x4d 0x51 0x58 (contains "Hello")
- A fragmented unmasked text message
  - 0x01 0x03 0x48 0x65 0x6c (contains "Hel")
  - 0x80 0x02 0x6c 0x6f (contains "lo")
- Unmasked Ping request and masked Ping response
  - 0x89 0x05 0x48 0x65 0x6c 0x6c 0x6f (contains a body of "Hello", but the contents of the body are arbitrary)
  - 0x8a 0x85 0x37 0xfa 0x21 0x3d 0x7f 0x9f 0x4d 0x51 0x58 (contains a body of "Hello", matching the body of the ping)
- 256 bytes binary message in a single unmasked frame
  - 0x82 0x7E 0x0100 [256 bytes of binary data]

Closing the connection

- An endpoint MUST send a Close control frame(opcode = 1000)
- Connection Close Code
  - 1000: a normal closure.
  - 1001: an endpoint is "going away".
  - 1002: a protocol error occurs.
  - 1003: an endpoint has received a type of data it cannot accept.
  - 1004, 1005, 1006, 1015: Reserved.
  - 1007: an endpoint has received an inconsistent type of data.
  - 1008: an endpoint has received a message violates its policy.
  - 1009: an endpoint has received a message that is too big to process.
  - 1010: the client has expected the server to negotiate one or more extensions but received no response about that.
  - 1011: the server encountered an unexpected condition.

Example 1: Mimic a WebSocket Server

```
import asyncio
import websockets

async def echo(websocket, path):
    async for message in websocket:
        message = "I got your message: {}".format(message)
        await websocket.send(message)

asyncio.get_event_loop().run_until_complete(
    websockets.serve(echo, '127.0.0.1', 8766))
asyncio.get_event_loop().run_forever()
```

Example 2: Mimic a WebSocket Client

```
import asyncio
import websockets

async def echo(uri):
    async with websockets.connect(uri) as websocket:
        while True:
            message = input("Write down your message:")
            await websocket.send(message)
            print("<", message)
            recv_text = await websocket.recv()
            print("> {}".format(recv_text))

asyncio.get_event_loop().run_until_complete(
    echo('ws://127.0.0.1:8766'))
```

Example 3: Use curl

> curl --include --no-buffer --header "Connection: Upgrade" --header "Upgrade: websocket" --header "Host: example.com:80" --header "Origin: http://example.com:80" --header "Sec-WebSocket-Key: GVsbG8sIHdvcmxkIQ==" --header "Sec-WebSocket-Version: 13" http://example.com:80/

```
C:\Users\wq>curl --include --no-buffer --header "Connection: Upgrade" --header "Upgrade: w
ebsocket" --header "Host: example.com:80" --header "Origin: http://example.com:80"  --head
er "Sec-WebSocket-Key: GVsbG8sIHdvcmxkIQ==" --header "Sec-WebSocket-Version: 13" http://ex
ample.com:80/
HTTP/1.1 200 OK
Accept-Ranges: bytes
Cache-Control: max-age=604800
Content-Type: text/html; charset=UTF-8
Date: Tue, 05 Apr 2022 03:56:49 GMT
Etag: "3147526947"
Expires: Tue, 12 Apr 2022 03:56:49 GMT
Last-Modified: Thu, 17 Oct 2019 07:18:26 GMT
Server: EOS (vny/0454)
Content-Length: 1256
Connection: close

<!doctype html>
<html>
<head>
```

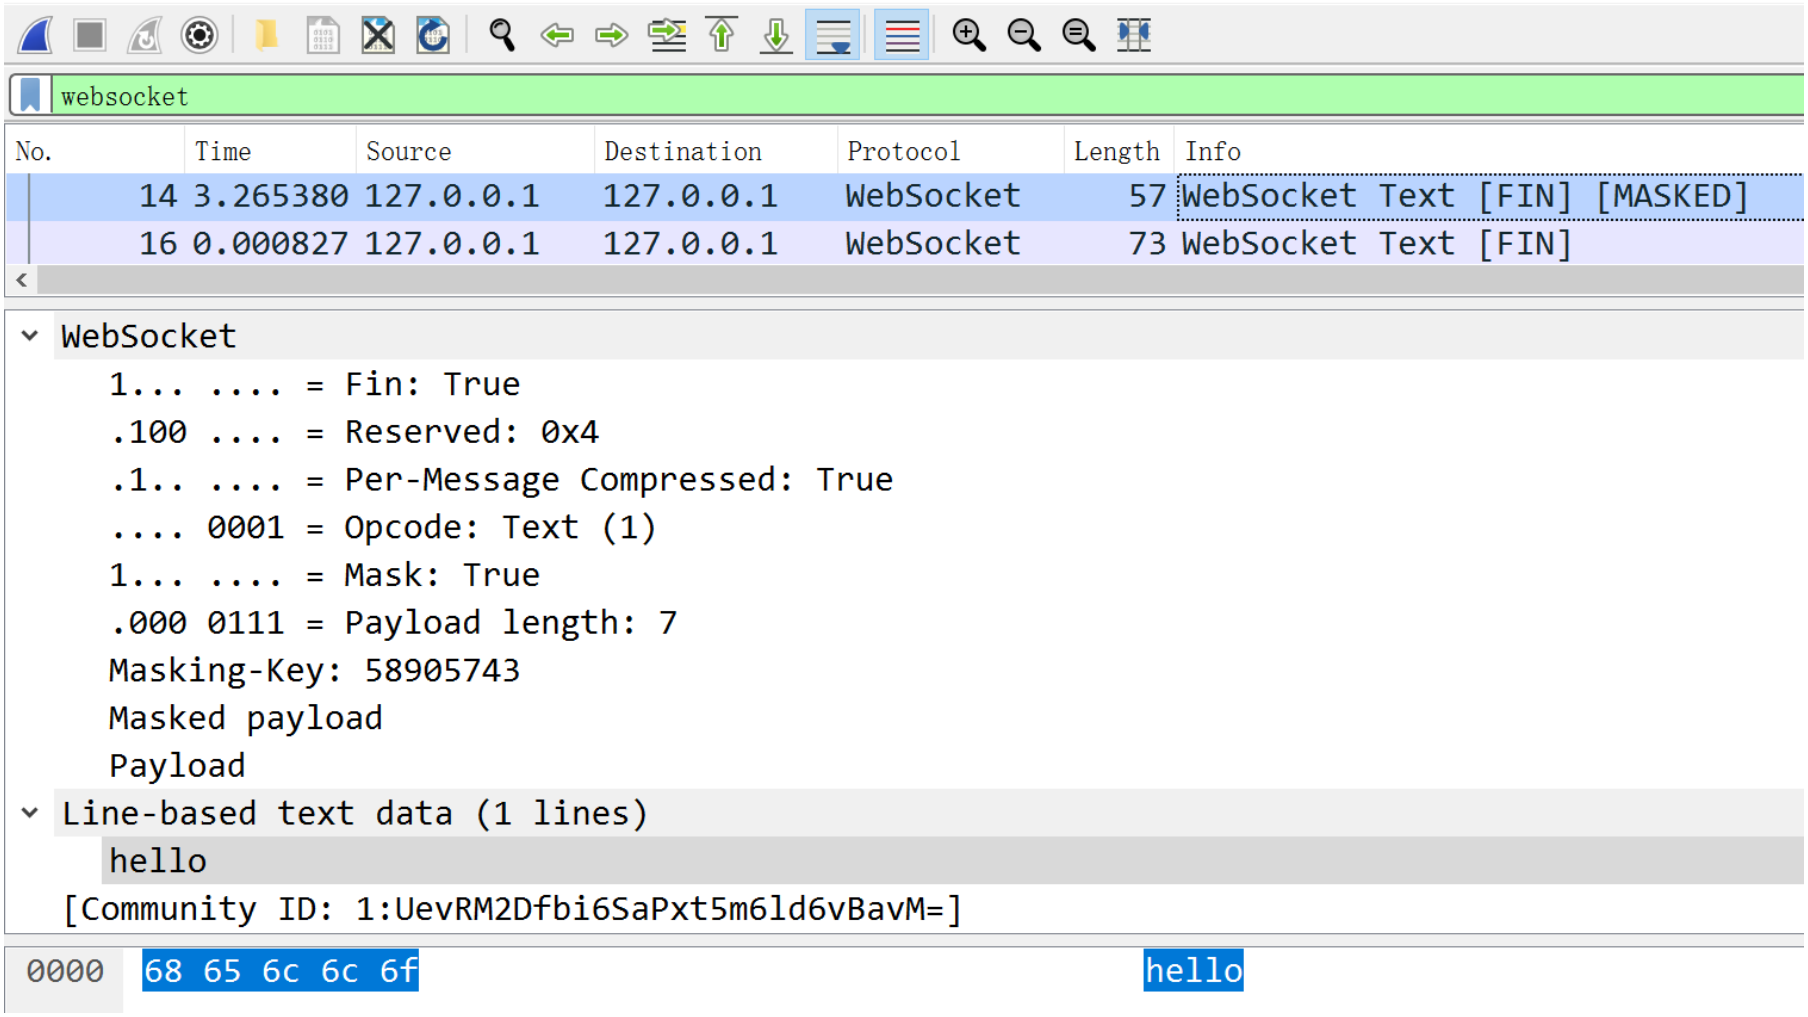

SUSTech
Southern University of Science and Technology

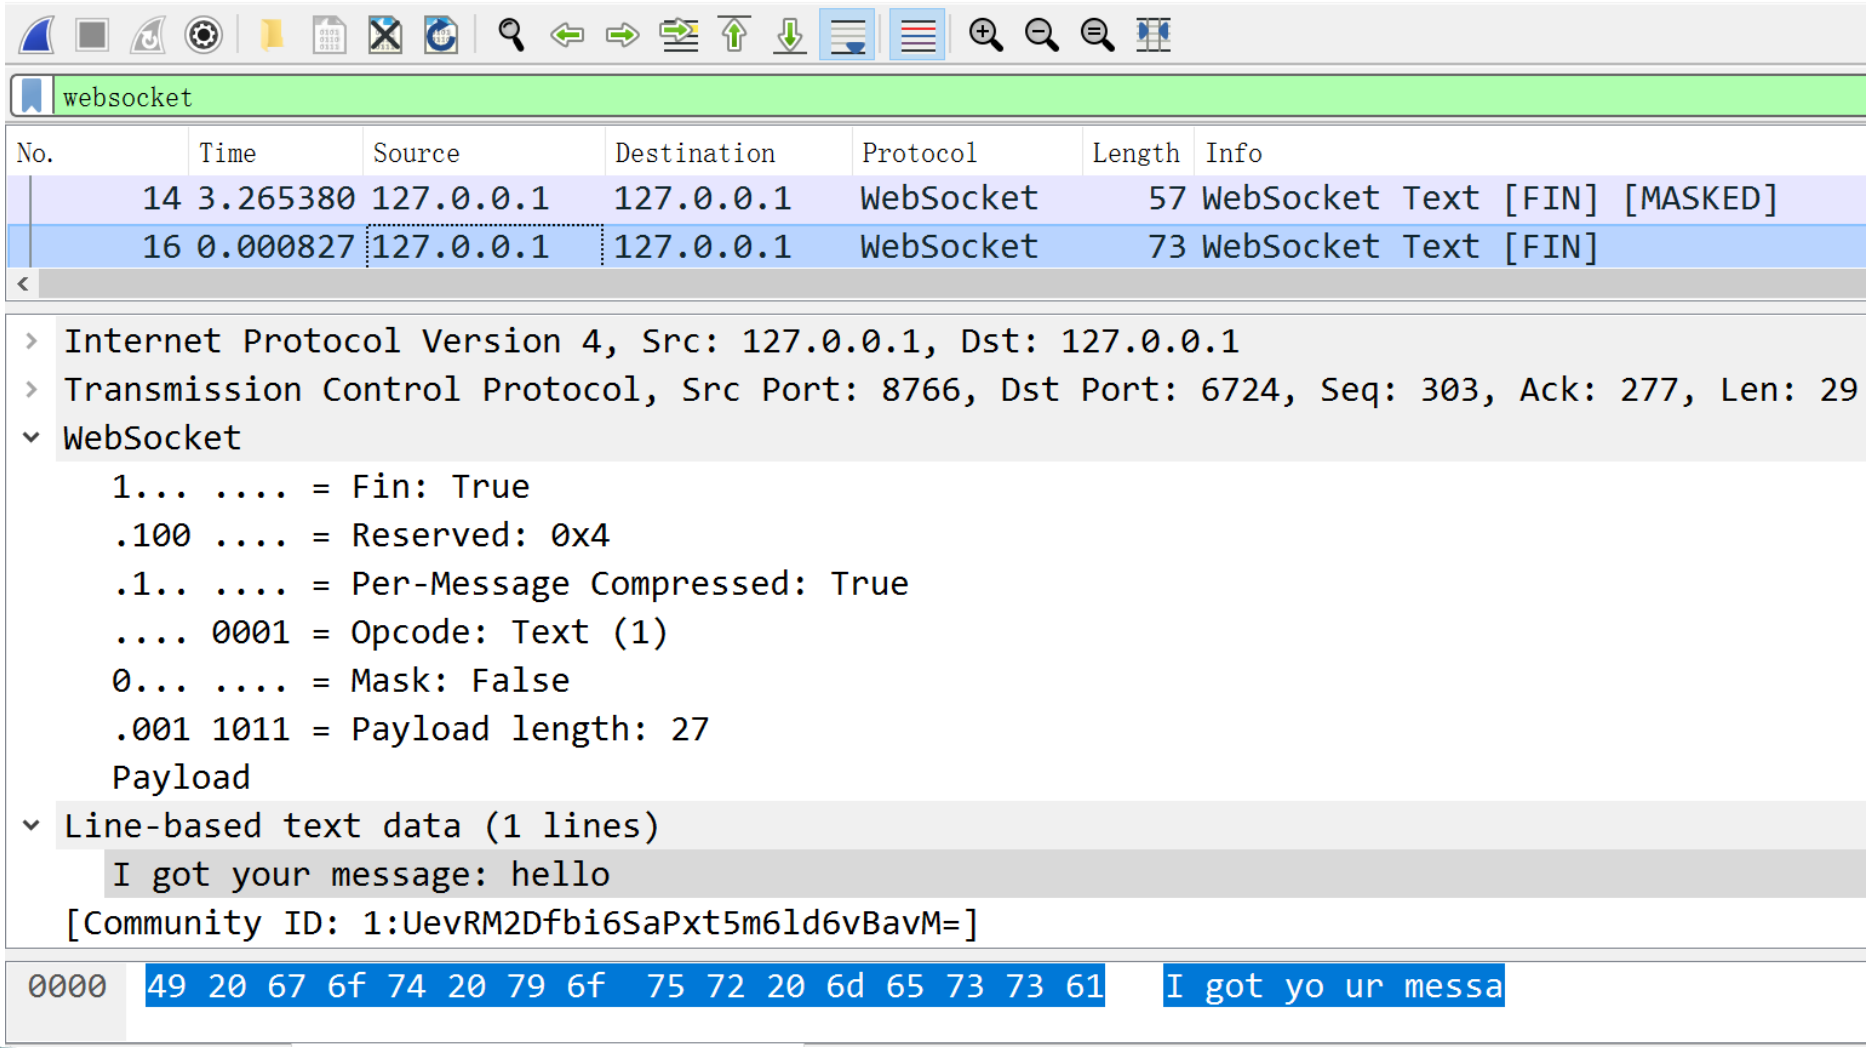

SUSTech
Southern University of Science and Technology

Practise 8.1

- Using Wireshark to capture and analysis the TCP stream.
- Invoke a HTTP request to get http://gaia.cs.umass.edu/wiresharklabs/alice.txt
- Analysis the TCP stream
  - Any duplicate ack, what's the possible reason?
  - Any TCP segment with sack permit option and sack option
    - select a tcp package which contained a sack option, find the segment ranges which is acked in this sack option
  - Any TCP retransmission? Is it retransmission or fast retransmission？
  - Any window size 0 segment, what does it mean? If the window size is 0, what would happened next on this tcp connection?
  - Any TCP window full segment, what does it mean?
- Tips: you can use some tools to cause your network congestion, such as clumsy-0.2-win64, which has been uploaded in Sakai site.

Tips on wireshark(1)

- How to get all the TCP segments of the TCP stream related to a http session in Wireshark:
  - 1st step: Find a HTTP packet in the HTTP session
  - 2nd step: right click the packet which will invoke a shortcut menu, then choose “follow ->TCP” in the shortcut menu
- How to find a special segment in a TCP stream
  - Using view filter in Wireshark
  - 1st step: if a TCP stream is filted, then a description such as “TCP.stream eq xx” (xx here is the id of this TCP stream) could be found in the view filter
  - 2nd step: in the same view filter, make a new filter rule description along with the original one,then press “Enter” key in your keyboard to make the new filter run
    - such as :

    TCP.stream eq 1 is the original one, to find a TCP zero window in this TCP stream,
    TCP.stream eq 1 && TCP.window_size_value==0 is the new view filter description

Tips on wireshark(2)

- Analysis -> expert info
- Statistic->TCP stream graphs

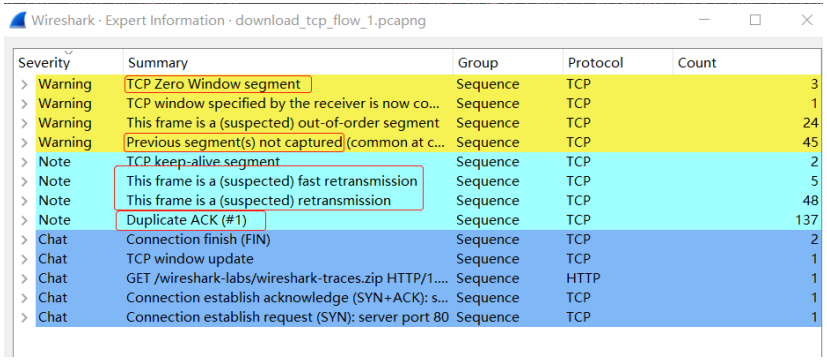

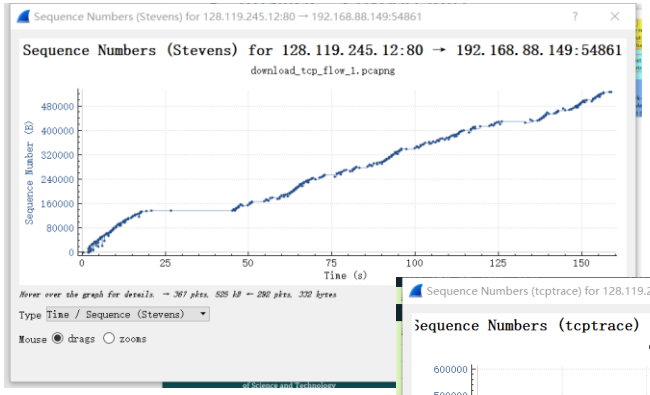

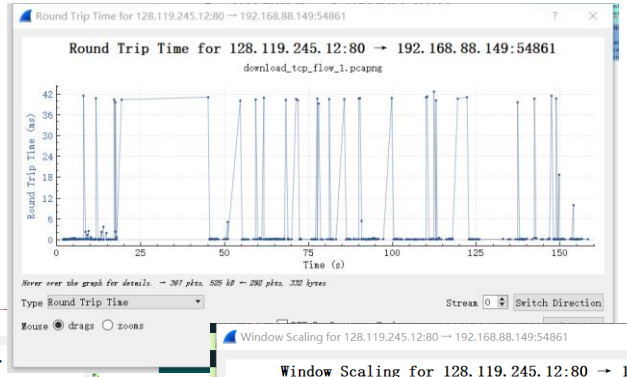

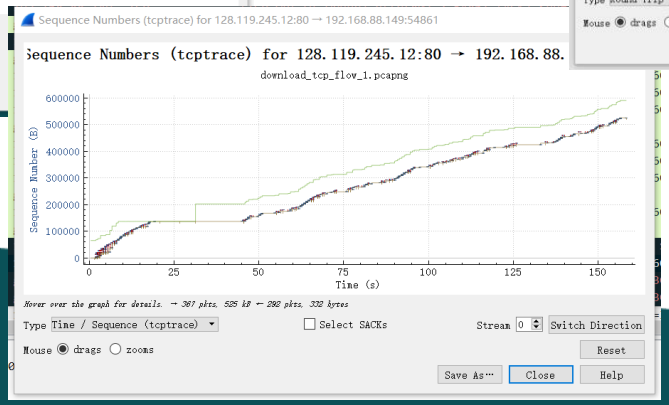

SUSTech
Southern University of Science and Technology