

# Lecture 3

# Processes

Prof. Yinqian Zhang

Summer 2022

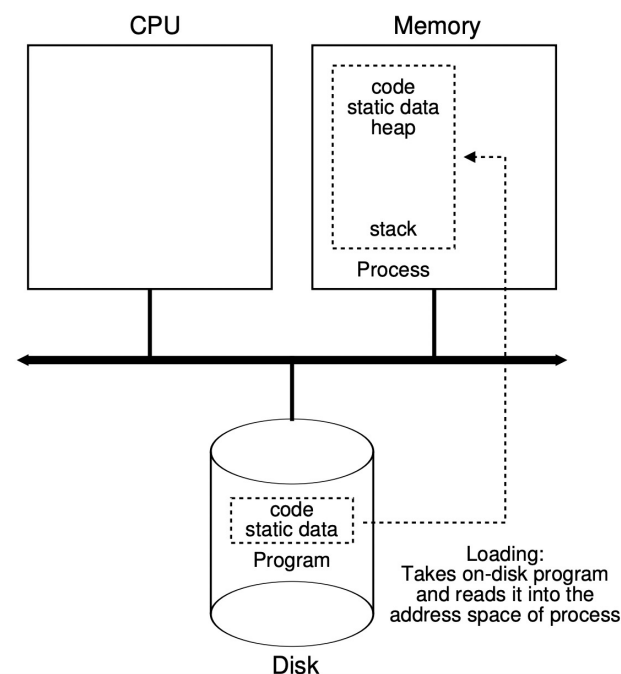
# Outline

- Process and system calls
- Process creation
- Kernel view of processes
- Kernel view of `fork()`, `exec()`, and `wait()`
- More about processes
- Threads

# Process and System Calls


# What Is a Process

- Process is a program in execution
- A program is a file on the disk
  - Code and static data
- A process is loaded by the OS
  - Code and static data are loaded from the program
  - Heap and stack are created by the OS



# What Is a Process (Cont'd)

- A process is an abstraction of machine states
  - Memory: address space
    - Register:
      - Program Counter (PC) or Instruction Pointer
      - Stack pointer
      - frame pointer
  - I/O: all files opened by the process



# Process Identification

- How can we distinguish processes from one to another?
  - Each process is given a unique ID number, and is called the process ID, or the PID.
  - The system call, `getpid()`, prints the PID of the calling process.

```
// compile to getpid
#include <stdio.h> // printf()
#include <unistd.h> // getpid()


int main(void) {
    printf("My PID is %d\n", getpid());
}
```

```
$ ./getpid
My PID is 1234
$ ./getpid
My PID is 1235
$ ./getpid
My PID is 1237
```

# Process Life Cycle

```
int main(void) {  
    int x = 1;  
    getchar();  
    return x;  
}
```

Process State

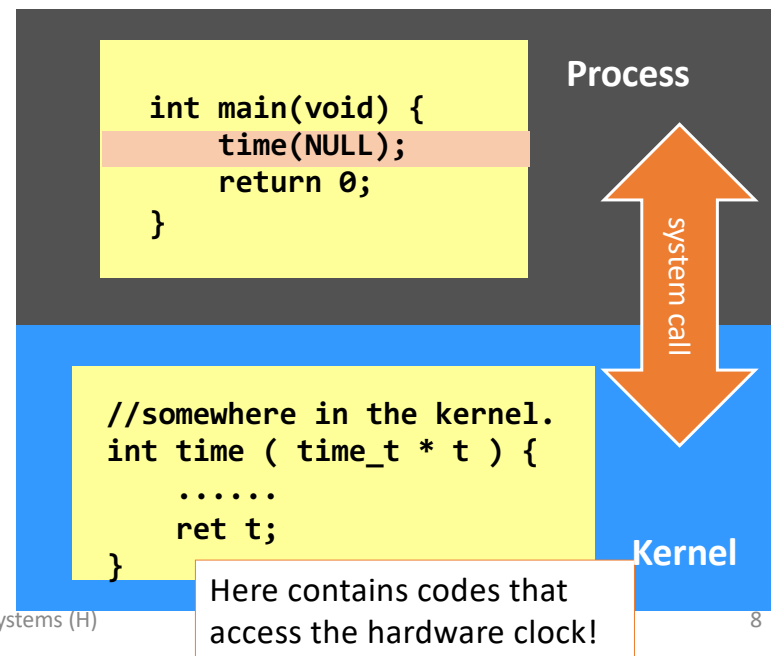


# System Call: Process-Kernel Interaction

- System call is a function call.
  - exposed by the **kernel**.
  - abstraction of kernel operations.

```
int add_function(int a, int b) {  
    return (a + b);  
}  
  
int main(void) {  
    int result;  
    result = add_function(a,b);  
    return 0;  
}  
  
// this is a dummy example...
```

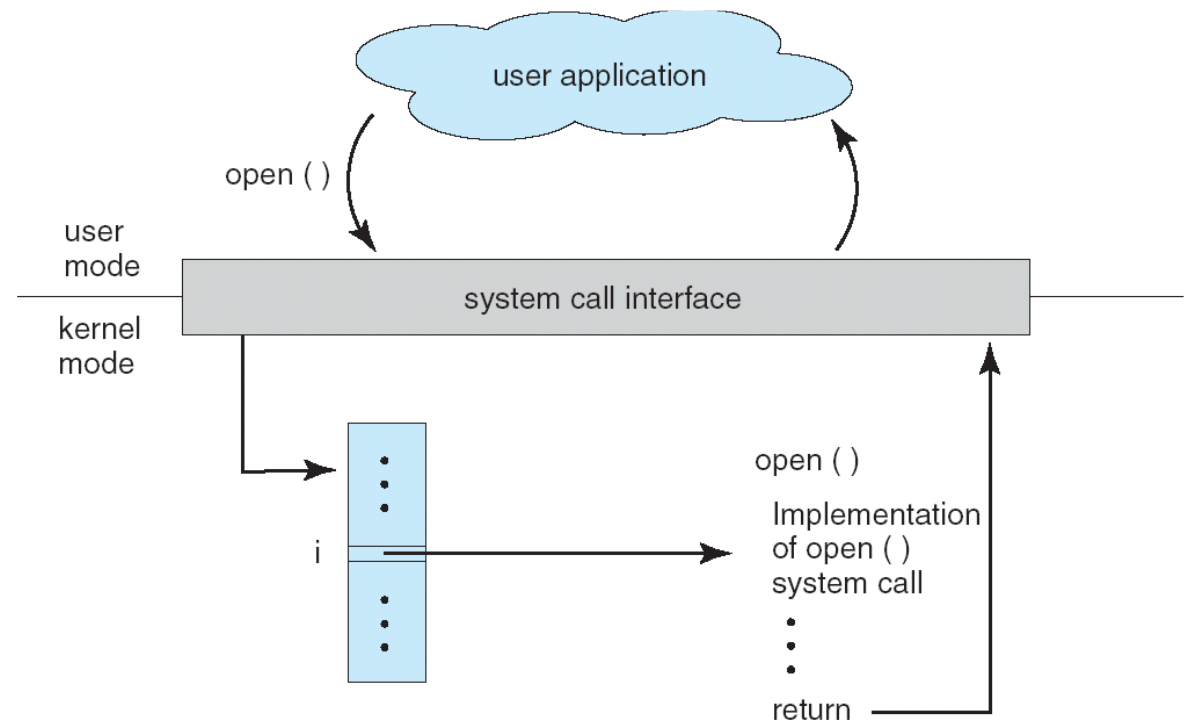
This is a  
function call.





# System Call: Call by Number

- System call is different from function call
- System call is a call by number



# System Call: Call by Number

- User-mode code from xv6-riscv

```
int main(void) {  
    .....  
    int fd = open("copyin1", O_CREATE|O_WRONLY);  
    .....  
    return 0;  
}
```

```
/* kernel/syscall.h */  
  
#define SYS_open 15
```

```
/* user/usys.S */  
.global open  
open:  
    li a7, SYS_open  
    ecall  
    ret
```

# System Call: Call by Number

- Kernel code from xv6-riscv

```
/* kernel/syscall.h */
```

```
#define SYS_open 15
```

```
/* kernel/file.c */
```

```
uint64 sys_open(void) {  
    .....  
    return fd;  
}
```

```
/* kernel/syscall.c */
```

```
static uint64 (*syscalls[])(void) = {  
    .....  
    [SYS_open] sys_open,  
    .....  
}
```

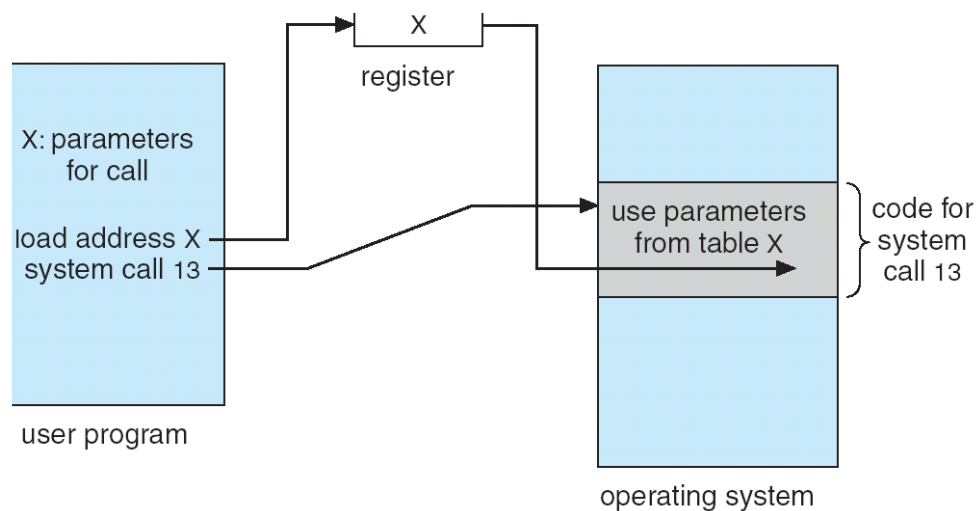
```
void syscall(void) (  
    struct proc *p = myproc();  
    num = p->trapframe->a7;  
    p->trapframe->a0 = syscalls[num]();  
}
```

# System Call: Parameter Passing

- Often, more information is required than the index of desired system call
  - Exact type and amount of information vary according to OS and call
- Three general methods used to pass parameters to the OS
  - **Registers:** pass the parameters in registers
    - In some cases, may be more parameters than registers
    - x86 and risc-v take this approach
  - **Blocks:** Parameters stored in a memory block and address of the block passed as a parameter in a register
  - **Stack:** Parameters placed, or pushed, onto the stack by the program and popped off the stack by the operating system
    - Block and stack methods do not limit the number or length of parameters being passed

# System Call: Parameter Passing


- Example: parameter passing via blocks



# System Call v.s. Library API Call

- Most operating systems provide standard C library to provide library API calls
  - A layer of indirection for system calls

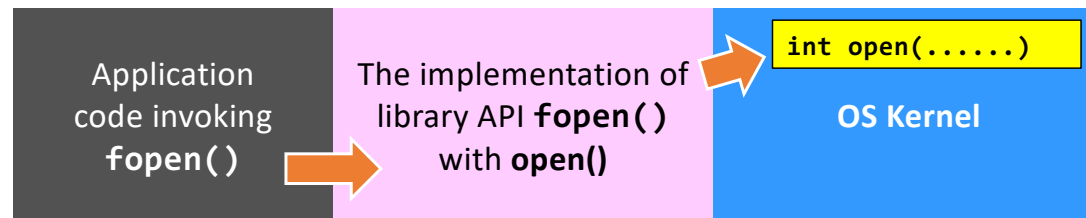
Name	System call?
printf() & scanf()	No
malloc() & free()	No
fopen() & fclose()	No
mkdir() & rmdir()	Yes
chown() & chmod()	Yes



# System Call v.s. Library API Call

- Take `fopen()` as an example.
  - `fopen()` invokes the system call `open()`.
  - `open()` is too primitive and is not programmer-friendly!

Library call	<code>fopen("hello.txt", "w");</code>
System call	<code>open("hello.txt", O_WRONLY   O_CREAT   O_TRUNC, 0666);</code>



# Process Creation



# Process Creation


- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Generally, process identified and managed via a process identifier (pid)
- Resource sharing
  - Parent and children share all resources
  - Children share subset of parent's resources
  - Parent and child share no resources
- Execution
  - Parent and children execute concurrently
  - Parent waits until children terminate

# Process Creation (Cont'd)

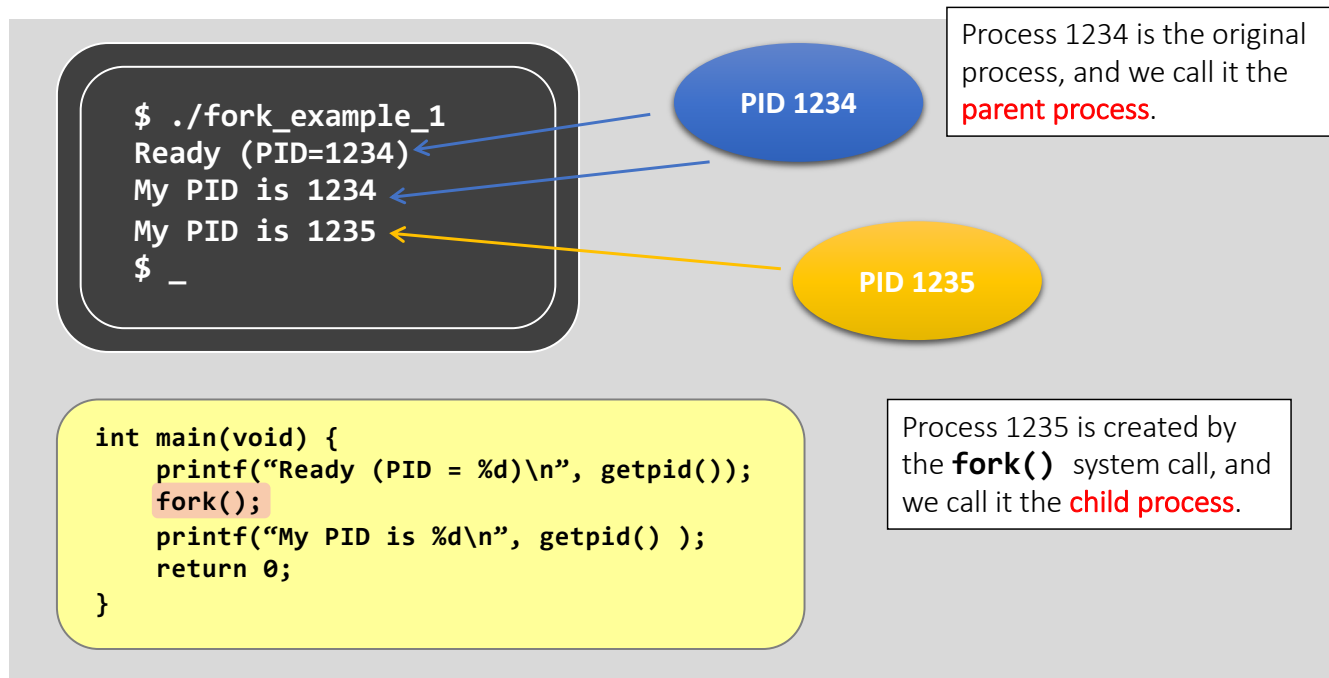
- Address space
  - Child duplicate of parent
  - Child has a program loaded into it
- UNIX examples
  - fork system call creates new process
  - exec system call used after a fork to replace the process' memory space with a new program

# Process Creation (Cont'd)

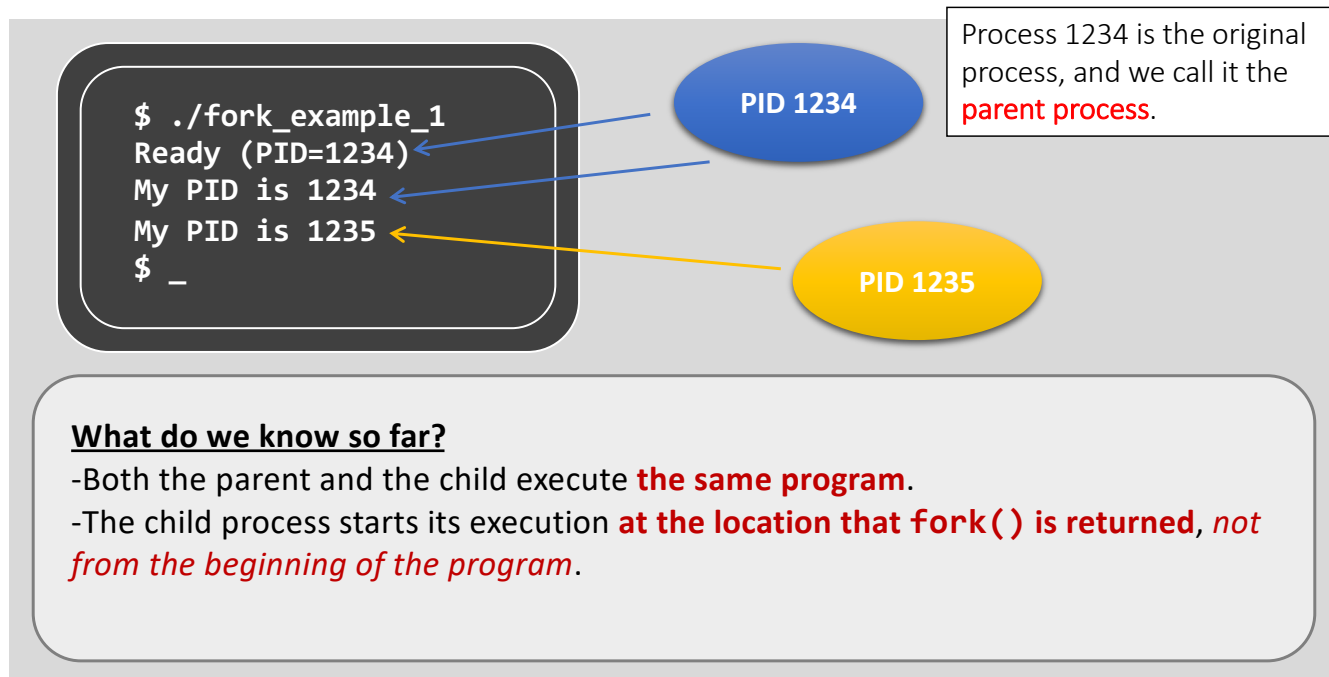
- A tree of processes in Linux




# Creating Processes with fork() System Call



# Creating Processes with fork() System Call



# Creating Processes with fork() System Call




```
1 int main(void) {
2     int result;
3     printf("before fork ...\n");
4     result = fork();
5     printf("result = %d.\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\n");
9         printf("My PID is %d\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\n");
13        printf("My PID is %d\n", getpid());
14    }
15
16    printf("program terminated.\n");
17 }
```

```
$ ./fork_example_2
before fork ...
```

PID 1234

# Creating Processes with fork() System Call



```
1 int main(void) {
2     int result;
3     printf("before fork ...\n");
4     result = fork();
5     printf("result = %d.\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\n");
9         printf("My PID is %d\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\n");
13        printf("My PID is %d\n", getpid());
14    }
15
16    printf("program terminated.\n");
17 }
```

```
$ ./fork_example_2
before fork ...
```

## Important

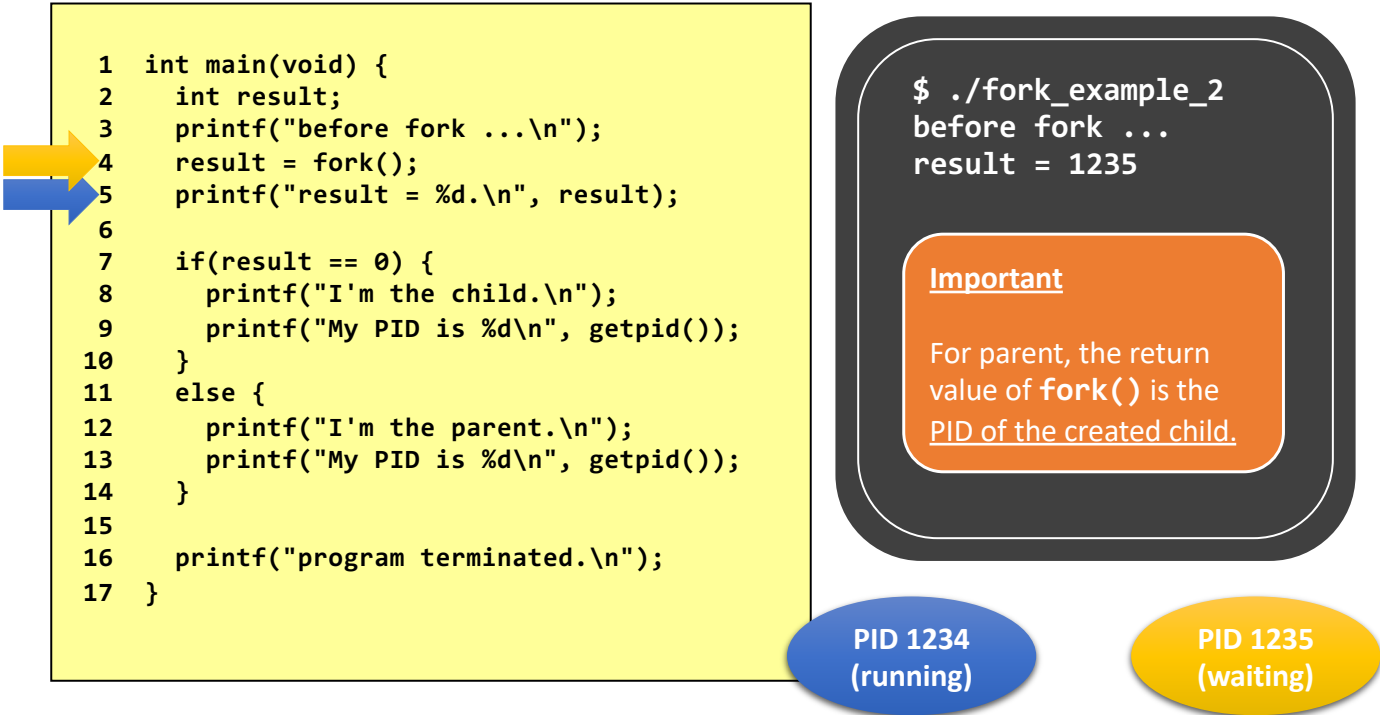
- Both parent and child need to return from fork().
- CPU scheduler decides which to run first.

PID 1234

fork()

PID 1235

# Creating Processes with fork() System Call



```
1 int main(void) {
2     int result;
3     printf("before fork ...\n");
4     result = fork();
5     printf("result = %d.\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\n");
9         printf("My PID is %d\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\n");
13        printf("My PID is %d\n", getpid());
14    }
15
16    printf("program terminated.\n");
17 }
```

```
$ ./fork_example_2
before fork ...
result = 1235
```

## Important

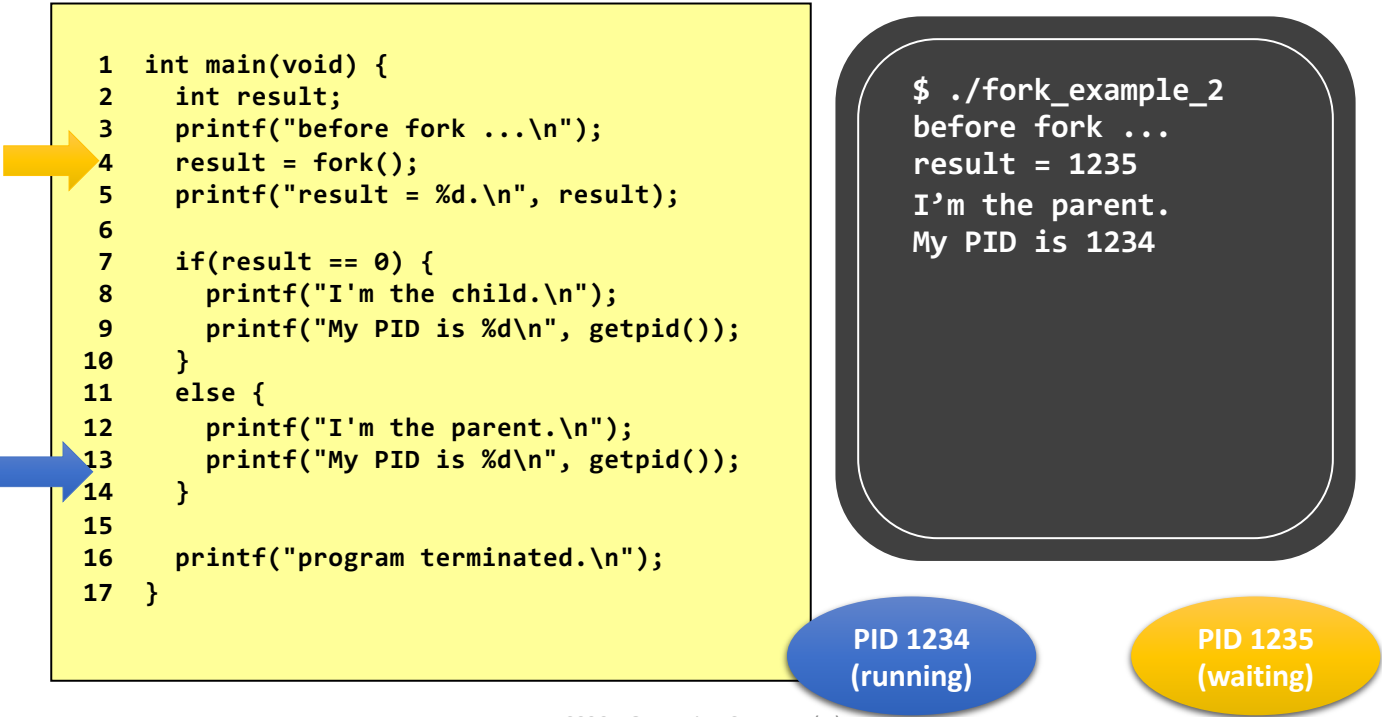
For parent, the return value of `fork()` is the PID of the created child.

PID 1234  
(running)

PID 1235  
(waiting)



# Creating Processes with fork() System Call



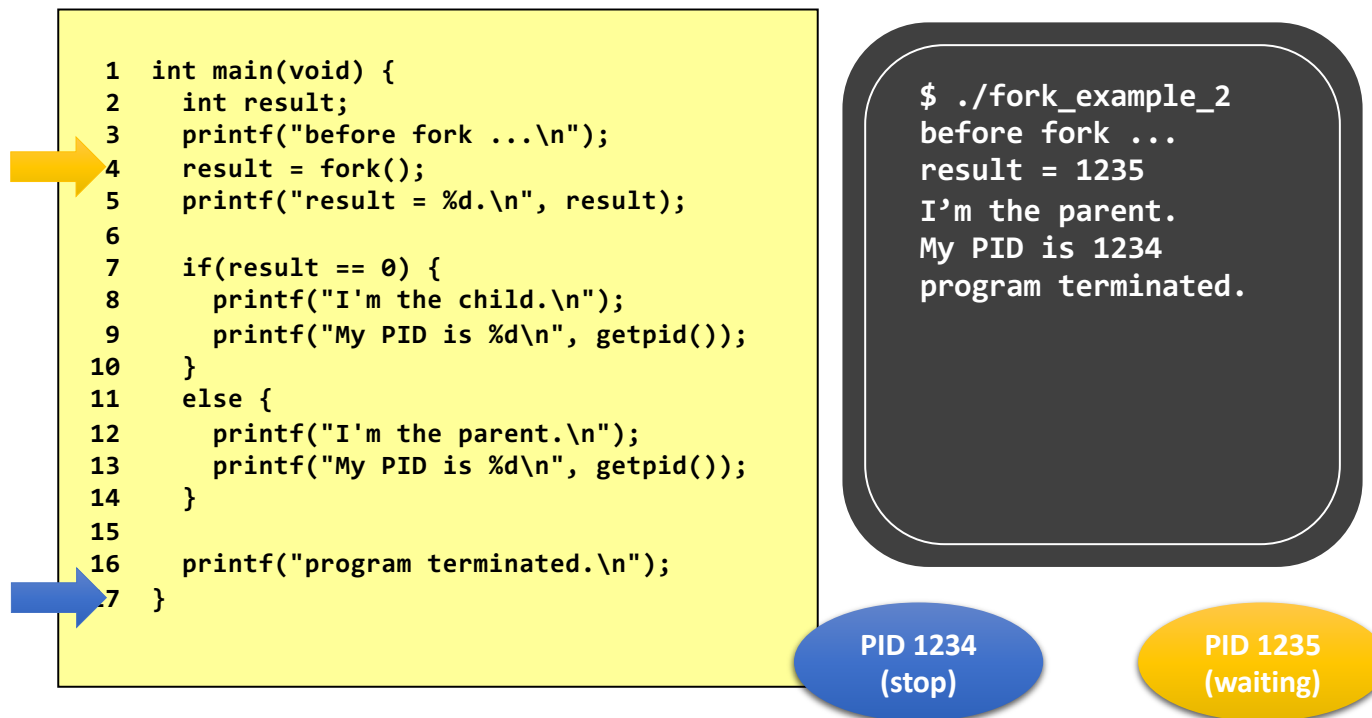
```
1 int main(void) {  
2     int result;  
3     printf("before fork ...\n");  
4     result = fork();  
5     printf("result = %d.\n", result);  
6  
7     if(result == 0) {  
8         printf("I'm the child.\n");  
9         printf("My PID is %d\n", getpid());  
10    }  
11    else {  
12        printf("I'm the parent.\n");  
13        printf("My PID is %d\n", getpid());  
14    }  
15  
16    printf("program terminated.\n");  
17 }
```

```
$ ./fork_example_2  
before fork ...  
result = 1235  
I'm the parent.  
My PID is 1234
```


PID 1234  
(running)

PID 1235  
(waiting)


# Creating Processes with fork() System Call



# Creating Processes with fork() System Call



```
1 int main(void) {
2     int result;
3     printf("before fork ...\n");
4     result = fork();
5     printf("result = %d.\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\n");
9         printf("My PID is %d\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\n");
13        printf("My PID is %d\n", getpid());
14    }
15
16    printf("program terminated.\n");
17 }
```



```
$ ./fork_example_2
before fork ...
result = 1235
I'm the parent.
My PID is 1234
program terminated.
result = 0
```

## Important

For child, the return value of `fork()` is 0.

PID 1234  
(stop)

PID 1235  
(running)

# Creating Processes with fork() System Call

```
1 int main(void) {  
2     int result;  
3     printf("before fork ...\n");  
4     result = fork();  
5     printf("result = %d.\n", result);  
6  
7     if(result == 0) {  
8         printf("I'm the child.\n");  
9         printf("My PID is %d\n", getpid());  
10    }  
11    else {  
12        printf("I'm the parent.\n");  
13        printf("My PID is %d\n", getpid());  
14    }  
15  
16    printf("program terminated.\n");  
17 }
```

```
$ ./fork_example_2  
before fork ...  
result = 1235  
I'm the parent.  
My PID is 1234  
program terminated.  
result = 0  
I'm the child.  
My PID is 1235
```

PID 1234  
(stop)

PID 1235  
(running)

# Creating Processes with fork() System Call

```
1 int main(void) {  
2     int result;  
3     printf("before fork ...\n");  
4     result = fork();  
5     printf("result = %d.\n", result);  
6  
7     if(result == 0) {  
8         printf("I'm the child.\n");  
9         printf("My PID is %d\n", getpid());  
10    }  
11    else {  
12        printf("I'm the parent.\n");  
13        printf("My PID is %d\n", getpid());  
14    }  
15  
16    printf("program terminated.\n");  
17 }
```

```
$ ./fork_example_2  
before fork ...  
result = 1235  
I'm the parent.  
My PID is 1234  
program terminated.  
result = 0  
I'm the child.  
My PID is 1235  
program terminated.  
$ _
```

PID 1234  
(stop)

PID 1235  
(stop)

# fork() System Call

- `fork()` behaves like “cell division”.
  - It creates the child process by **cloning** from the parent process, including all user-space data, e.g.,

Cloned items	Descriptions
Program counter [CPU register]	That’s why they both execute from the same line of code after <code>fork()</code> returns.
Program code [File & Memory]	They are sharing the same piece of code.
Memory	Including local variables, global variables, and dynamically allocated memory.
Opened files [Kernel’s internal]	If the parent has opened a file “fd”, then the child will also have file “fd” opened automatically.

# fork() System Call

- fork() does not clone the following...

Distinct items	Parent	Child
Return value of fork()	PID of the child process.	0
PID	Unchanged.	Different, not necessarily be "Parent PID + 1"
Parent process	Unchanged.	Parent.
Running time	Cumulated.	Just created, so should be 0.
<b>[Advanced]</b> File locks	Unchanged.	None.

# fork() System Call

- If a process can only duplicate itself and always runs the same program, it's not quite meaningful
  - how can we execute other programs?
- **exec()**
  - The **exec\*()** system call family.



# exec()

- `execl()` – a member of the `exec` system call family (`execl`, `execle`, `execlp`, `execv`, `execve`, `execvp`).

```
int main(void) {  
    printf("before execl ...\n");  
    execl("/bin/ls", "/bin/ls", NULL);  
    printf("after execl ...\n");  
    return 0;  
}
```

```
$ ./exec_example  
before execl ...
```

## Arguments of the `execl()` call

1<sup>st</sup> argument: the program name, `"/bin/ls"` in the example.

2<sup>nd</sup> argument: `argument[0]` to the program.

3<sup>rd</sup> argument: `argument[1]` to the program.

# exec()

- `execl()` – a member of the `exec` system call family (and the family has 6 members).

```
int main(void) {  
    printf("before execl ...\n");  
    → execl("/bin/ls", "/bin/ls", NULL);  
    printf("after execl ...\n");  
    return 0;  
}
```

```
$/exec_example  
before execl ...  
exec_example  
exec_example.c
```

What is the output?

The same as the output of running “ls” in the shell.

# exec()

- Example #1: run the command **"/bin/ls"**

```
execl("/bin/ls", "/bin/ls", NULL);
```

Argument Order	Value in above example	Description
1	<b>"/bin/ls"</b>	The file that the programmer wants to execute.
2	<b>"/bin/ls"</b>	When the process switches to <b>"/bin/ls"</b> , this string is the <b>program argument[0]</b> .
3	<b>NULL</b>	This states the end of the program argument list.

# exec()

- Example #2: run the command **"/bin/ls -l"**

```
execl("/bin/ls", "/bin/ls", "-l", NULL);
```

Argument Order	Value in above example	Description
1	<b>"/bin/ls"</b>	The file that the programmer wants to execute.
2	<b>"/bin/ls"</b>	When the process switches to <b>"/bin/ls"</b> , this string is the <b>program argument[0]</b> .
3	<b>"-l"</b>	When the process switches to <b>"/bin/ls"</b> , this string is the <b>program argument[1]</b> .
4	<b>NULL</b>	This states the end of the program argument list.

# exec()

- The `exec` system call family is not simply a function that “invokes” a command.

```
int main(void) {  
    printf("before execl ...\n");  
    execl("/bin/ls", "/bin/ls", NULL);  
    printf("after execl ...\n");  
    return 0;  
}
```

WHAT?!  
The shell prompt appears!

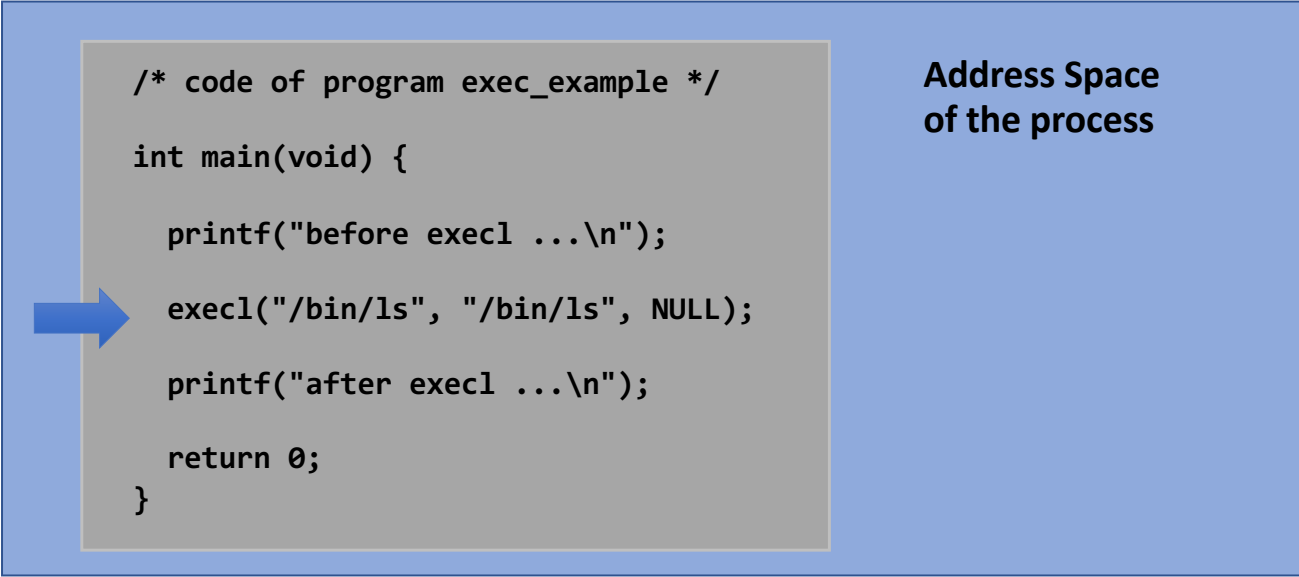
```
./exec_example  
before execl ...  
exec_example  
exec_example.c  
$ _
```

The output says:  
(1) The gray code block **is not reached!**  
(2) The process is **terminated!**

WHY IS THAT?!

# exec()

- The `exec` system call family is not simply a function that “invokes” a command.



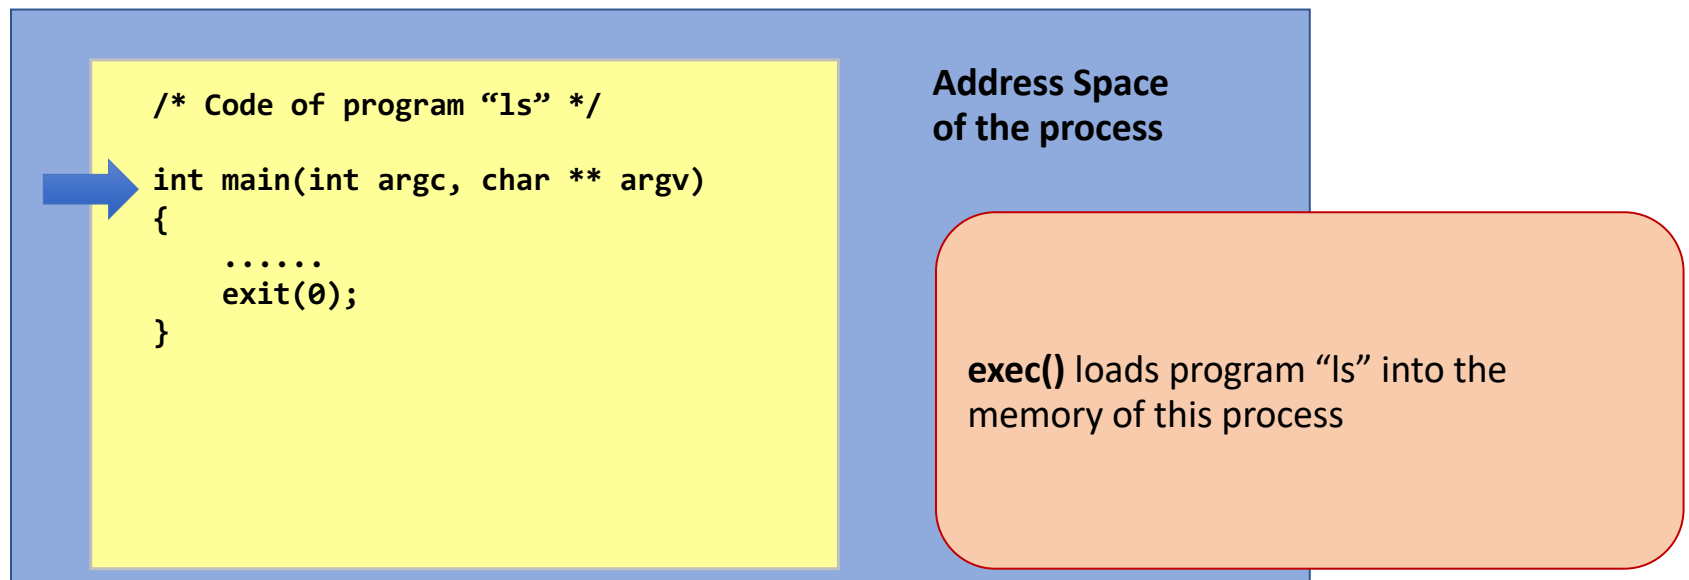
The diagram illustrates the address space of a process. It is represented by a large light blue rectangle. Inside this rectangle, on the left, is a smaller gray rectangle containing C code. A blue arrow points from the left edge of the gray rectangle to the `execl` function call. To the right of the gray rectangle, within the light blue area, is the text "Address Space of the process".

```
/* code of program exec_example */
int main(void) {
    printf("before execl ...\n");
    execl("/bin/ls", "/bin/ls", NULL);
    printf("after execl ...\n");
    return 0;
}
```

Address Space  
of the process

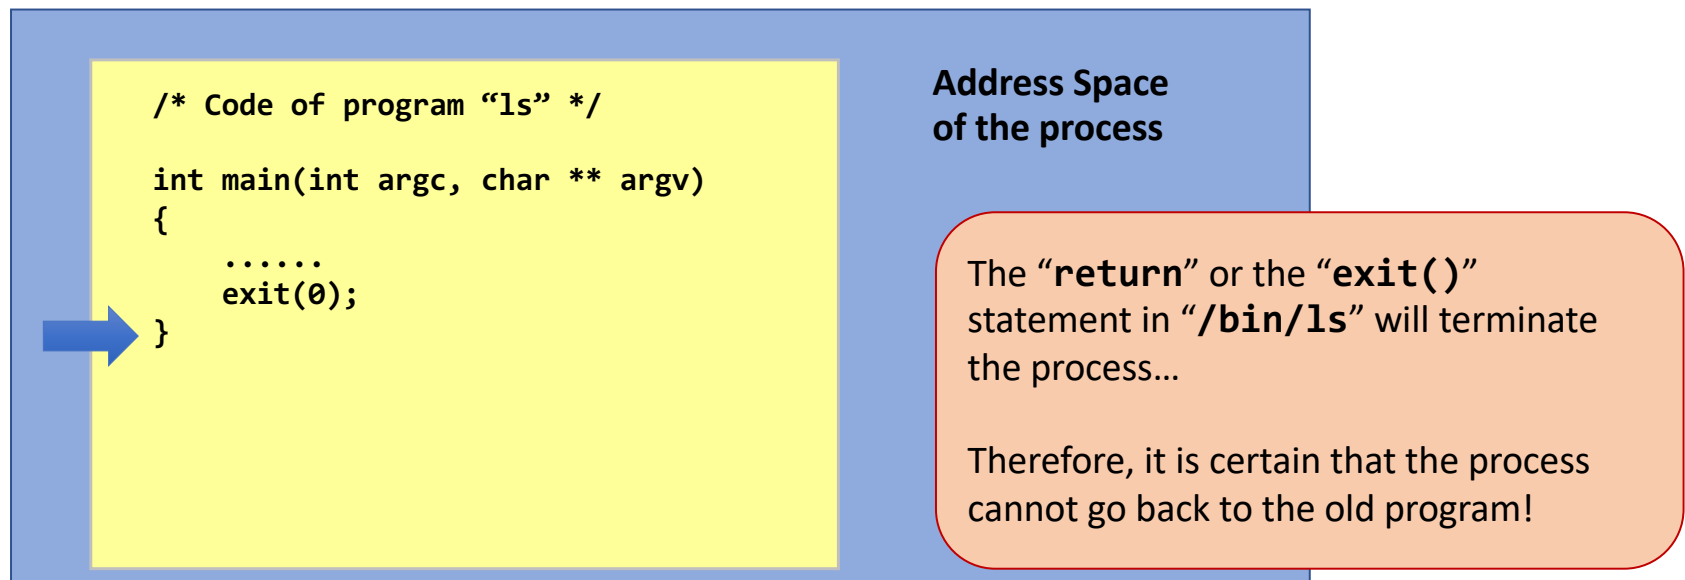
# exec()

- The `exec` system call family is not simply a function that “invokes” a command.



# exec()

- The `exec` system call family is not simply a function that “invokes” a command.





# exec() Summary

- The process is changing the code that is executing and never returns to the original code.
  - The last two lines of codes are therefore not executed.
- The process that calls an `exec*` system call will replace user-space info, e.g.,
  - Program Code
  - Memory: local variables, global variables, and dynamically allocated memory;
  - Register value: e.g., the program counter;
- But, the kernel-space info of that process is preserved, including:
  - PID;
  - Process relationship;
  - etc.

# CPU Scheduler and fork()

```
1 int main(void) {
2     int result;
3     printf("before fork ...\n");
4     result = fork();
5     printf("result = %d.\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\n");
9         printf("My PID is %d\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\n");
13        printf("My PID is %d\n", getpid());
14    }
15
16    printf("program terminated.\n");
17 }
```

Parent return  
from fork() first

```
$ ./fork_example_2
before fork ...
result = 1235
I'm the parent.
My PID is 1234
program terminated.
result = 0
I'm the child.
My PID is 1235
program terminated.
$ _
```

Child return  
from fork() first

```
$ ./fork_example_2
before fork ...
result = 0
I'm the child.
My PID is 1235
result = 1235
program terminated.
I'm the parent.
My PID is 1234
program terminated.
$ _
```

# wait(): Sync Parent with Child

```
1 int main(void) {
2     int result;
3     printf("before fork ...\n");
4     result = fork();
5     printf("result = %d.\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\n");
9         printf("My PID is %d\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\n");
13        wait(NULL);
14        printf("My PID is %d\n", getpid());
15    }
16
17    printf("program terminated.\n");
18 }
```

Parent return  
from fork() first

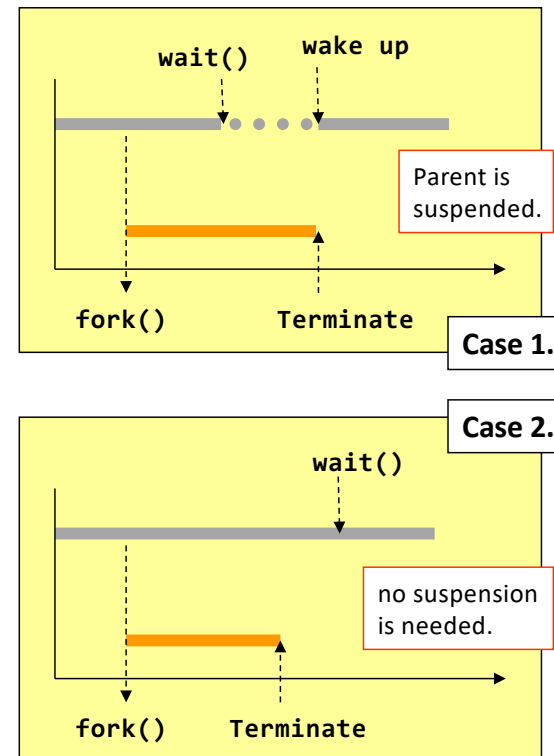
```
$ ./fork_example_2
before fork ...
result = 1235
I'm the parent.
result = 0
I'm the child.
My PID is 1235
program terminated.
My PID is 1234
program terminated.
$ _
```

Child return  
from fork() first

```
$ ./fork_example_2
before fork ...
result = 0
I'm the child.
My PID is 1235
result = 1235
program terminated.
I'm the parent.
My PID is 1234
program terminated.
$ _
```

# wait()

- wait() suspends the calling process to **waiting**
- wait() returns when
  - one of its child processes changes from running to terminated.
- Return immediately (i.e., does nothing) if
  - It has no children
  - Or a child terminates before the parent calls wait for

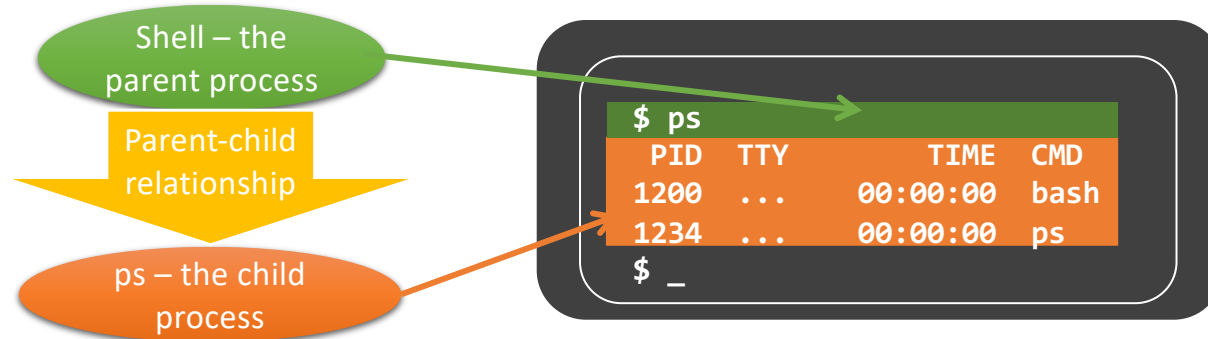


# **wait() v.s. waitpid()**

- **wait()**
  - Wait for any one of the child processes
  - Detect child termination only
- **waitpid()**
  - Depending on the parameters, waitpid() will wait for a particular child only
  - Depending on the parameters, waitpid() can detect different status changes of the child (resume/stop by a signal)

# Implement Shell with fork(), exec(), and wait()

- A shell is a CLI
  - Bash in linux
  - invokes a function fork() to create a new process
  - Ask the the child process to exec() the target program
  - Use wait() to wait until the child process terminates

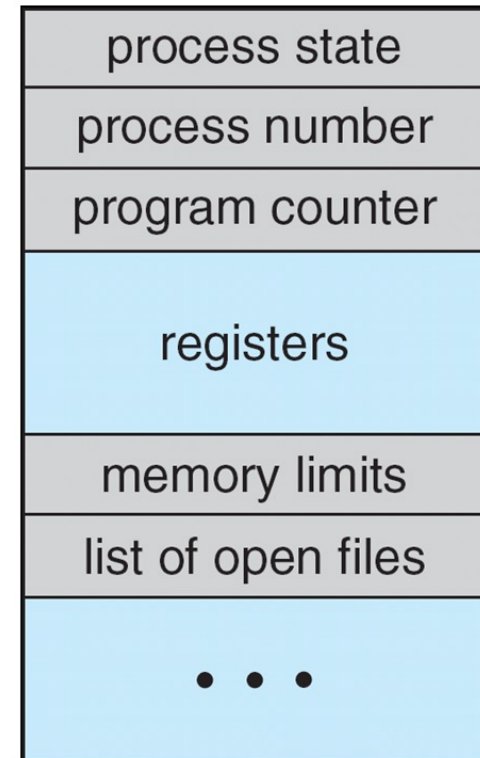


# Processes: Kernel View

# Process Control Block (PCB)

Information associated with each process

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information





# PCB Example: uCore

```
/* kern/process/proc.h in ucore */

struct proc_struct {
    enum proc_state state;           // Process state
    int pid;                         // Process ID
    int runs;                        // the running times of Process
    uintptr_t kstack;               // Process kernel stack
    volatile bool need_resched;      // bool value: need to be rescheduled to release CPU?
    struct proc_struct *parent;      // the parent process
    struct mm_struct *mm;           // Process's memory management field
    struct context context;         // Switch here to run process
    struct trapframe *tf;           // Trap frame for current interrupt
    uintptr_t cr3;                  // CR3 register: the base addr of Page Directroy Table(PDT)
    uint32_t flags;                 // Process flag
    char name[PROC_NAME_LEN + 1];   // Process name
    list_entry_t list_link;         // Process link list
}
```

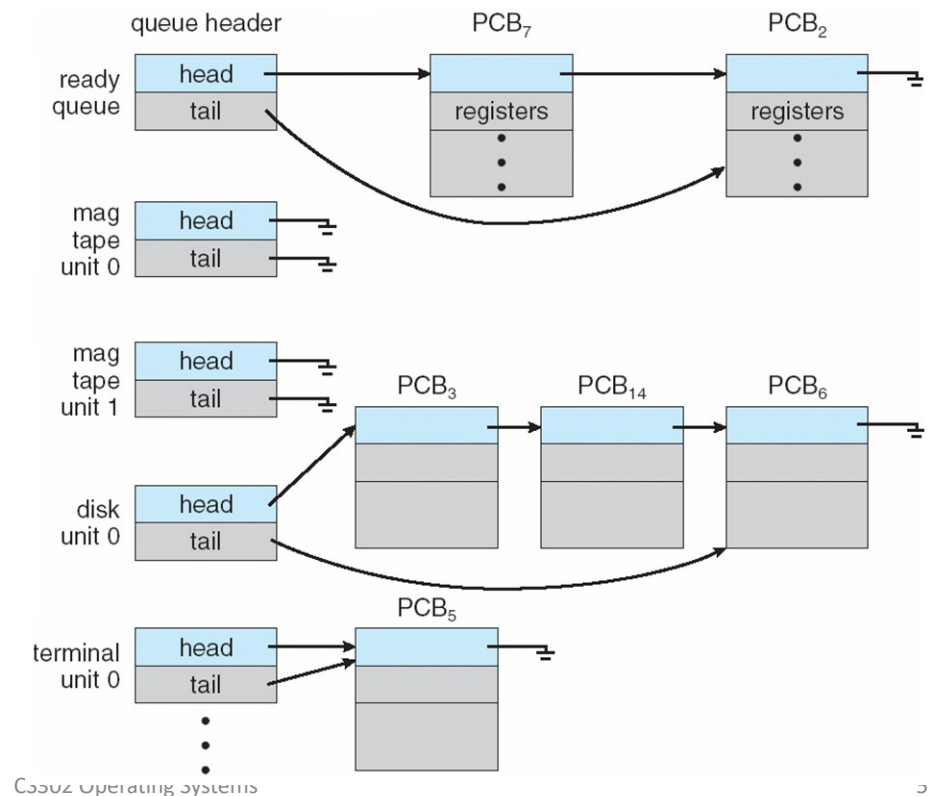
# PCB Example: uCore

```
/* kern/process/proc.h in ucore */

list_entry_t hash_link;           // Process hash list
int exit_code;                    // exit code (be sent to parent proc)
uint32_t wait_state;              // waiting state
struct proc_struct *cptr, *yptr, *optr; // relations between processes
struct run_queue *rq;             // running queue contains Process
list_entry_t run_link;            // the entry linked in run queue
int time_slice;                   // time slice for occupying the CPU
struct files_struct *filesp;      // the file related info of process
};
```

# Ready Queue And I/O Device Queues

- PCBs are linked in multiple queues
  - Ready queue contains all processes in the ready state (to run on this CPU)
  - Device queue contains processes waiting for I/O events from this device
  - Process may migrate among these queues

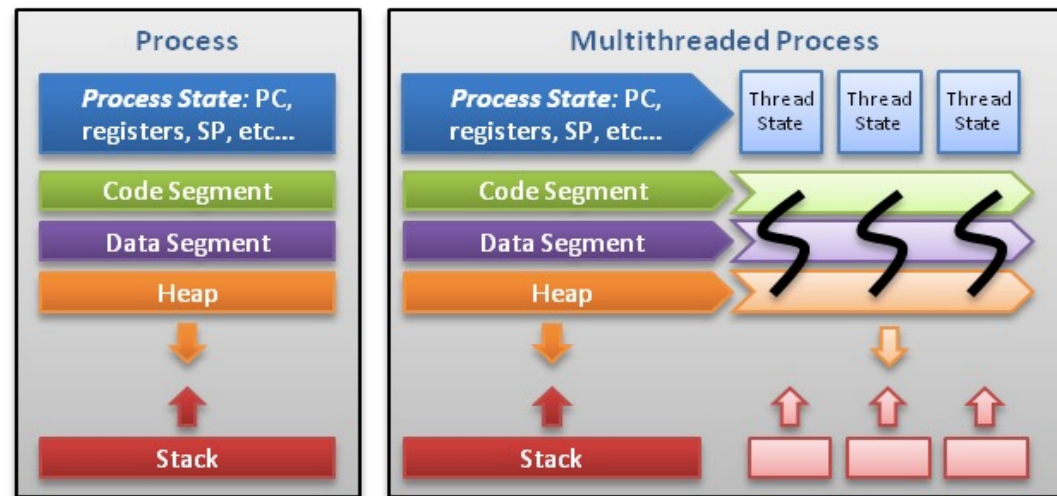


# Threads

- One process may have more than one threads
  - A single-threaded process performs a single thread of execution
  - A multi-threaded process performs multiple threads of execution “concurrently”, thus allowing short response time to user’s input even when the main thread is busy
- PCB is extended to include information about each thread

# Process and Thread

- Single threaded process and multi-threaded process



Threads contain only necessary information, such as a stack (for local variables, function arguments, return values), a copy of the registers, program counter and any thread-specific data to allow them to be scheduled individually. Other data is shared within the process between all threads.

# Switching Between Processes

- Once a process runs on a CPU, it only gives back the control of a CPU
  - when it makes a system call
  - when it raises an exception
  - when an interrupt occurs
- What if none of these would happen for a long time?
  - Cooperative scheduling: OS will have to wait
    - Early Macintosh OS, old Alto system
  - Non-cooperative scheduling: timer interrupts
    - Modern operating systems

# Switching Between Processes (Cont'd)

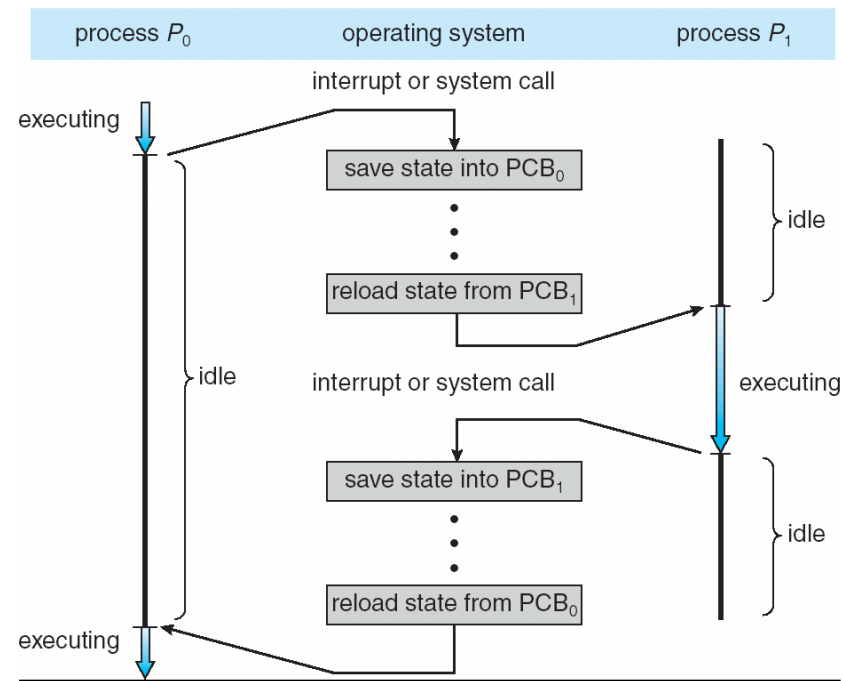
- When OS kernel regains the control of CPU
  - It first completes the task
    - Serve system call, or
    - Handle interrupt/exception
  - It then decides which process to run next
    - by asking its **CPU scheduler**
    - How does it make decisions?
    - More about CPU scheduler later
  - It performs a **context switch** if the soon-to-be-executing process is different from the previous one

# Context Switch

- During context switch, the system must save the state of the old process and load the saved state for the new process
- Context of a process is represented in the PCB
- The time used to do context switch is an overhead of the system; the system does no useful work while switching
  - Time of context switch depends on hardware support
  - Context switch cannot be too frequent



# Context Switch (Cont'd)

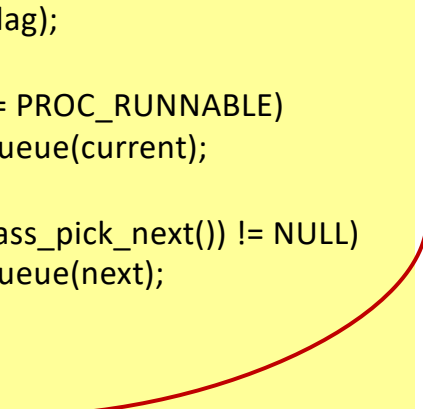


# Context Switch: uCore

```
/* kern/schedule/sched.c */
void schedule(void) {
    bool intr_flag;
    struct proc_struct *next;
    local_intr_save(intr_flag);
    {
        if (current->state == PROC_RUNNABLE)
            sched_class_enqueue(current);

        if ((next = sched_class_pick_next()) != NULL)
            sched_class_dequeue(next);

        if (next != current)
            proc_run(next);
    }
    local_intr_restore(intr_flag);
}
```



```
/* kern/process/proc.c */

void proc_run(struct proc_struct *proc) {
    if (proc != current) {
        bool intr_flag;
        struct proc_struct *prev = current, *next = proc;
        local_intr_save(intr_flag);
        {
            current = proc;
            lcr3(next->cr3);
            switch_to(&(prev->context), &(next->context));
        }
        local_intr_restore(intr_flag);
    }
}
```

# Context Switch: uCore (Cont'd)

```
/* kern/process/switch.S */
```

```
.globl switch_to
```

```
switch_to:
```

```
    # save from's registers
```

```
    STORE ra, 0*REGBYTES(a0)
```

```
    STORE sp, 1*REGBYTES(a0)
```

```
    STORE s0, 2*REGBYTES(a0)
```

```
    STORE s1, 3*REGBYTES(a0)
```

```
    STORE s2, 4*REGBYTES(a0)
```

```
    STORE s3, 5*REGBYTES(a0)
```

```
    STORE s4, 6*REGBYTES(a0)
```

```
    STORE s5, 7*REGBYTES(a0)
```

```
    STORE s6, 8*REGBYTES(a0)
```

```
    STORE s7, 9*REGBYTES(a0)
```

```
    STORE s8, 10*REGBYTES(a0)
```

```
    STORE s9, 11*REGBYTES(a0)
```

```
    STORE s10, 12*REGBYTES(a0)
```

```
    STORE s11, 13*REGBYTES(a0)
```

```
    # restore to's registers
```

```
    LOAD ra, 0*REGBYTES(a1)
```

```
    LOAD sp, 1*REGBYTES(a1)
```

```
    LOAD s0, 2*REGBYTES(a1)
```

```
    LOAD s1, 3*REGBYTES(a1)
```

```
    LOAD s2, 4*REGBYTES(a1)
```

```
    LOAD s3, 5*REGBYTES(a1)
```

```
    LOAD s4, 6*REGBYTES(a1)
```

```
    LOAD s5, 7*REGBYTES(a1)
```

```
    LOAD s6, 8*REGBYTES(a1)
```

```
    LOAD s7, 9*REGBYTES(a1)
```

```
    LOAD s8, 10*REGBYTES(a1)
```

```
    LOAD s9, 11*REGBYTES(a1)
```

```
    LOAD s10, 12*REGBYTES(a1)
```

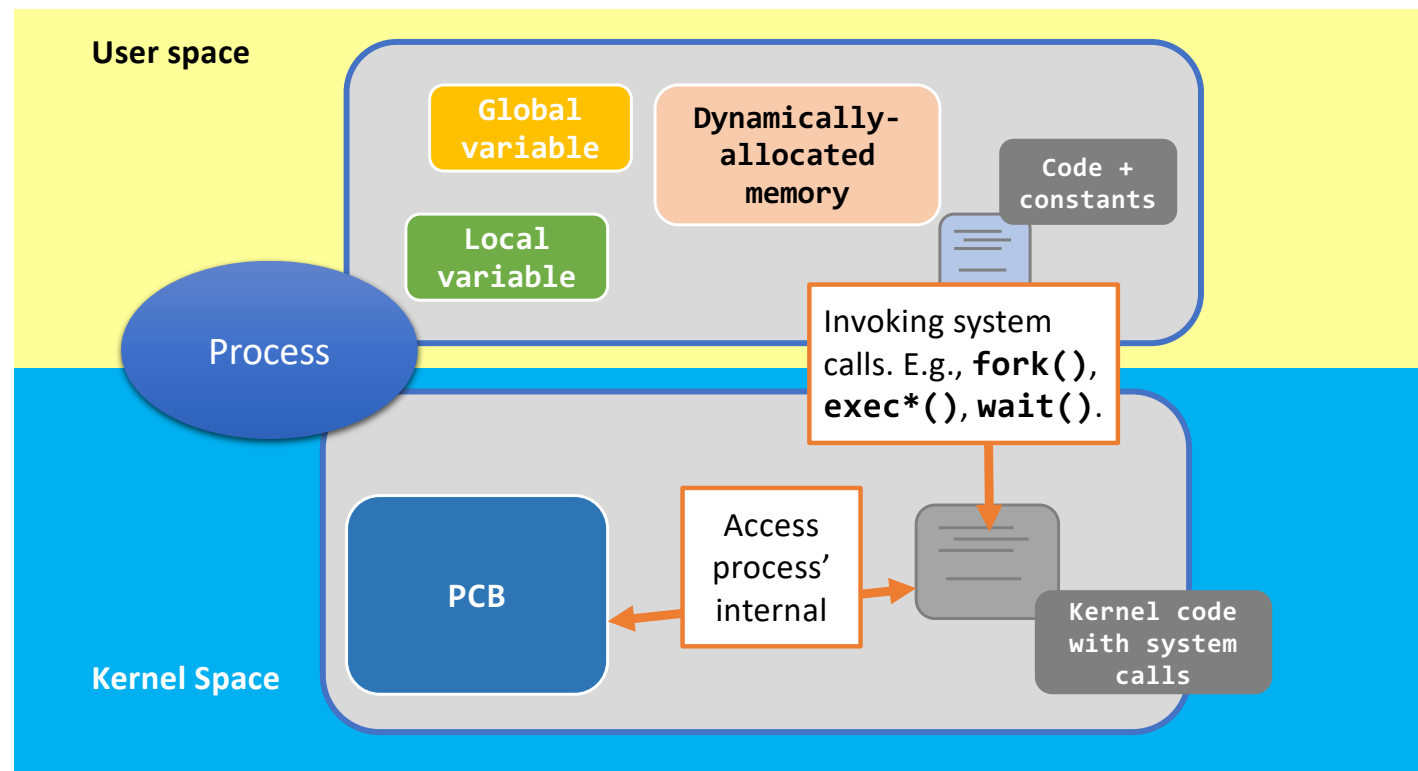
```
    LOAD s11, 13*REGBYTES(a1)
```

```
ret
```


# **fork(), exec(), wait()**

## **Kernel View**

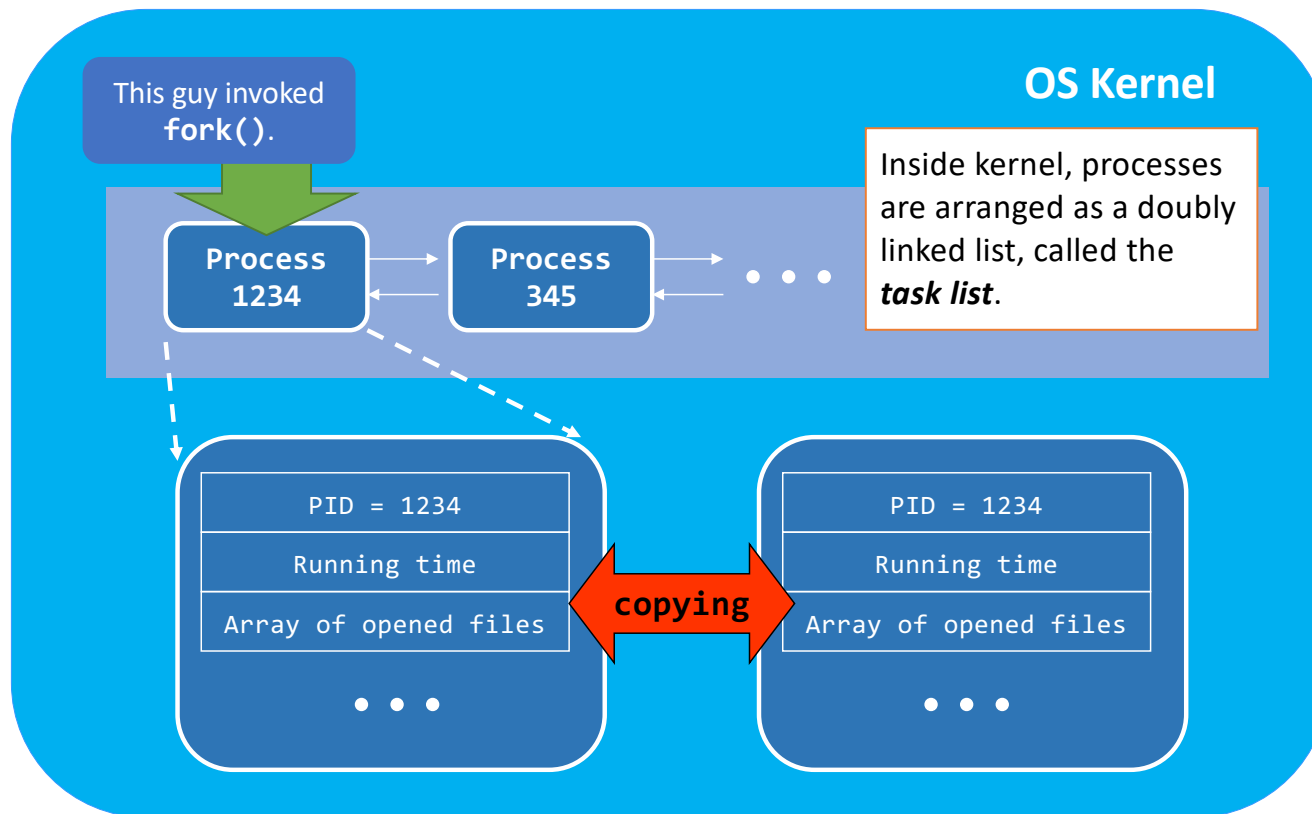
# Recall: `fork()`, `exec()`, and `wait()`



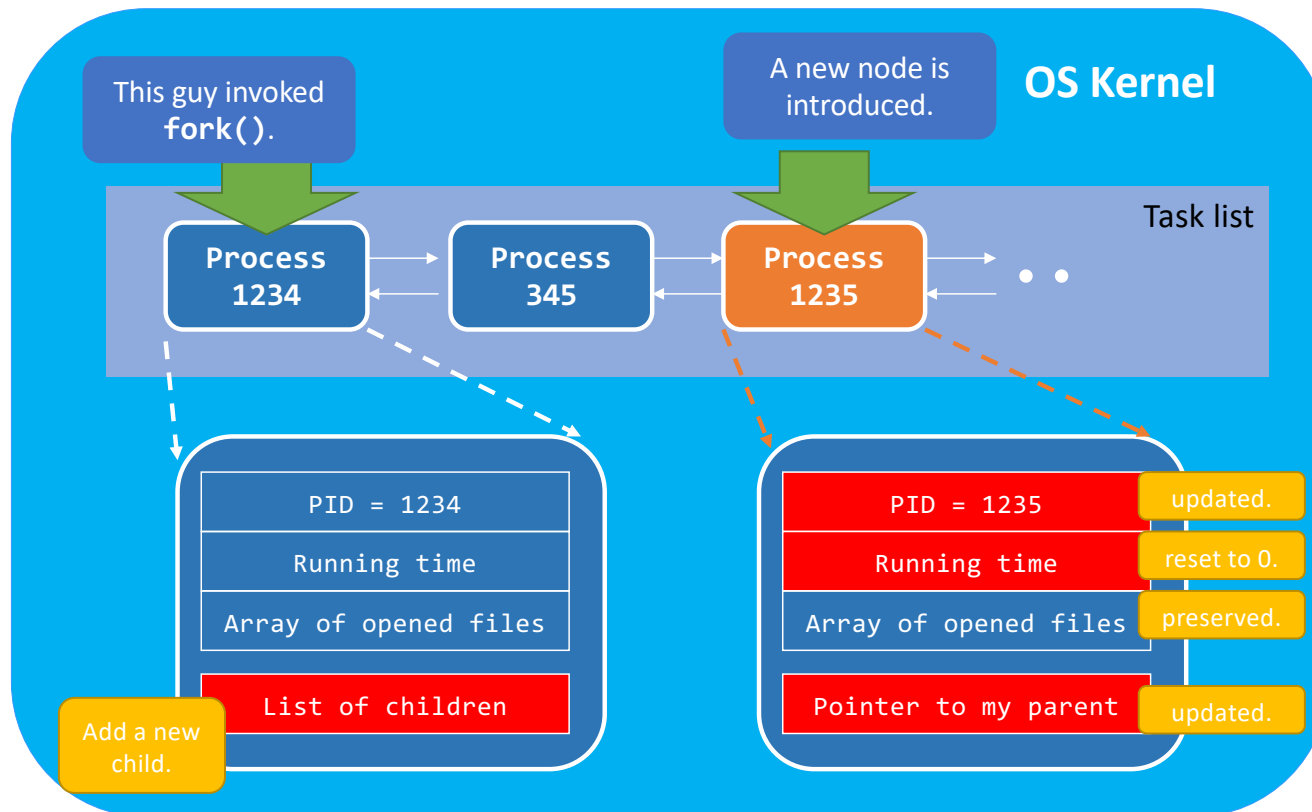
# Fork() in User Mode



# fork(): Kernel View

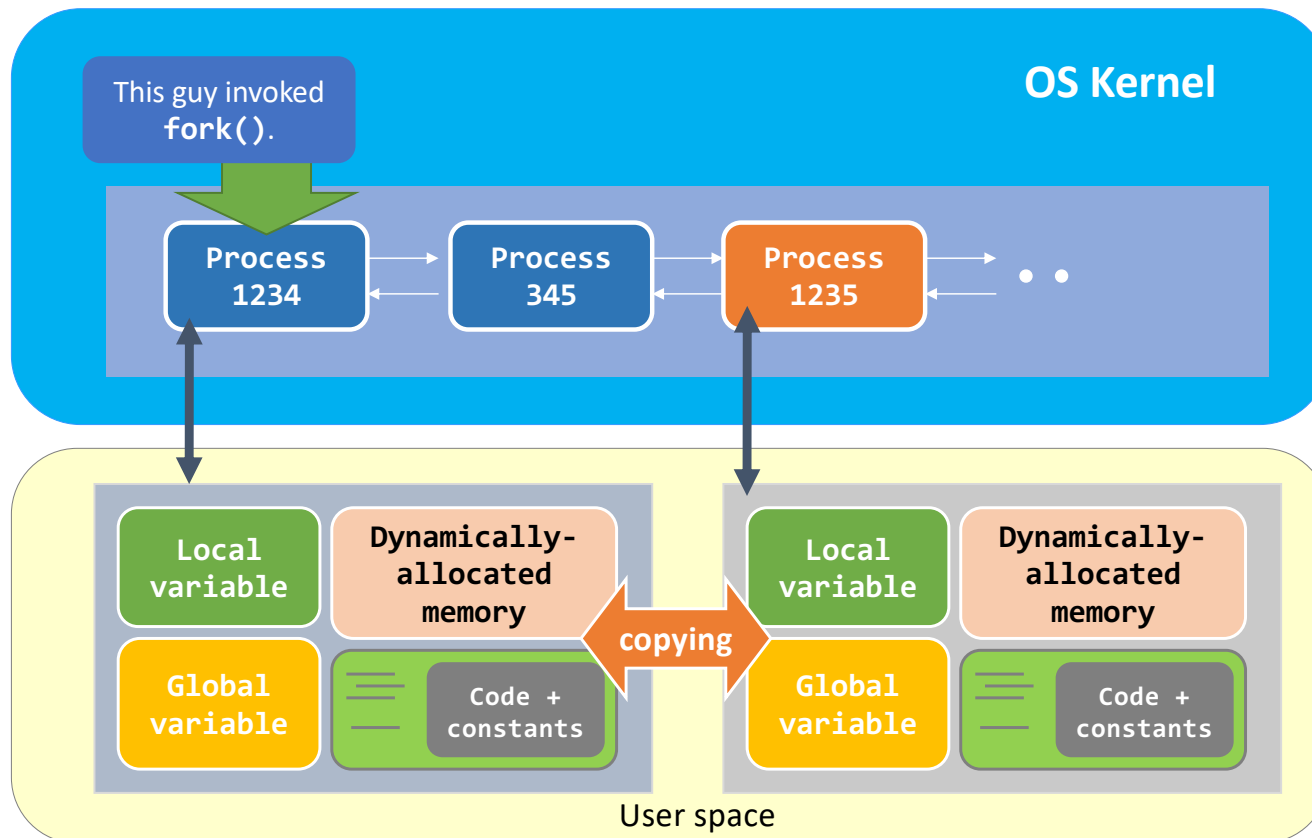


# fork(): Kernel View

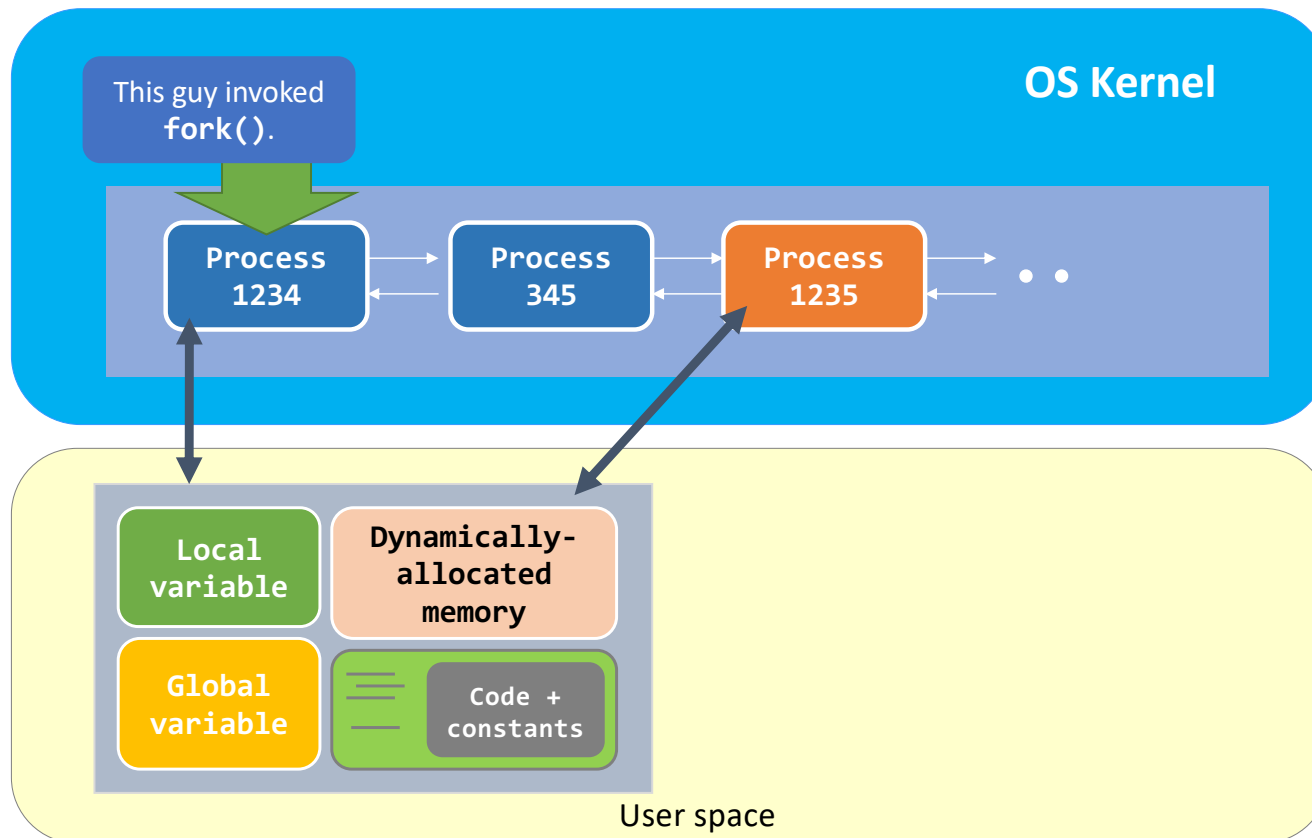




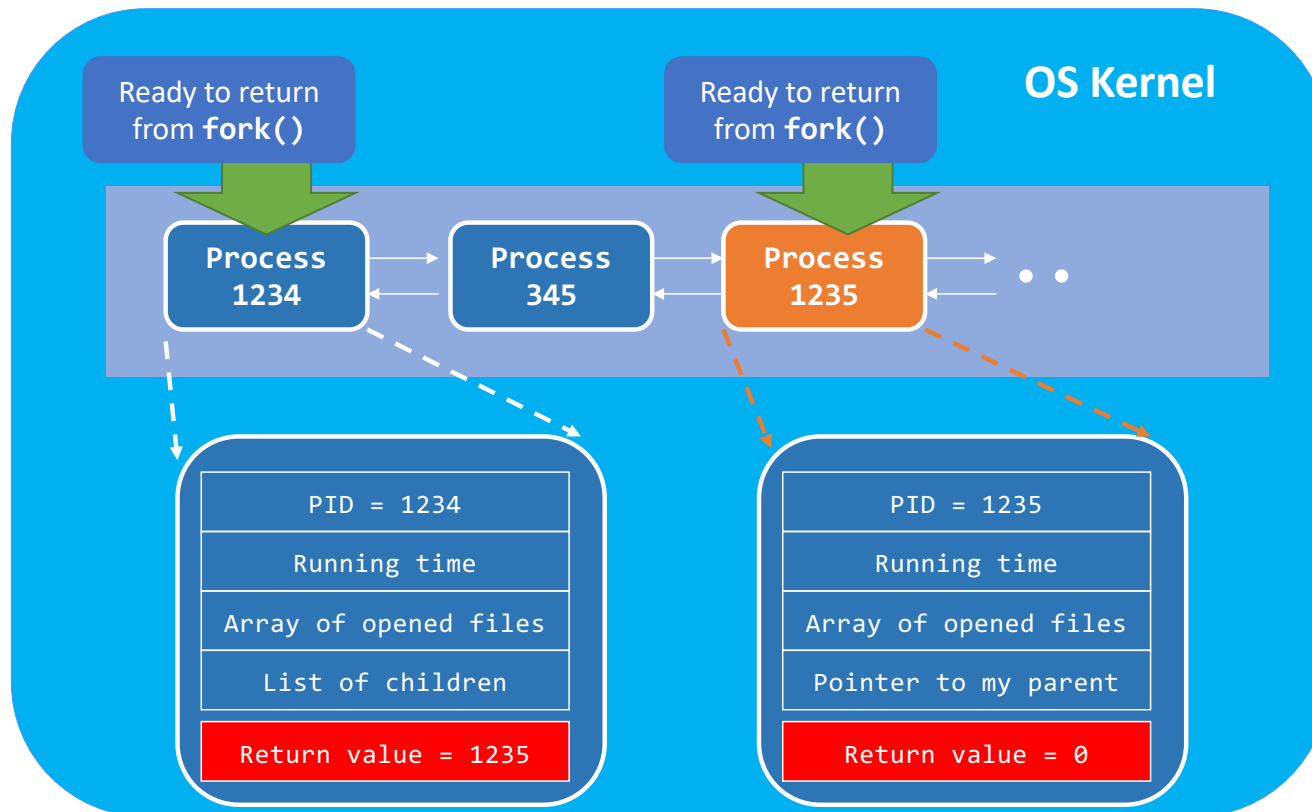
# Case 1: Duplicate Address Space



## Case 2: Copy on Write



# fork(): Kernel View



# fork(): Opened Files


- Array of opened files contains:

Array Index	Description
0	Standard Input Stream; <b>FILE *stdin;</b>
1	Standard Output Stream; <b>FILE *stdout;</b>
2	Standard Error Stream; <b>FILE *stderr;</b>
3 or beyond	Storing the files you opened, e.g., <b>fopen()</b> , <b>open()</b> , etc.

- That's why a parent process shares the same terminal output stream as the child process.

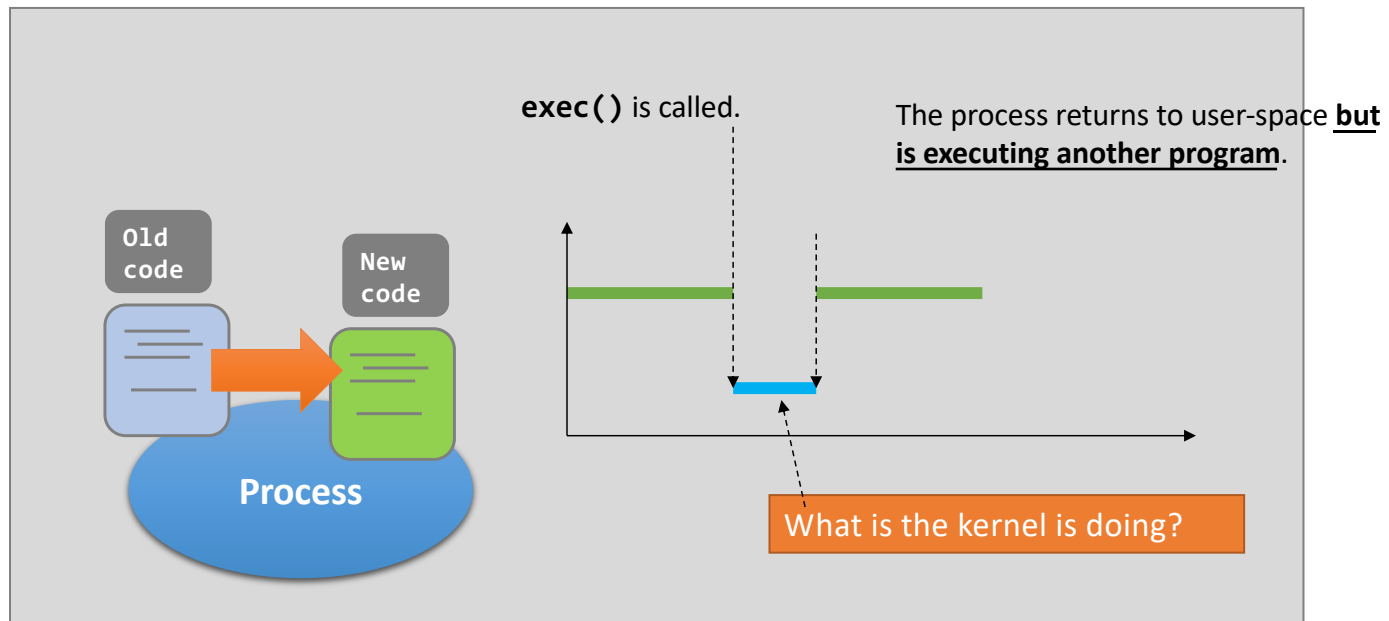
# fork(): Opened Files

- What if two processes, sharing the same opened file, write to that file together?

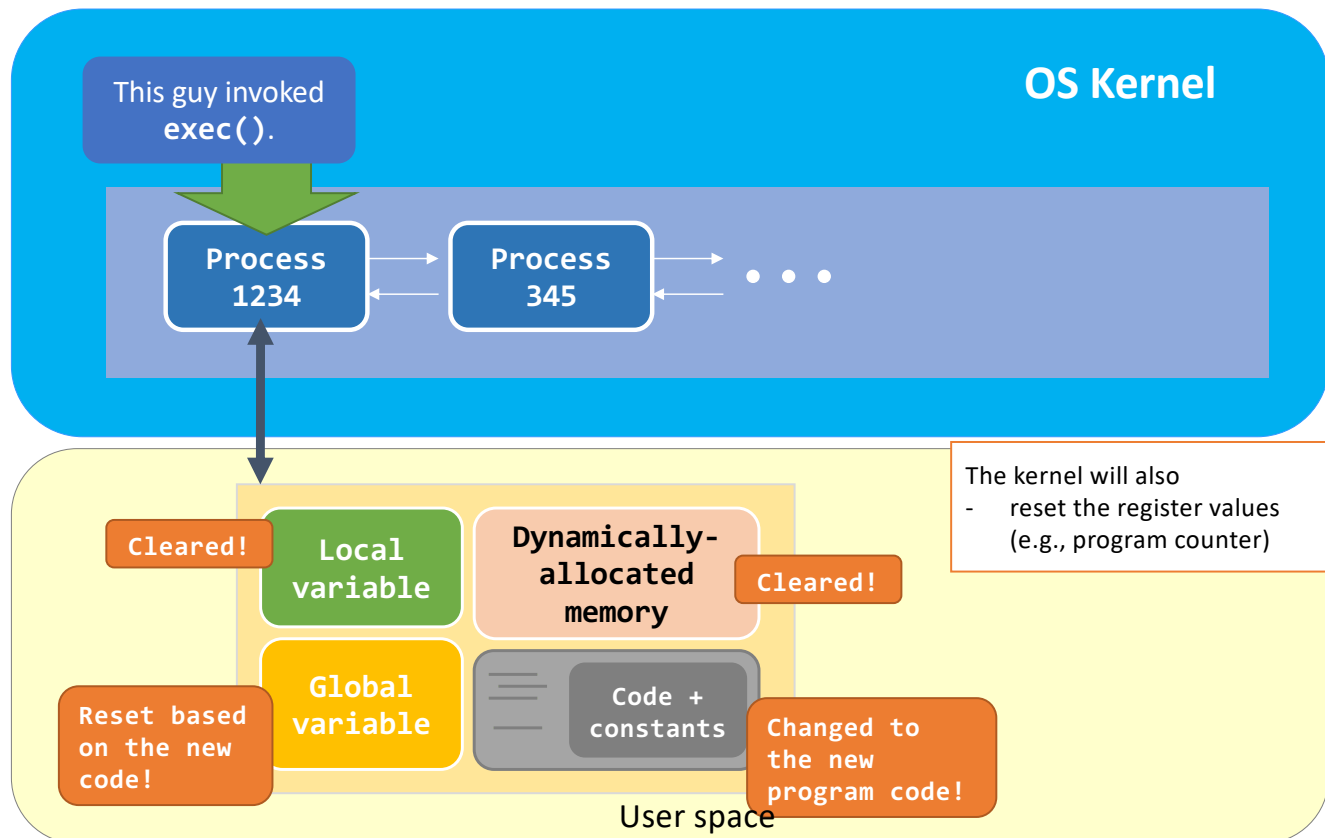


Let's see what will happen when the program finishes running!

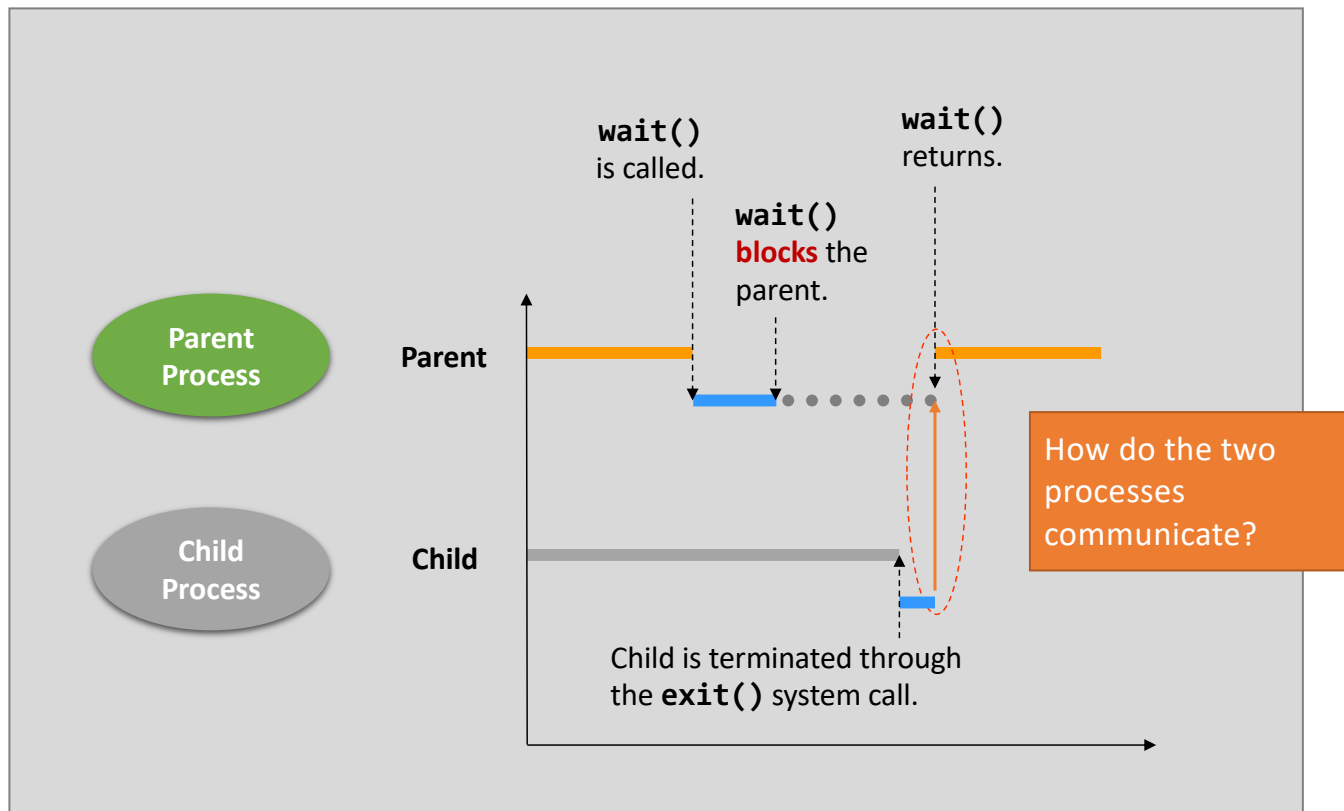
# exec() in the User Mode



# exec(): Kernel View

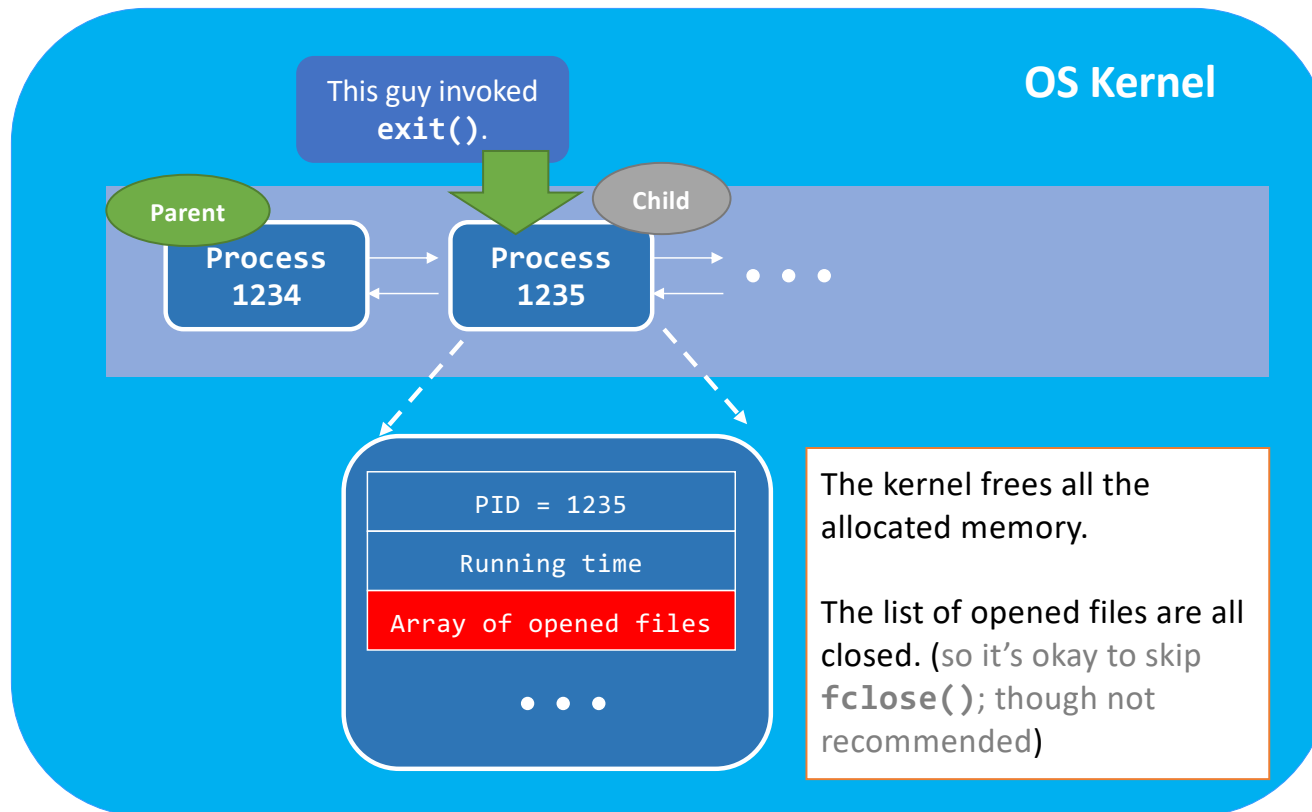


# wait() and exit()

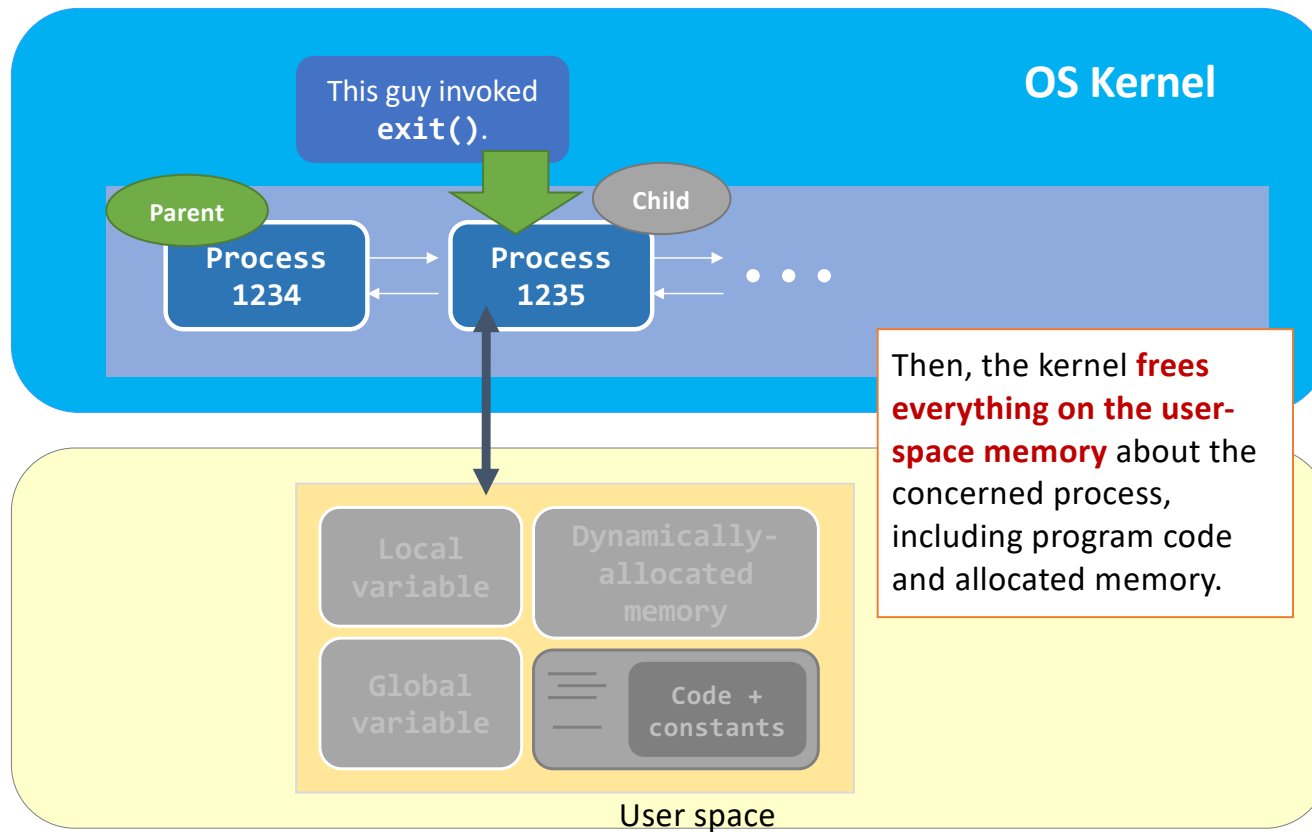




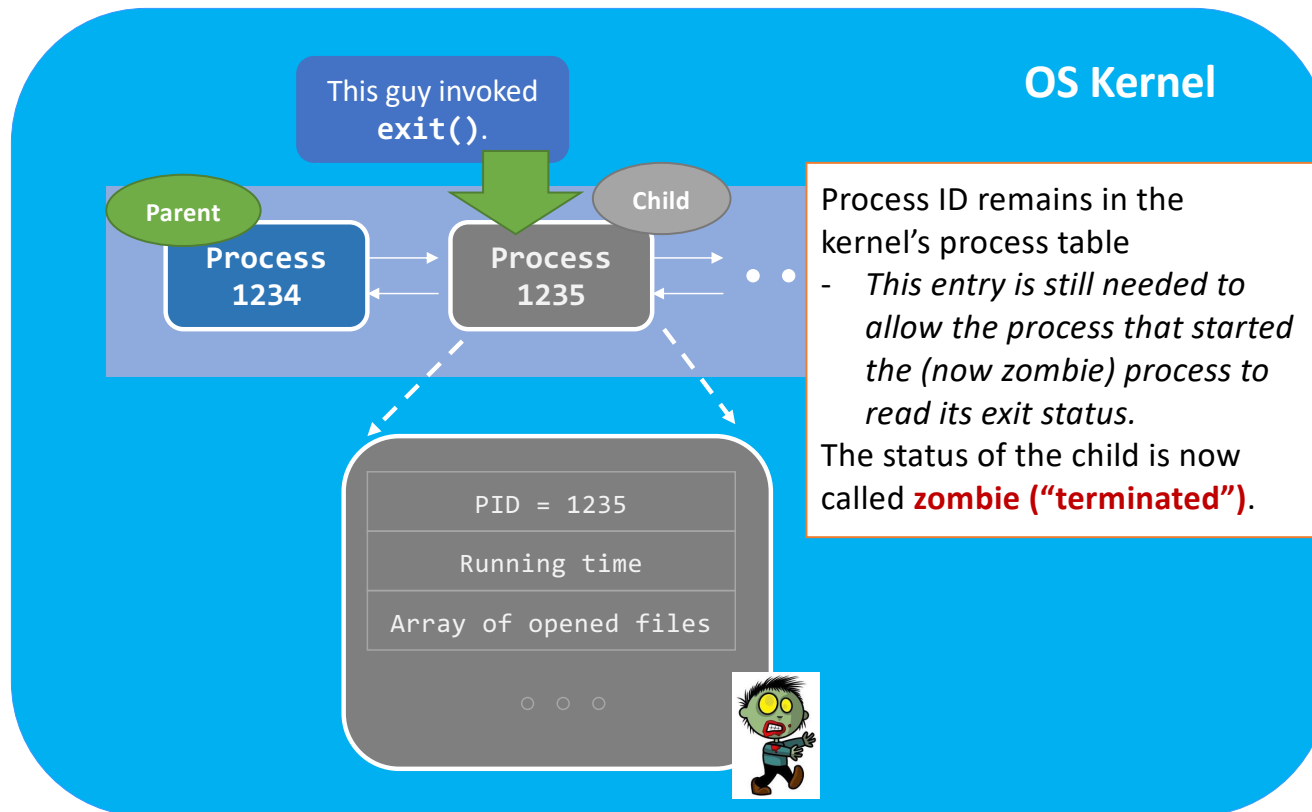
# exit(): Kernel View



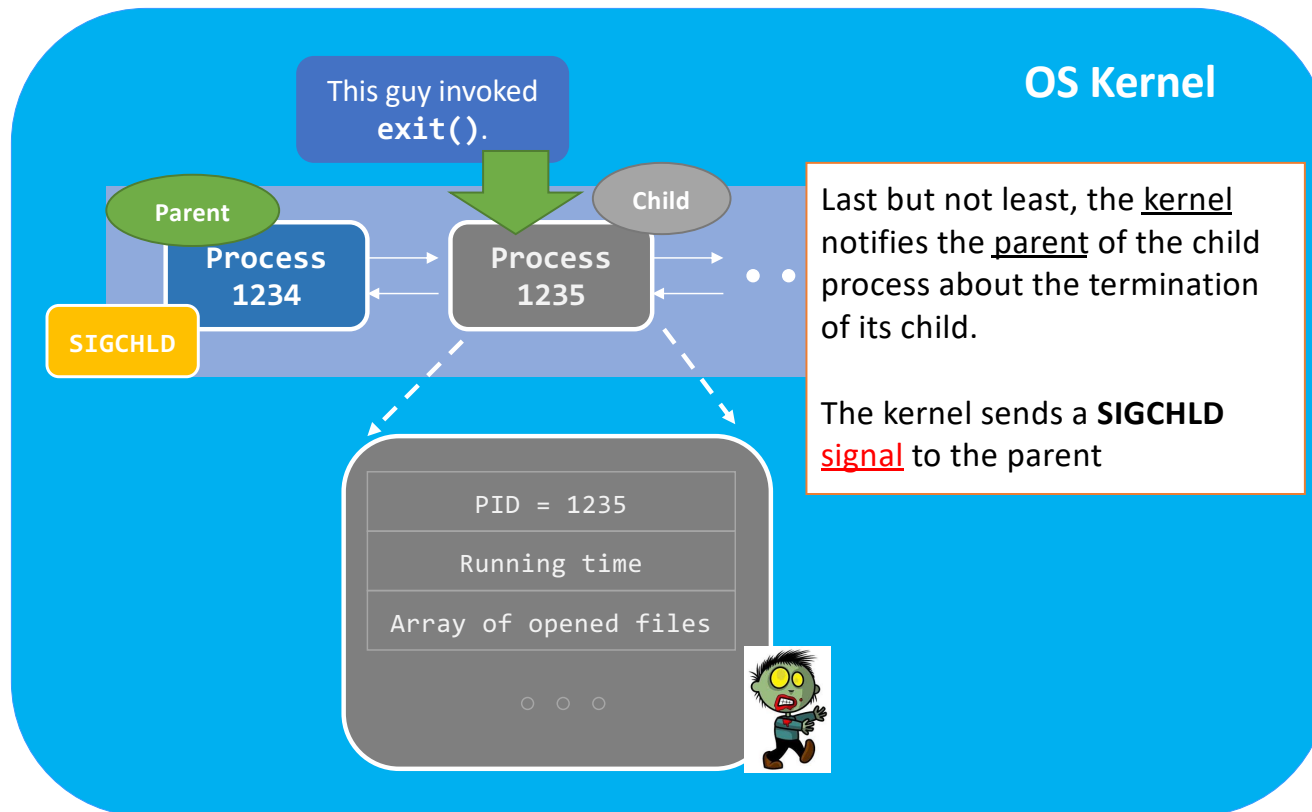
# exit(): Kernel View



# exit(): Kernel View



# exit(): Kernel View




# exit(): Summary

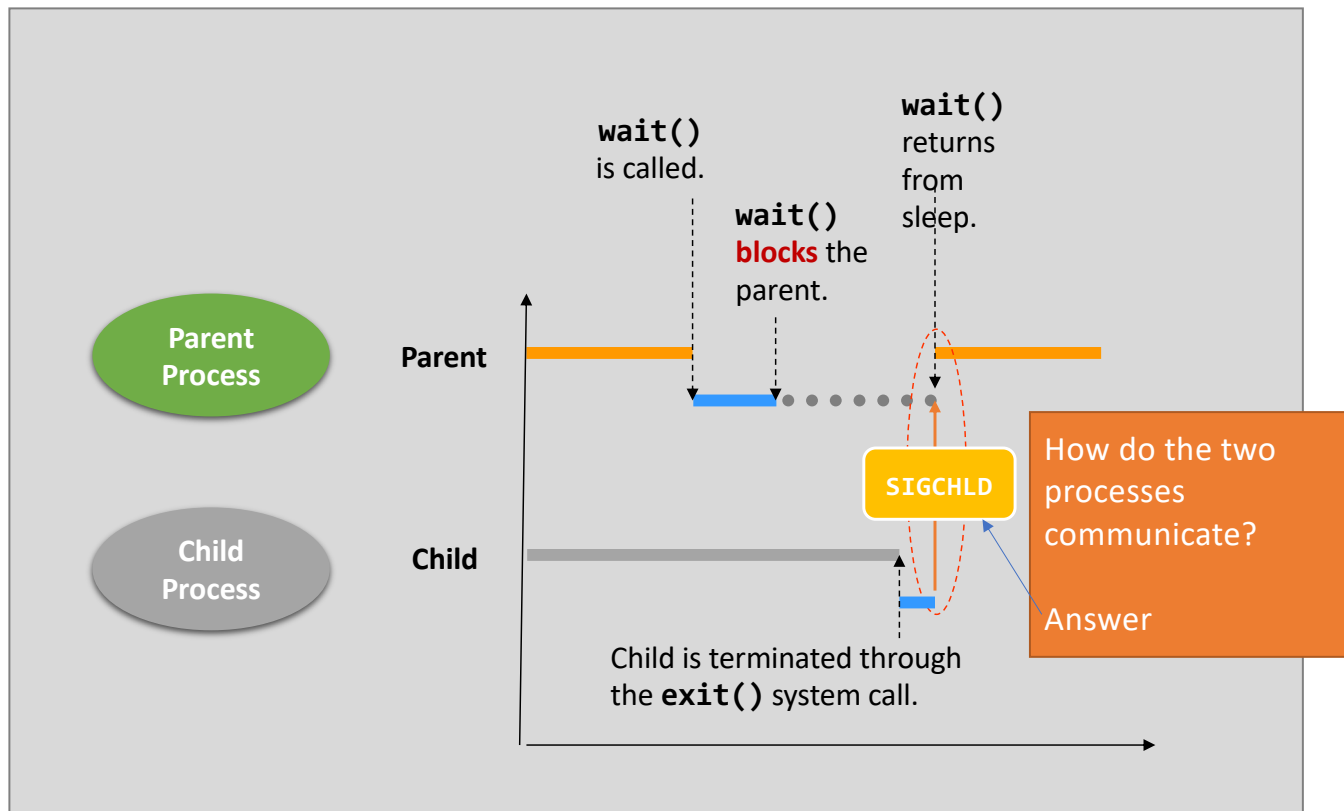
Step (1) Clean up most of the allocated kernel-space memory (e.g., process's running time info).

Step (2) Clean up the exit process's user-space memory.

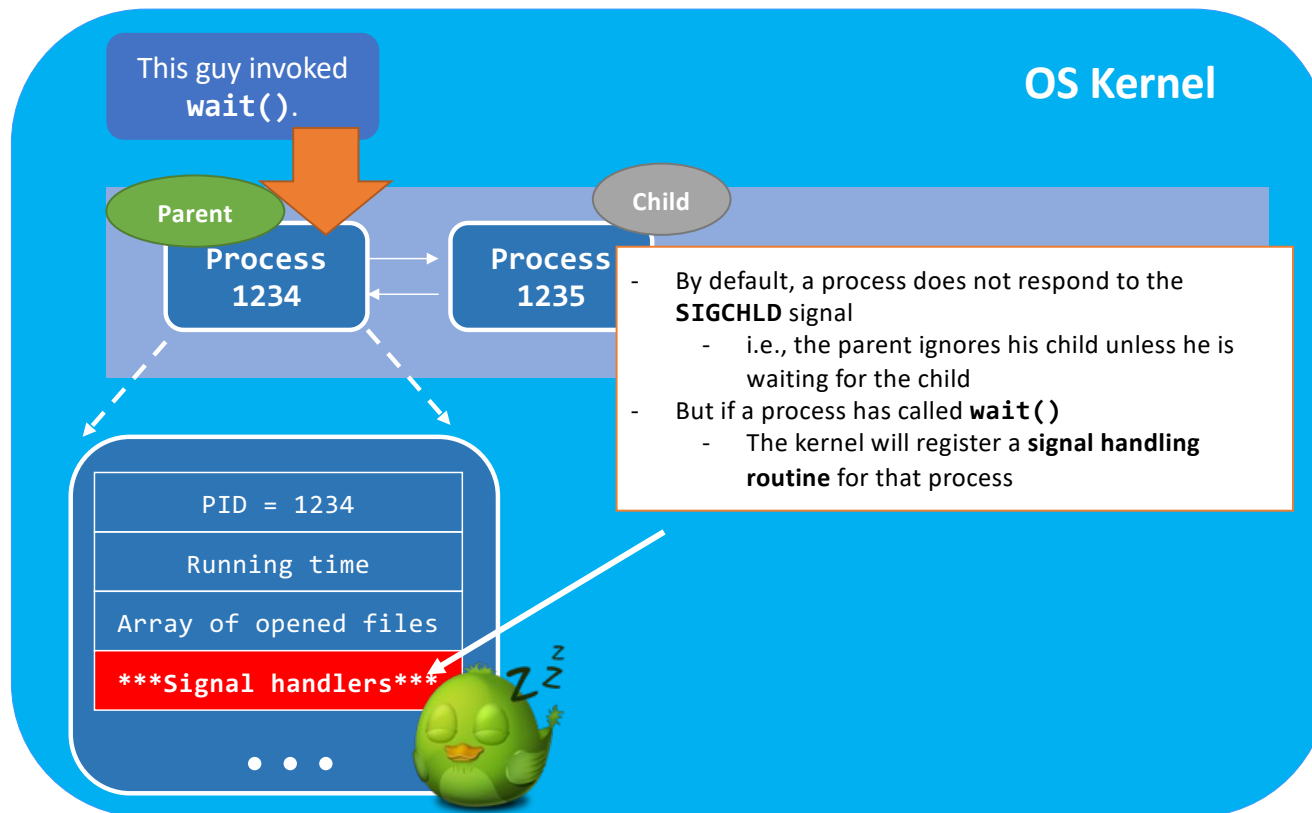
Step (3) Notify the parent with SIGCHLD.



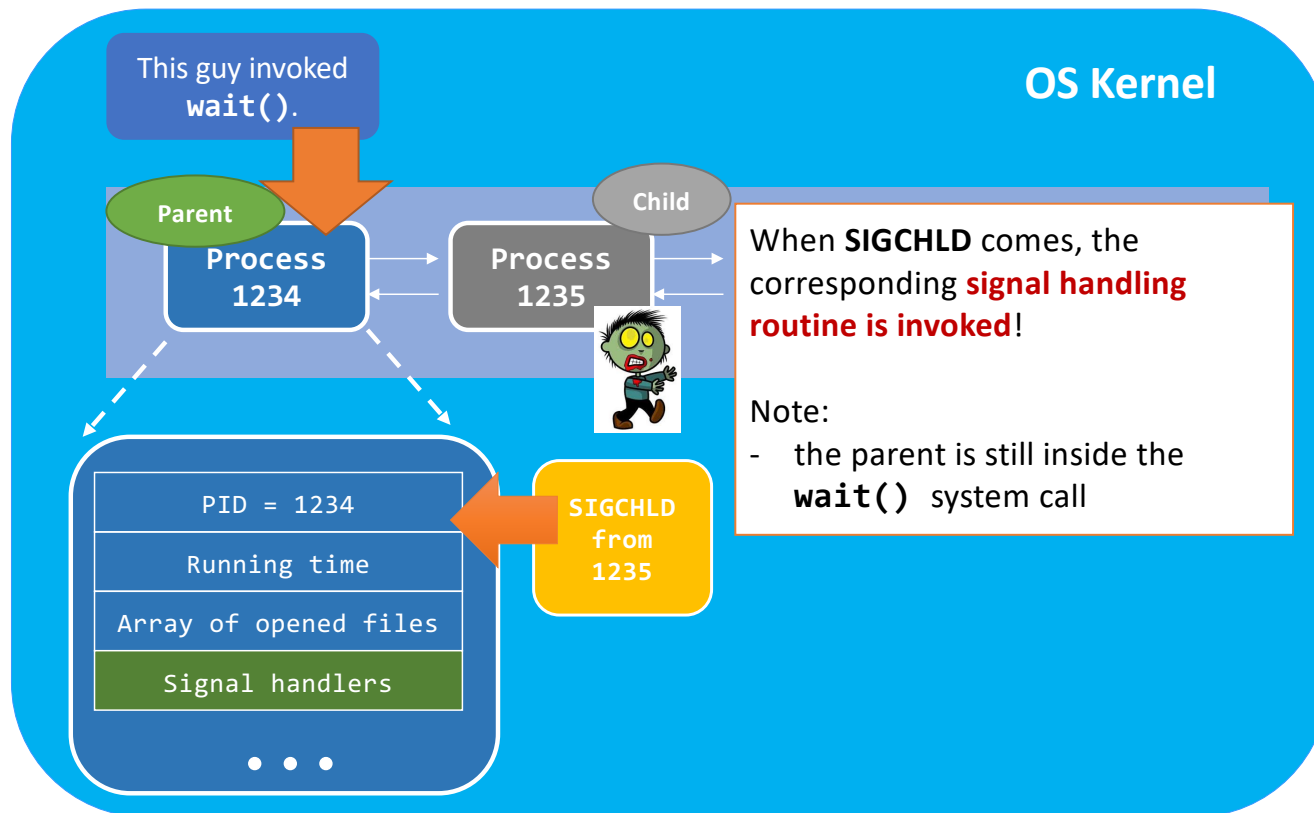
# wait() and exit()



# wait() Kernel View

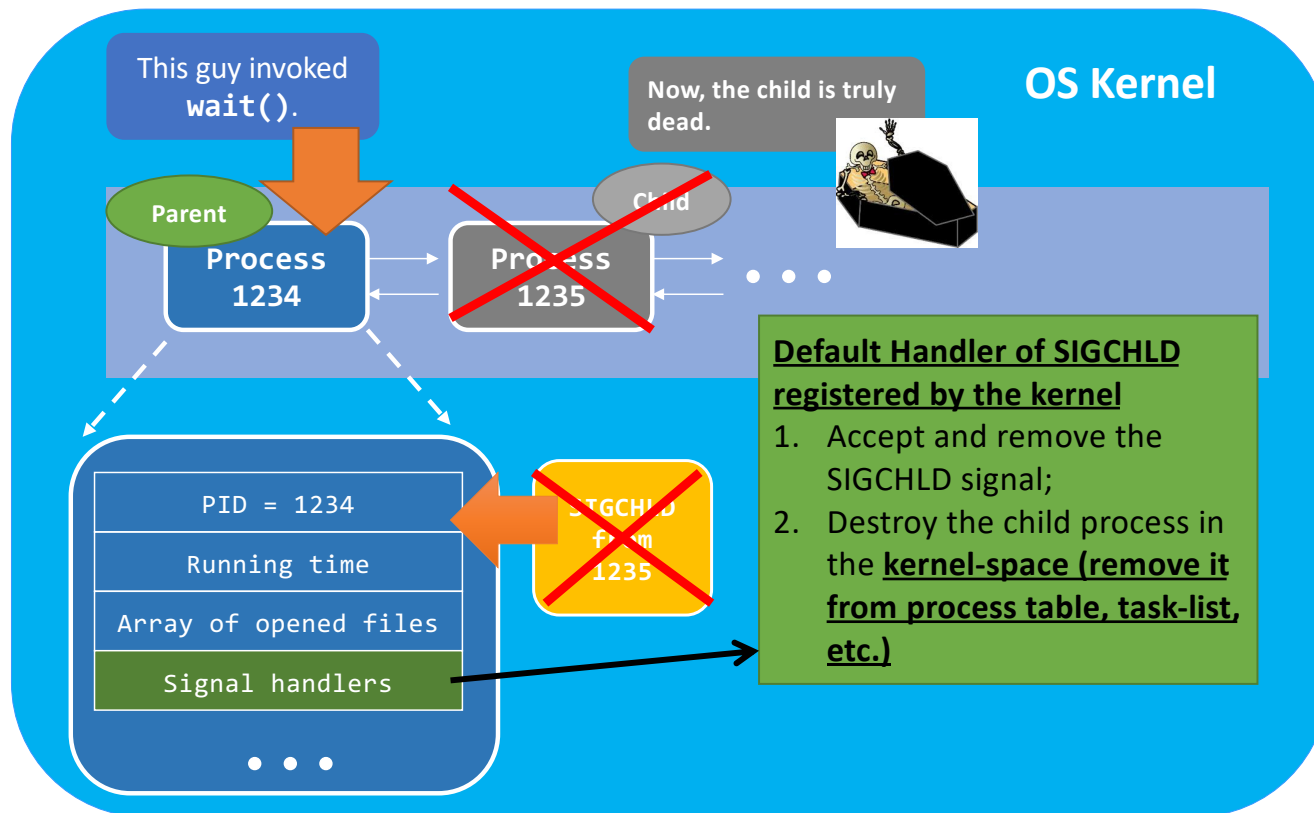


# wait() Kernel View

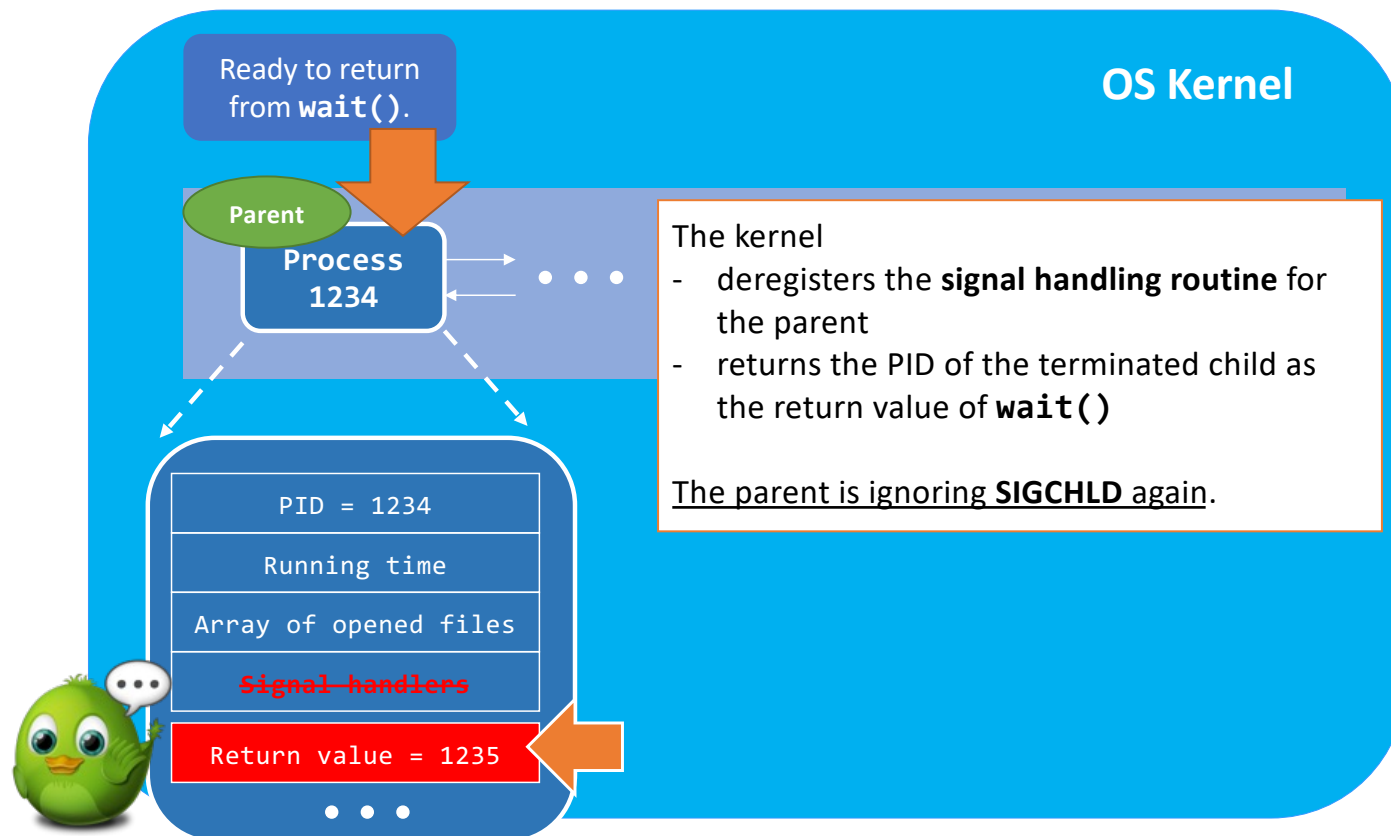




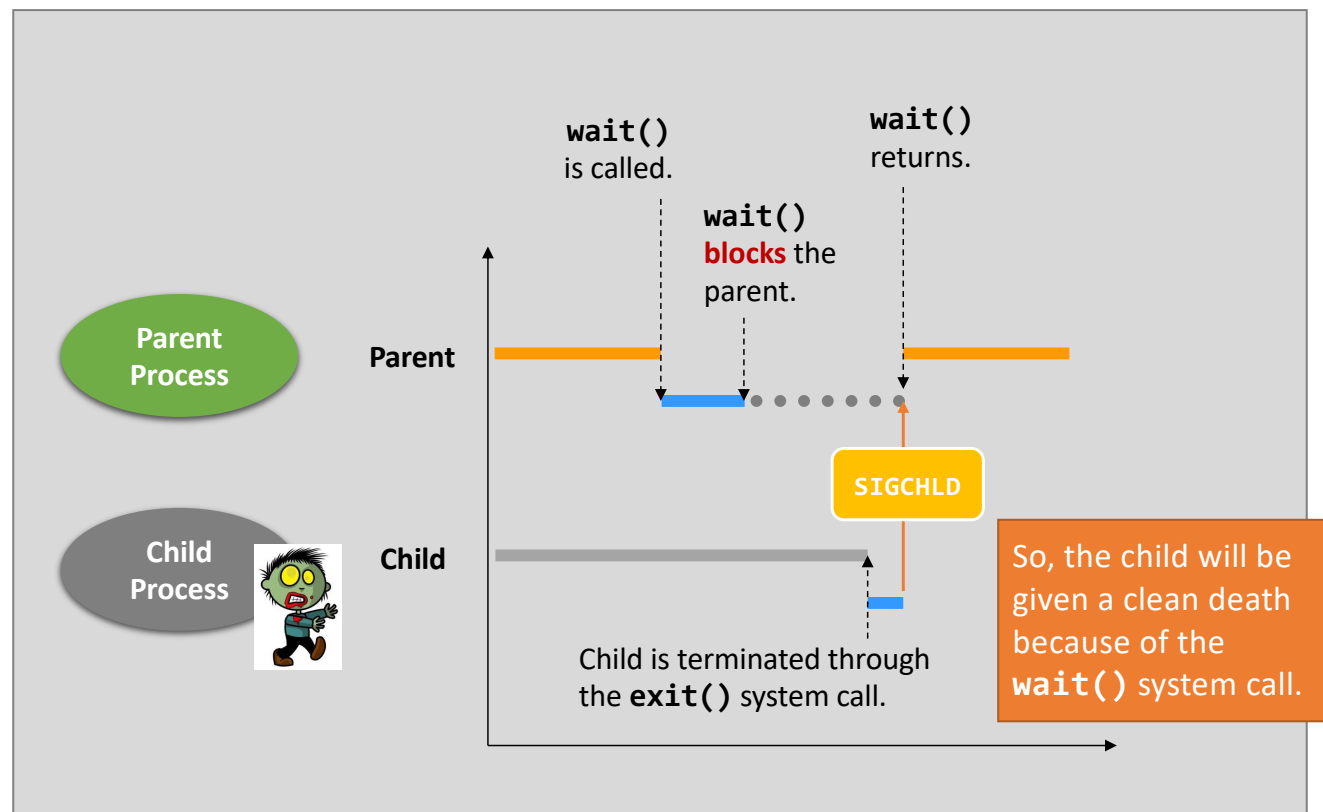
# wait() Kernel View



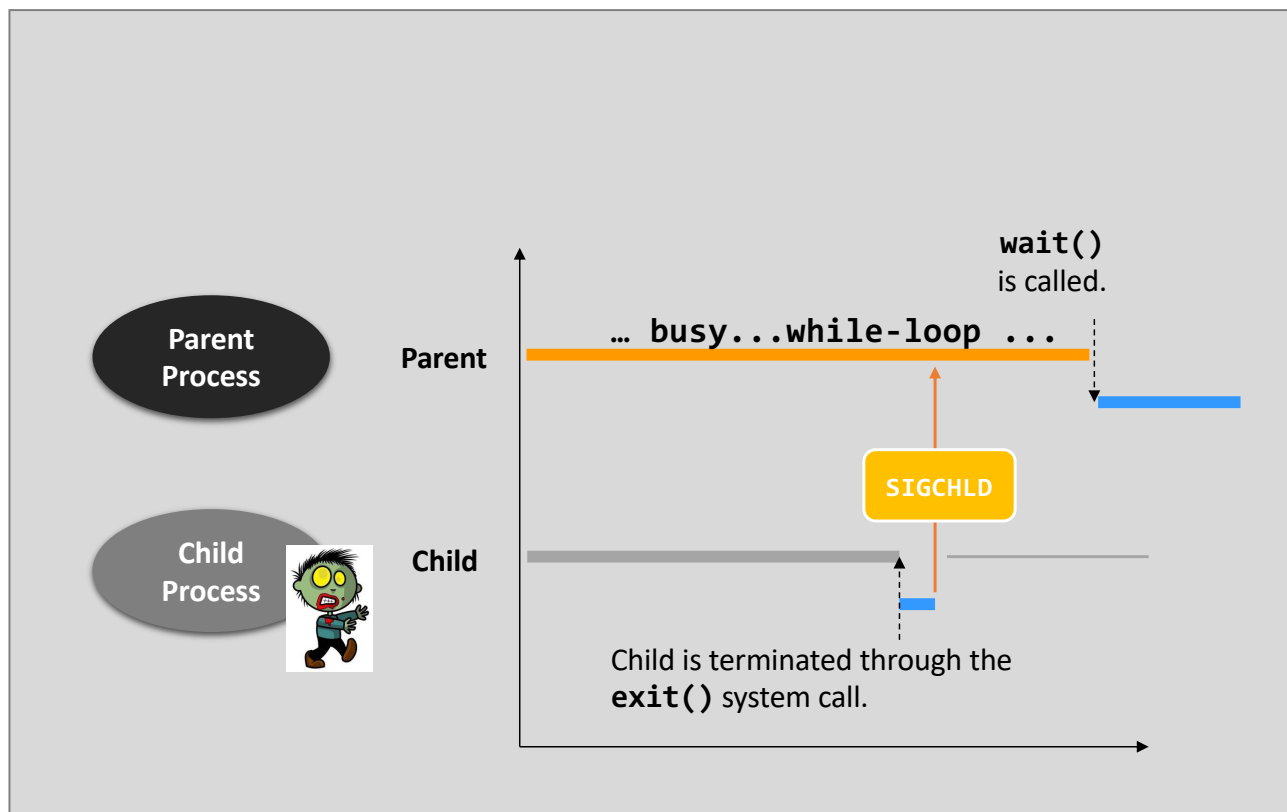
# wait() Kernel View



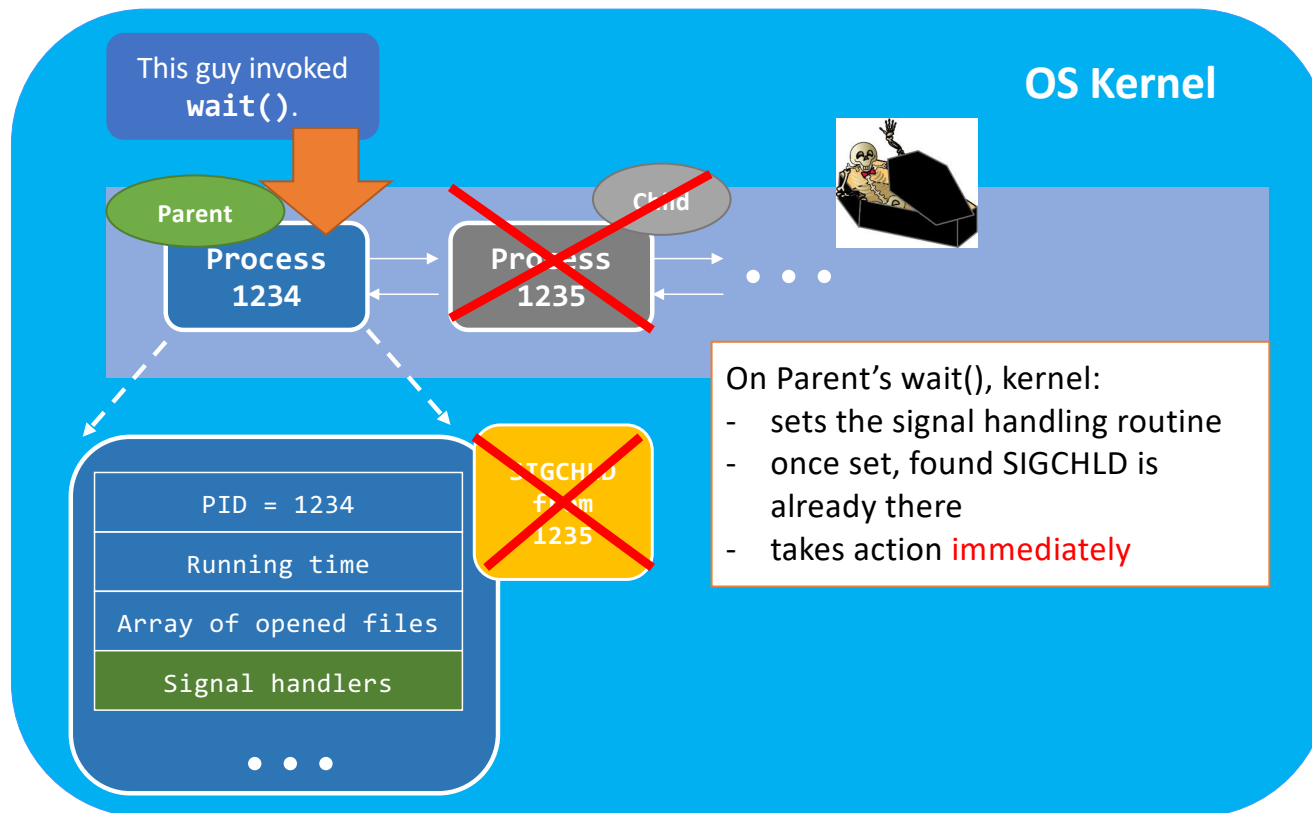
# Normal Case



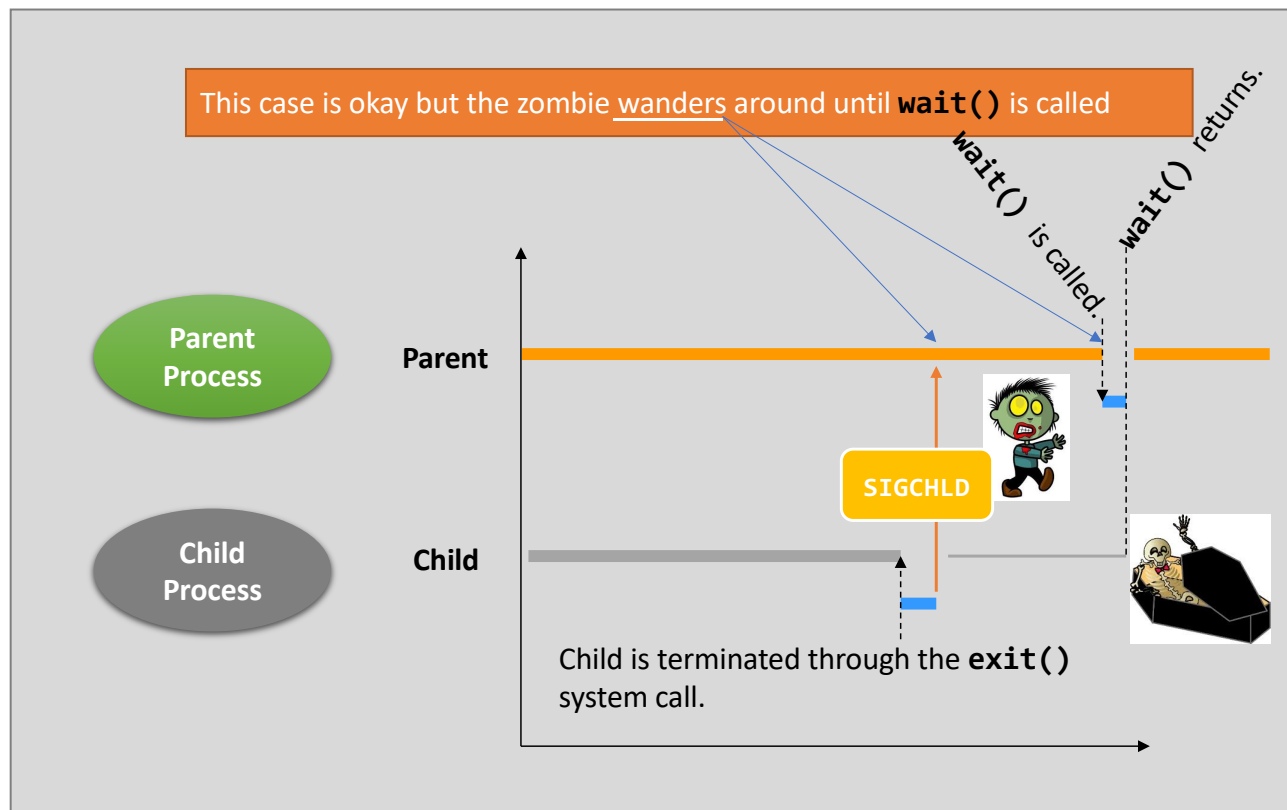
# Parent's wait() after Child's exit()



# Parent's Wait() after Child's exit()



# Parent's wait() after Child's exit()



# Summary of wait() and exit()

- `exit()` system call turns a process into a zombie when...
  - The process calls `exit()`.
  - The process returns from `main()`.
  - The process terminates abnormally.
    - The kernel knows that the process is terminated abnormally. Hence, the kernel invokes `exit()` for it.

# Summary of wait() and exit()

- `wait()` & `waitpid()` reap zombie child processes.
  - It is a must that you should never leave any zombies in the system.
  - `wait()` & `waitpid()` pause the caller until
    - A child terminates/stops, OR
    - The caller receives a signal (i.e., the signal interrupted the `wait()`)
- Linux will label zombie processes as "<defunct>".
  - To look for them:

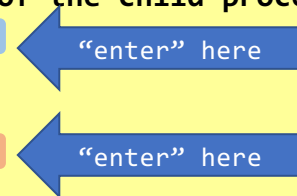
```
$ ps aux | grep defunct
..... 3150 ... [ls] <defunct>
$ _
```

PID of the  
process



# Summary of wait() and exit()

```
1 int main(void)
2 {
3     int pid;
4     if( (pid = fork()) !=0 ) {
5         printf("Look at the status of the child process %d\n", pid);
6         while( getchar() != '\n' );
7         wait(NULL);
8         printf("Look again!\n");
9         while( getchar() != '\n' );
10    }
11    return 0;
12 }
```



This program requires you to type “enter” twice before the process terminates.

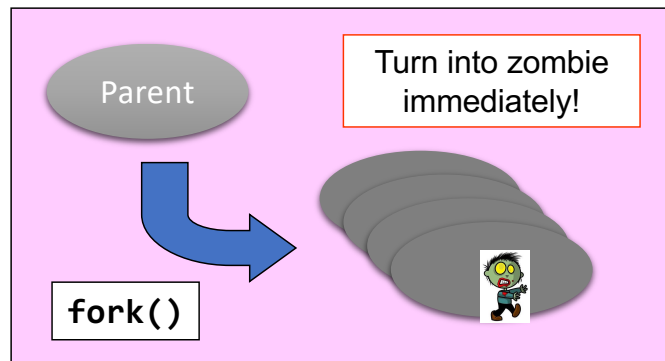
You are expected to see **the status of the child process changes (ps aux [PID])** between the 1<sup>st</sup> and the 2<sup>nd</sup> “enter”.

# Using wait() for Resource Management

- It is not only about process execution / suspension...
- It is about system resource management.
  - A zombie takes up a PID;
  - The total number of PIDs are limited;
    - Read the limit: `cat /proc/sys/kernel/pid_max`
    - It is 32,768.
  - What will happen if we don't clean up the zombies?

# Using wait() for Resource Management

```
int main(void) {  
    while( fork() );  
    return 0;  
}
```



```
$ ./interesting
```

```
-
```

Terminal A

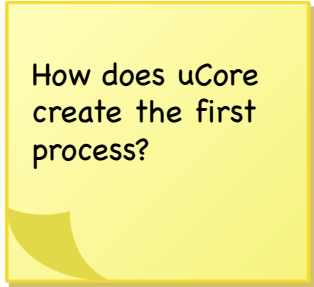
```
$ ls  
No process left.  
$ poweroff  
No process left.  
$ ==  
No process left.  
$ -
```

Terminal B

# More about Processes

# The first process

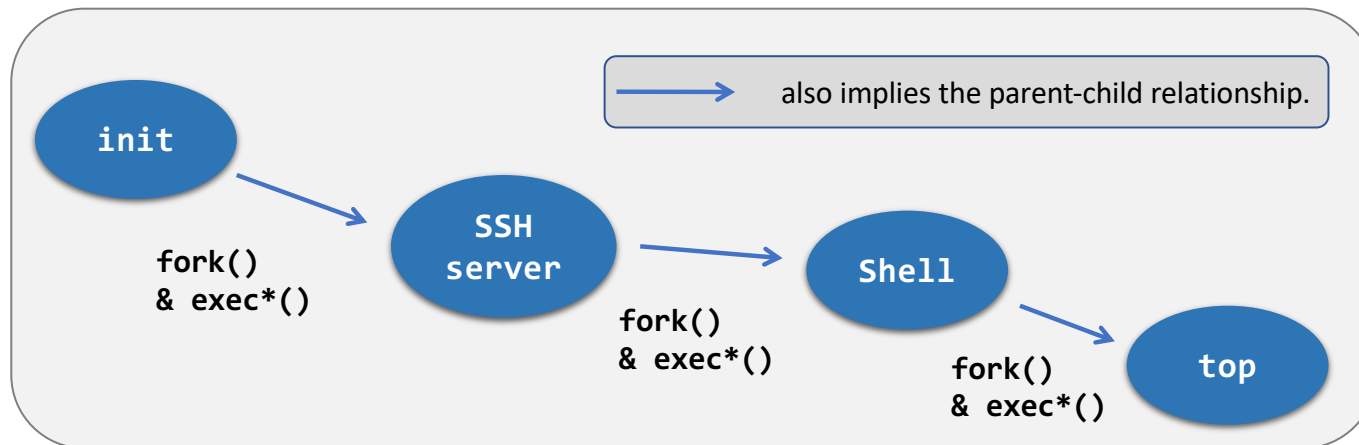
- We now focus on the process-related events.
  - The kernel, while it is booting up, creates the first process – init.
- The “init” process:
  - has PID = 1, and
  - is running the program code “/sbin/init”.
- Its first task is to create more processes...
  - Using fork() and exec().



How does uCore  
create the first  
process?

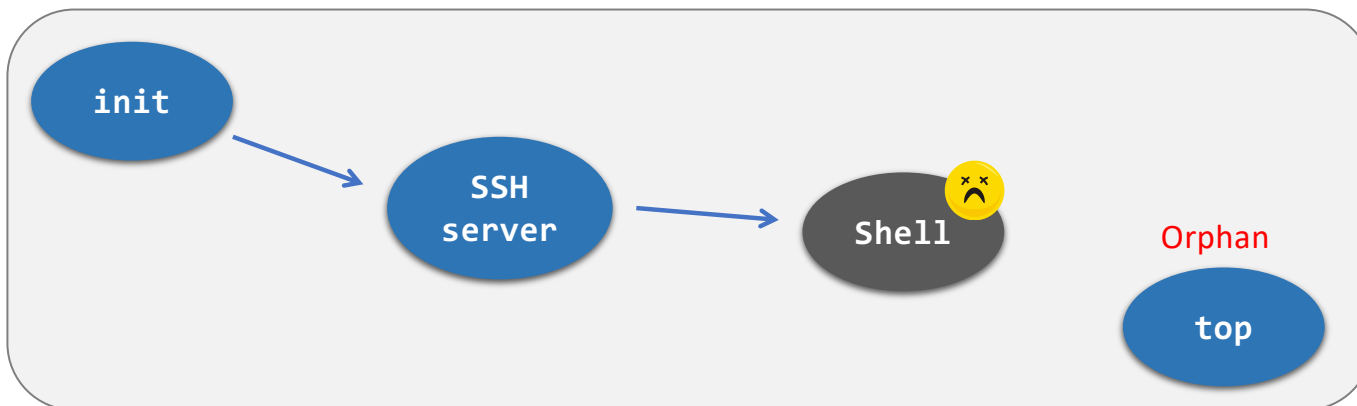
# A Tree of Processes

- You can view the tree with the command:
  - “pstree”; or
  - “pstree -A” for ASCII-character-only display.



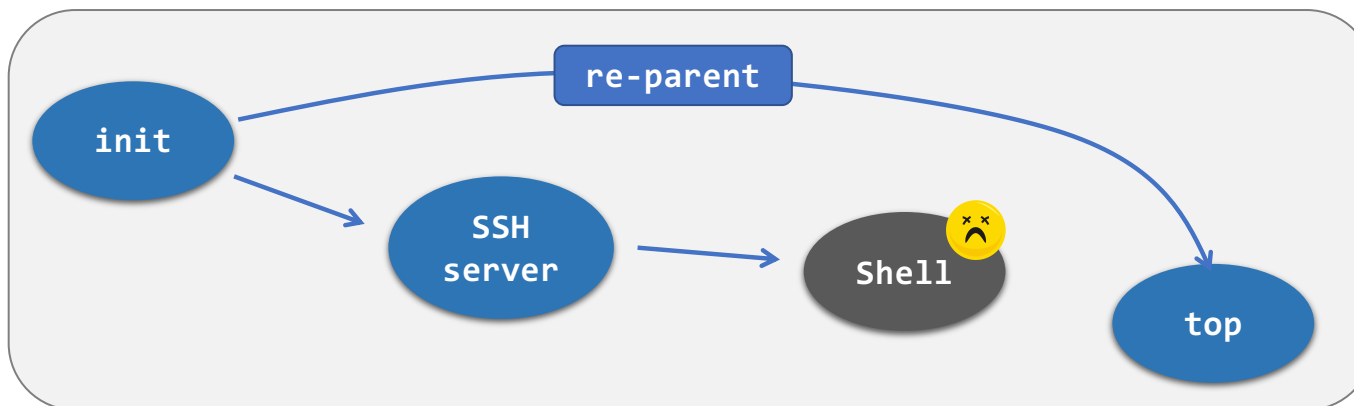
# Orphans

- However, termination can happen, at any time and in any place...
  - This is no good because an orphan turns the hierarchy from a **tree** into a **forest**!
  - Plus, no one would know the termination of the orphan.



# Re-parent

- In Linux
  - The “init” process will become the step-mother of all orphans
  - It's called **re-parenting**
- In Windows
  - It maintains a *forest-like process hierarchy*.....





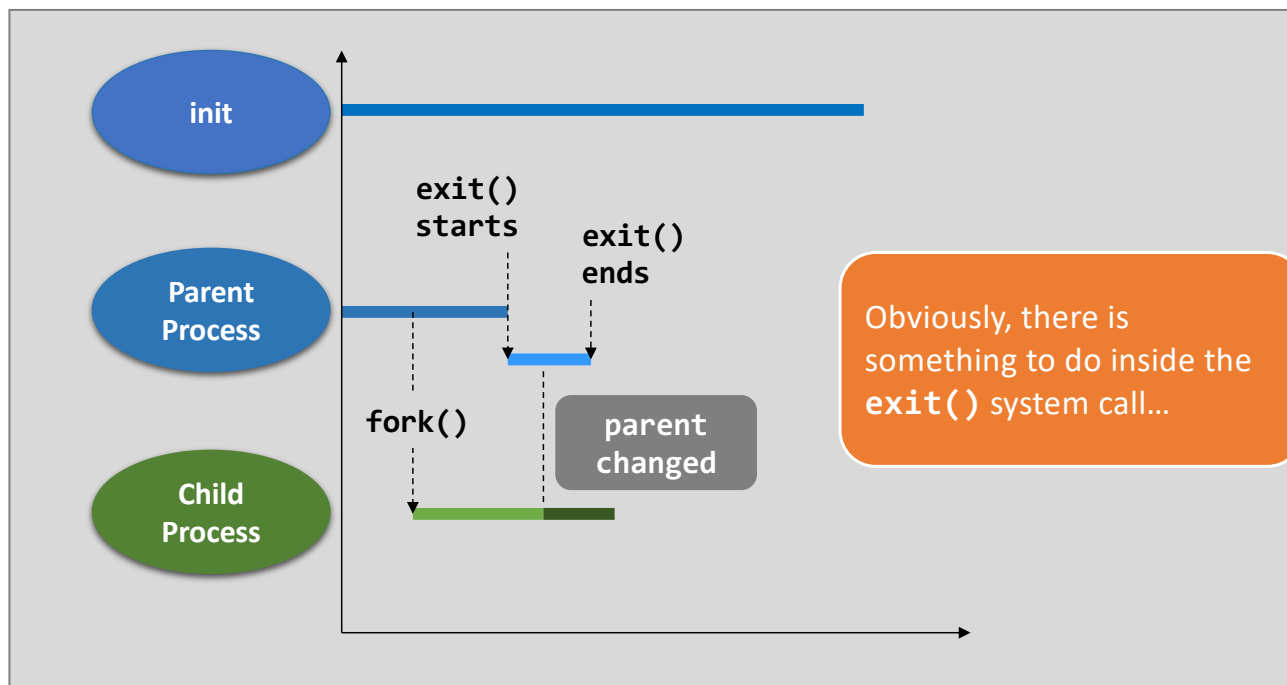
# An Example

```
1  int main(void) {
2      int i;
3      if(fork() == 0) {
4          for(i = 0; i < 5; i++) {
5              printf("(%d) parent's PID = %d\n",
6                  getpid(), getppid() );
7              sleep(1);
8          }
9      }
10     else
11         sleep(1);
12     printf("(%d) bye.\n", getpid());
13 }
```

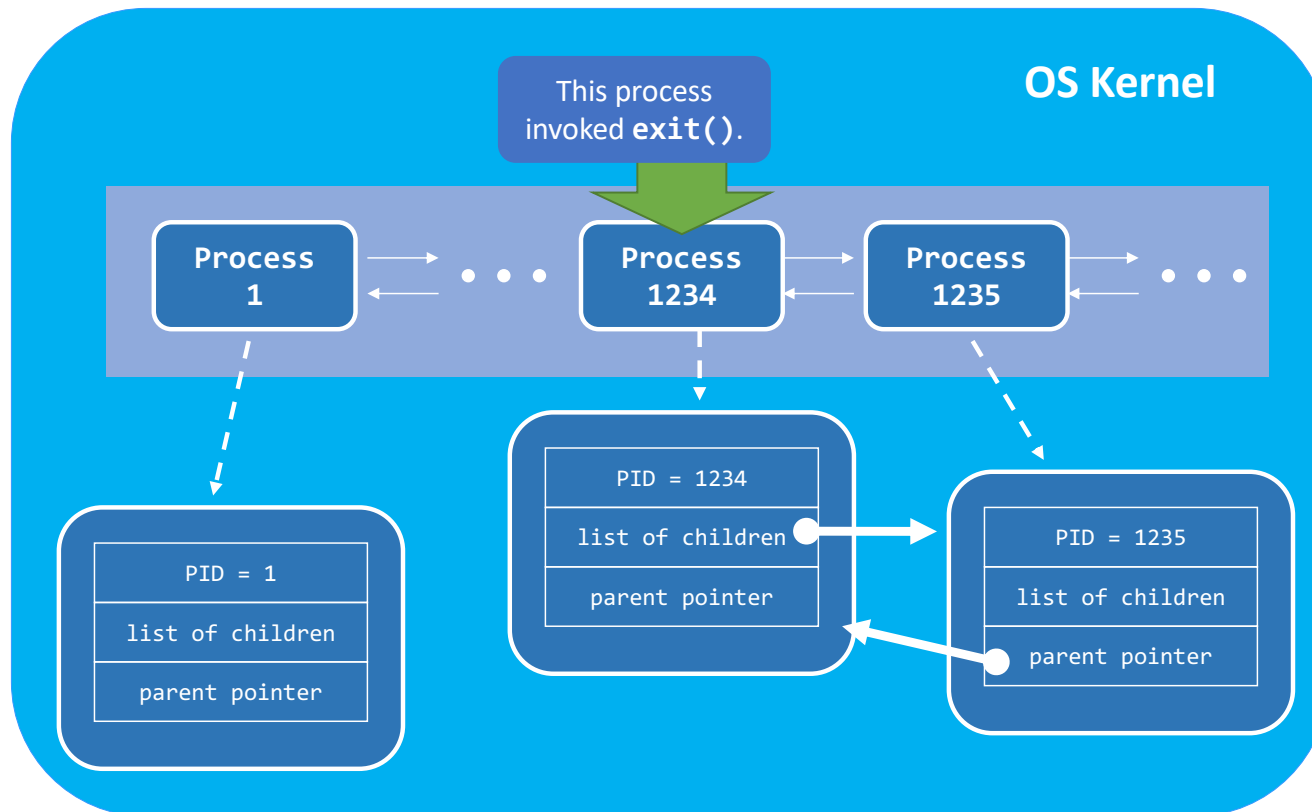
**getppid()** is the system call that returns the parent's PID of the calling process.

```
$ ./reparent
(1235) parent's PID = 1234
(1235) parent's PID = 1234
(1234) bye.
$ (1235) parent's PID = 1
(1235) parent's PID = 1
(1235) parent's PID = 1
(1235) bye.
$ _
```

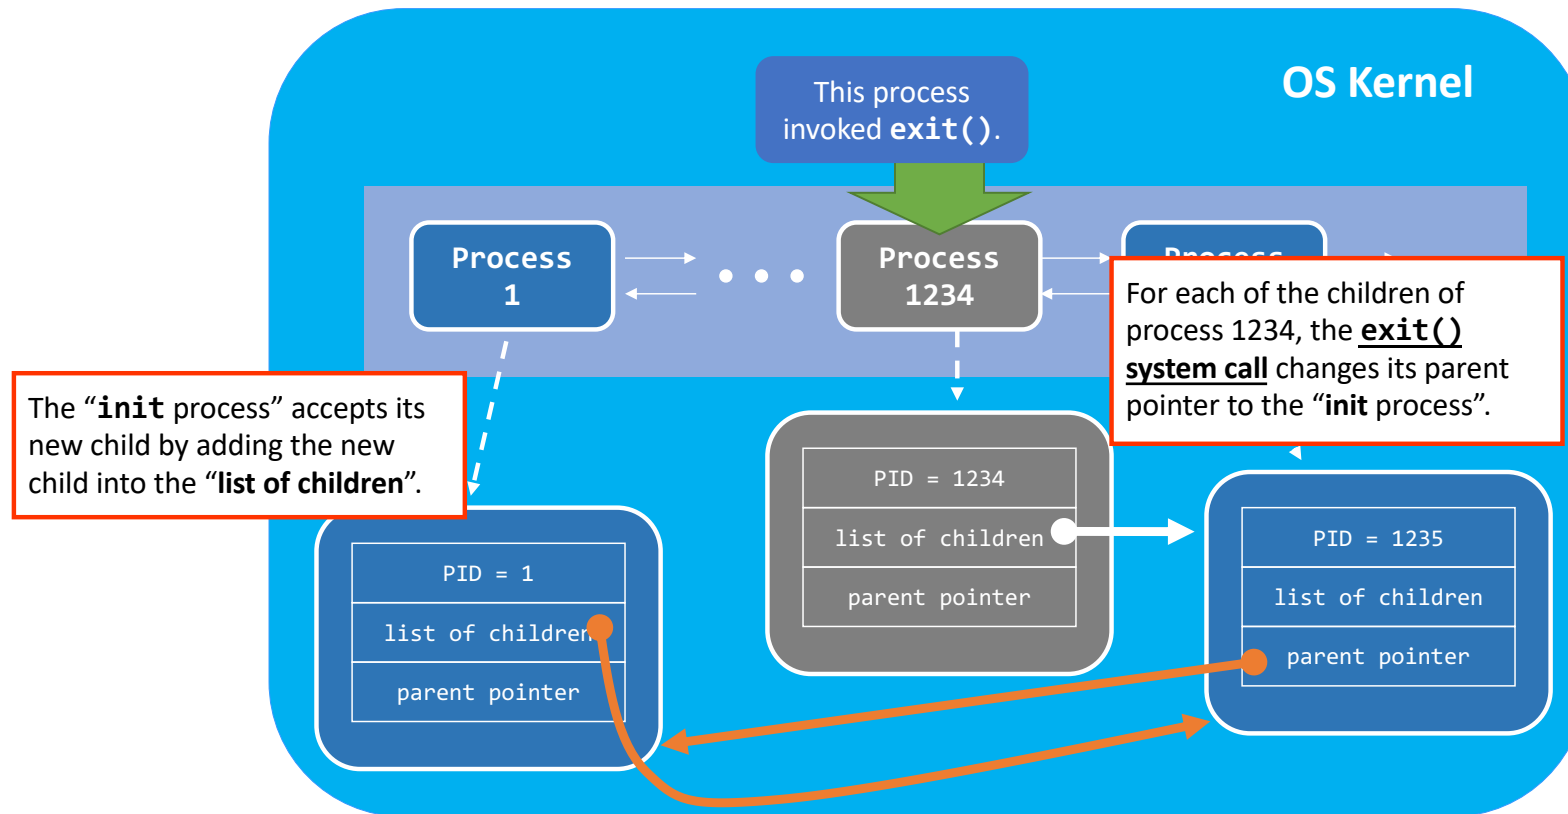
# Re-parenting Explained



# Re-parenting Explained (Cont'd)



# Re-parenting Explained (Cont'd)



# Background Jobs


- The re-parenting operation enables something called **background jobs** in Linux
  - It allows a process runs **without a parent terminal/shell**

[Back to home](#)

```
$ ./infinite_loop &  
$ exit  
  
[ The shell is gone ]
```

```
$ ps -C infinite_loop  
PID  TTY  
1234  ... ./infinite_loop  
$ _
```

# Measure Process Time



# User Time v.s. System Time (Case 1)

```
$ time ./time_example
```

```
real    0m0.001s
user    0m0.000s
sys     0m0.000s
$ _
```

Real-time elapsed when “./time\_example” terminates.

The user time of “./time\_example”.

The sys time of “./time\_example”.

It's possible:  
real > user + sys  
real < user + sys

Why?

- real > user + sys  
  **I/O intensive**
- real < user + sys  
  **multi-core**

```
int main(void) {
    int x = 0;
    for(i = 1; i <= 10000; i++) {
        x = x + i;
        // printf("x = %d\n", x);
    }
    return 0;
}
```

# User Time v.s. System Time (Case 1)

```
$ time ./time_example
```

```
real    0m0.001s
user    0m0.000s
sys     0m0.000s
$ _
```

```
int main(void) {
    int x = 0;
    for(i = 1; i <= 10000; i++) {
        x = x + i;
        // printf("x = %d\n", x);
    }
    return 0;
}
```

Commented on purpose.

```
$ time ./time_example
```

```
real 0m2.795s
user 0m0.084s
sys 0m0.124s
$ _
```

See? Accessing hardware costs the process more time.

```
int main(void) {
    int x = 0;
    for(i = 1; i <= 10000; i++) {
        x = x + i;
        printf("x = %d\n", x);
    }
    return 0;
}
```



# User Time v.s. System Time (Case 2)

- The user time and the sys time together **define the performance of an application.**
  - When writing a program, you must consider both the user time and the sys time.
    - E.g., the output of the following two programs are exactly the same. But, their running time is not.

```
#define MAX 1000000

int main(void) {
    int i;
    for(i = 0; i < MAX; i++)
        printf("x\n");
    return 0;
}
```

```
#define MAX 1000000

int main(void) {
    int i;
    for(i = 0; i < MAX / 5 ; i++)
        printf("x\nx\nx\nx\nx\n");
    return 0;
}
```

# User Time v.s. System Time (Case 2)

```
#define MAX 1000000

int main(void) {
    int i;
    for(i = 0; i < MAX; i++)
        printf("x\n");
    return 0;
}
```

```
$ time ./time_example_slow

real 0m1.562s
user 0m0.024s
sys 0m0.108s
$ _
```

```
#define MAX 1000000

int main(void) {
    int i;
    for(i = 0; i < MAX / 5 ; i++)
        printf("x\nx\nx\nx\nx\n");
    return 0;
}
```

```
$ time ./time_example_fast

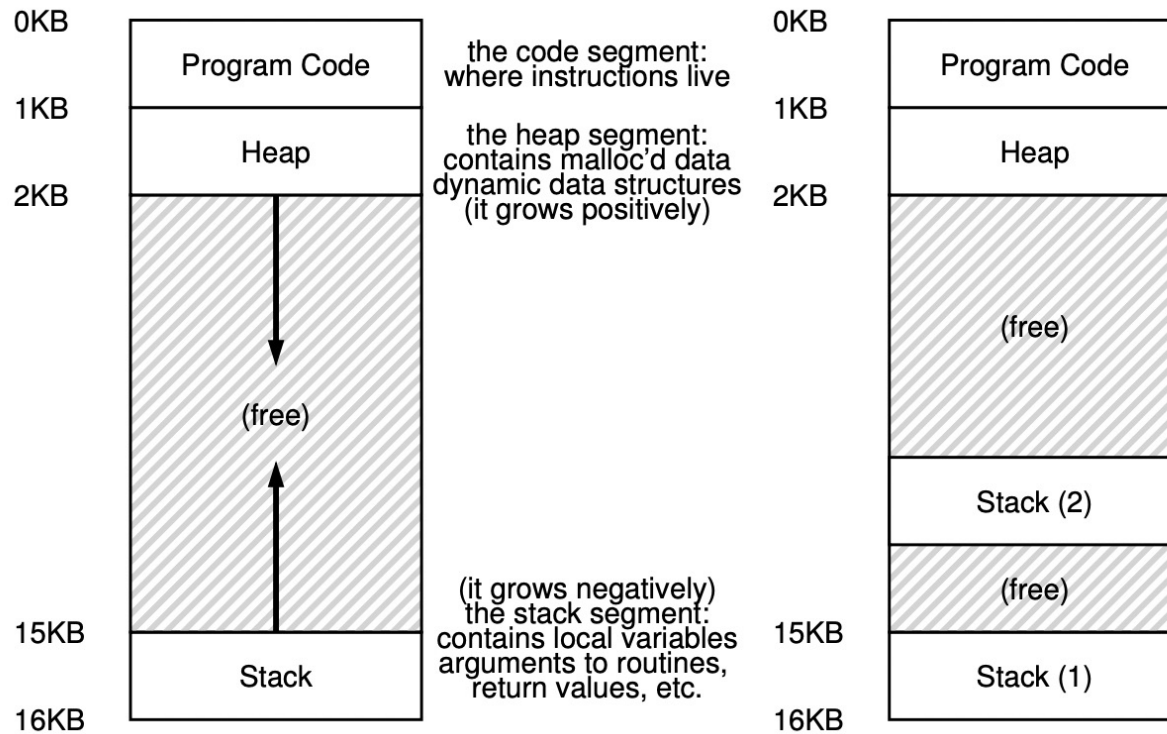
real 0m1.293s
user 0m0.012s
sys 0m0.084s
$ _
```

# Threads

# What is a Thread?

- Thread is an **abstraction** of the execution of a program
  - A single-threaded program has one point of execution
  - A multi-threaded program has more than one points of execution
- Each thread has its own **private** execution state
  - Program counter and a private set of registers
  - A private stack for thread-local storage
  - CPU switching from one thread to another requires context switch
- Threads in the same process **share** computing resources
  - Address space, files, signals, etc.

# Single-Threaded and Multi-Threaded



# Why Use Thread?


- Increase parallelism
  - One thread per CPU makes better use of multiple CPUs to improve efficiency
- Avoid blocking program progress due to slow I/O
  - Threading enables overlap of I/O with other activities within a single program
  - e.g., many modern server-based applications (web servers, database management systems, and the like) make use of threads
- And allow resource sharing !!!

# Thread Implementation


- User-level thread
  - Thread management (e.g., creating, scheduling, termination) done by user-level threads library
  - OS does not know about user-level thread
- Kernel-level thread
  - Thread management done by kernel
  - OS is aware of each kernel-level thread

# Thread Models


- One-to-one mapping
  - One user-level thread to one kernel-level thread
- Many-to-one mapping
  - Many user-level thread to one kernel-level thread
- Many-to-many mapping
  - Many user-level thread to many kernel-level thread



One-to-One



Many-to-One



Many-to-Many



# Pros and Cons

- Many-to-one mapping
  - Pros: context switch between threads is cheap
  - Cons: When one thread blocks on I/O, all threads block
- One-to-one mapping
  - Pros: Every thread can run or block independently
  - Cons: Need to make a crossing into kernel mode to schedule
- Many-to-many mapping
  - Many user-level threads multiplexed on less or equal number of kernel-level threads
  - Pros: best of the two worlds, more flexible
  - Cons: difficult to implement

# Thank you!

